

GETTING STARTED WITH

clickHouse

GoldenClover
Education

www.learn-olapdb.com

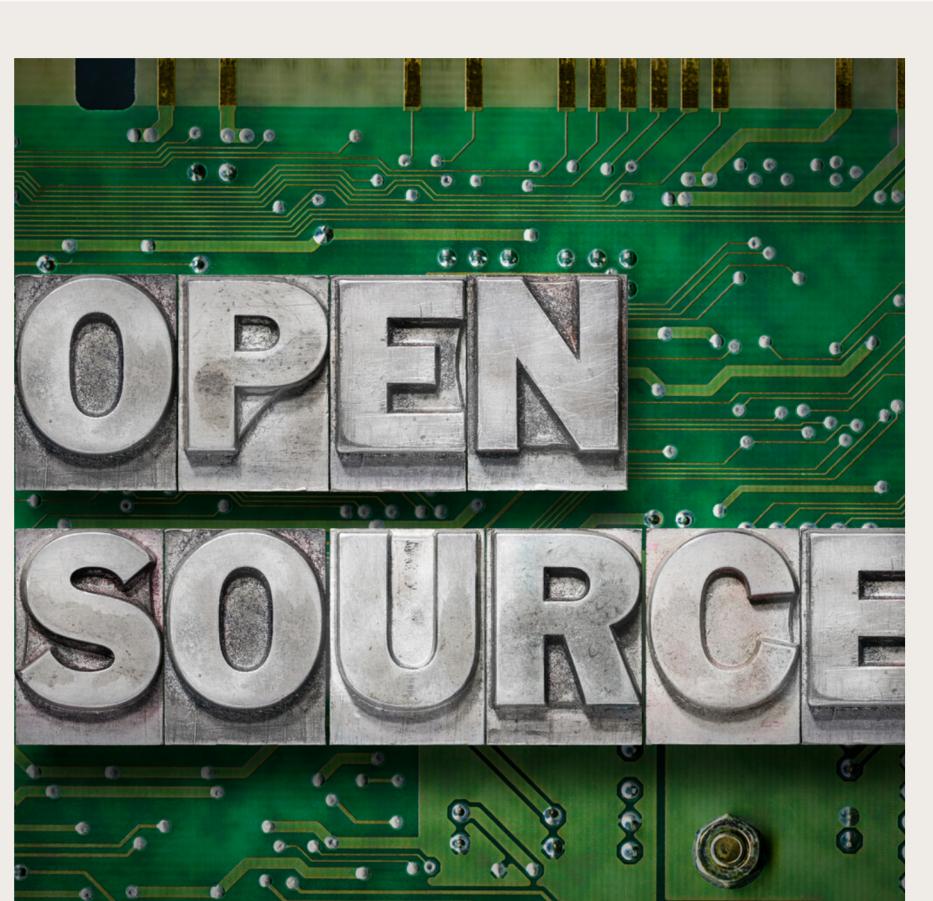






ClickHouse Database

CLICKHOUSE DATABASE



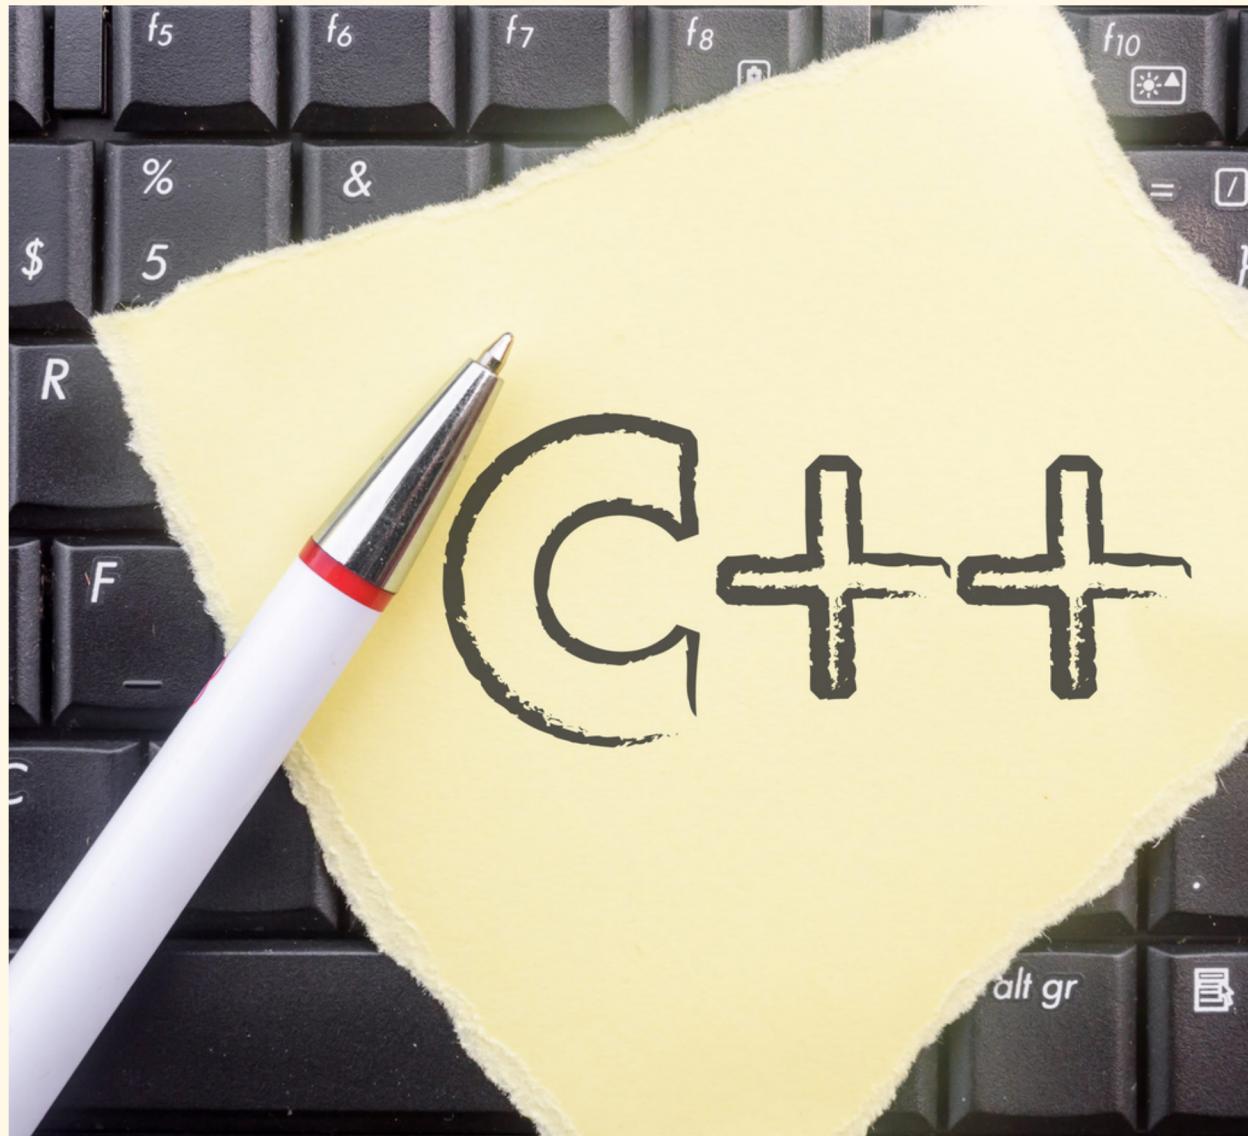
Apache 2.0 Licence



Super Fast



Cost Effective



6K+

Forks

1200+

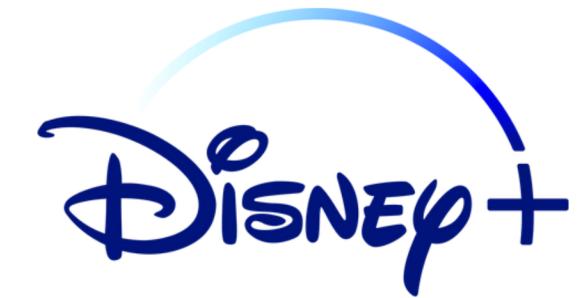
Contributors

35K +

Pull Requests

30K +

Github Stars



All product names, trademarks, and registered trademarks are the property of their respective owners. All company, product, and service names used are for identification purposes only

GETTING STARTED WITH

clickHouse

GoldenClover
Education

www.learn-olapdb.com

HARDWARE

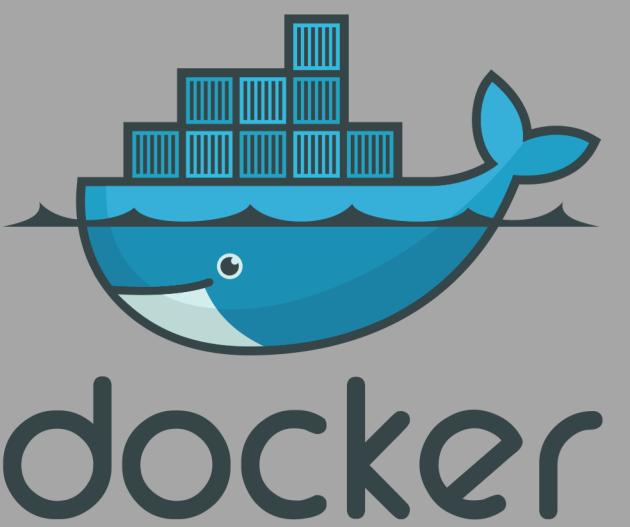
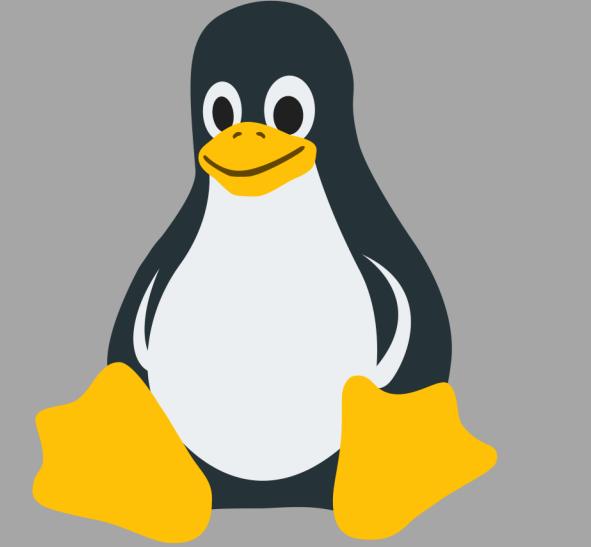
- 8 GB RAM
- 50 GB Hard Disk Space
- Linux (Ubuntu/Debian or Fedora/CentOS) or Mac OSX

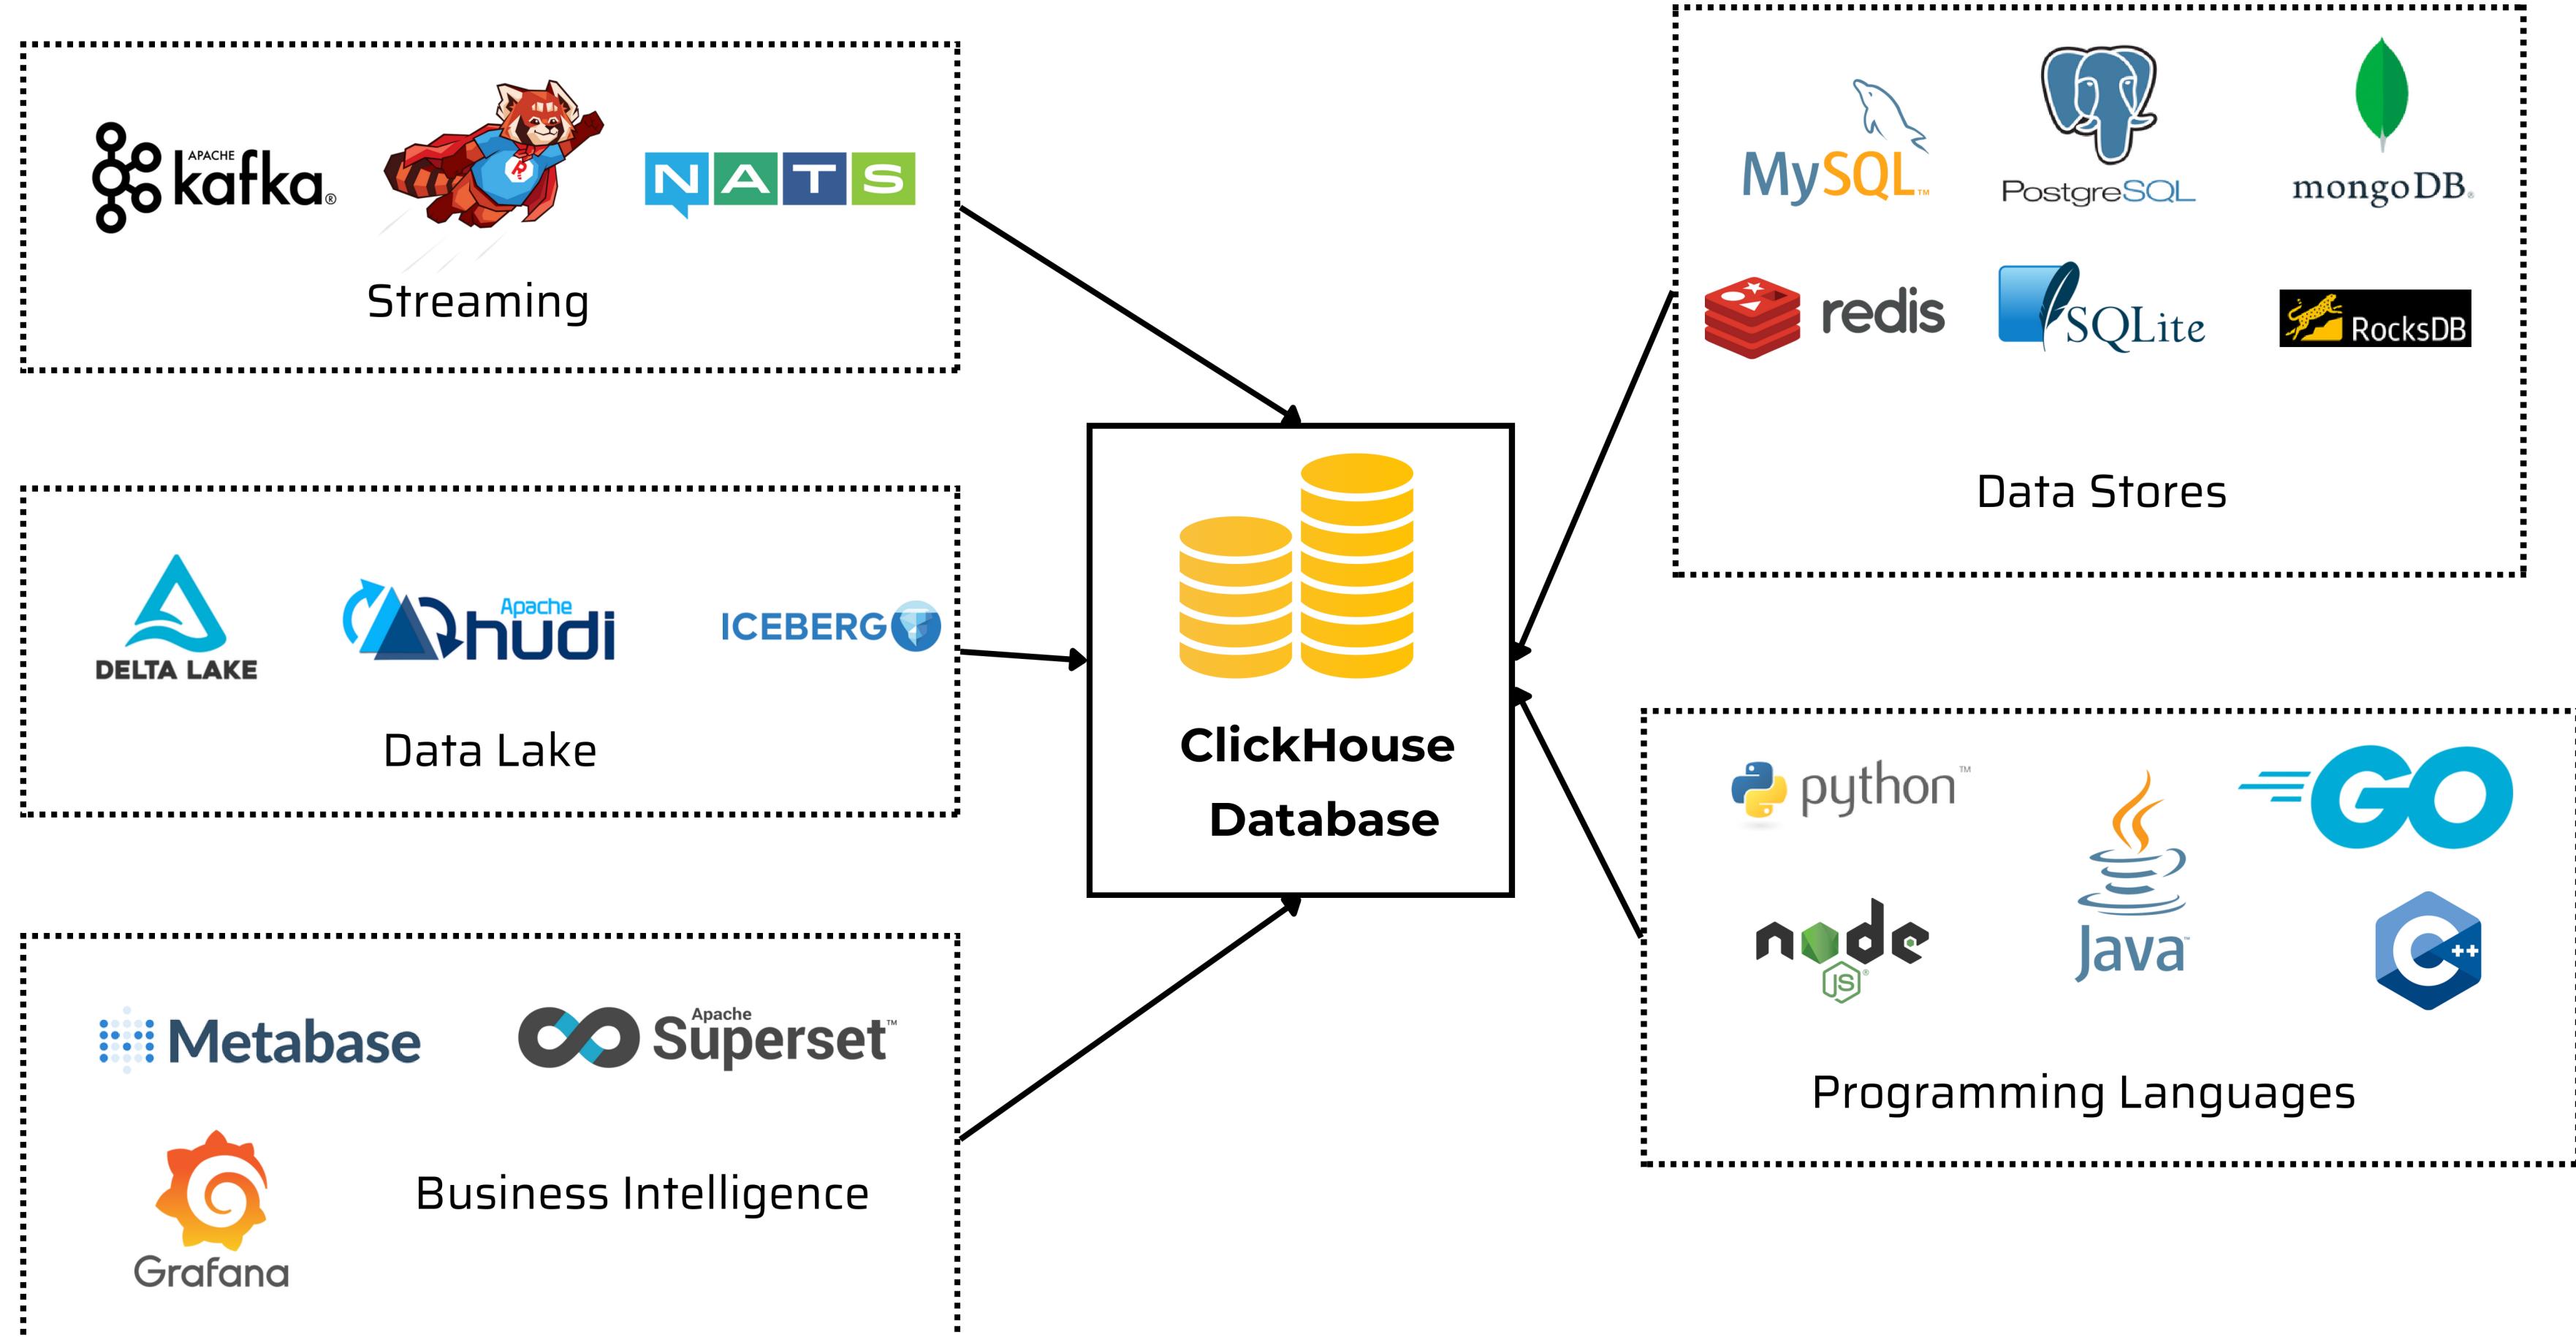


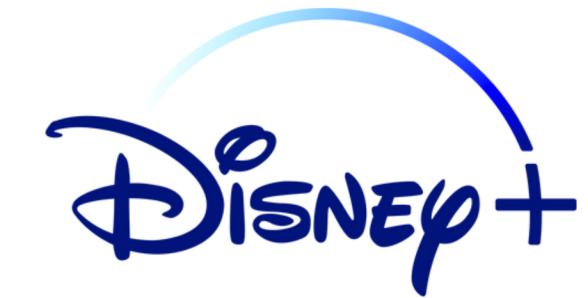


SOFTWARE

- Linux commands
- SQL
- Docker and Docker Compose







All product names, trademarks, and registered trademarks are the property of their respective owners. All company, product, and service names used are for identification purposes only

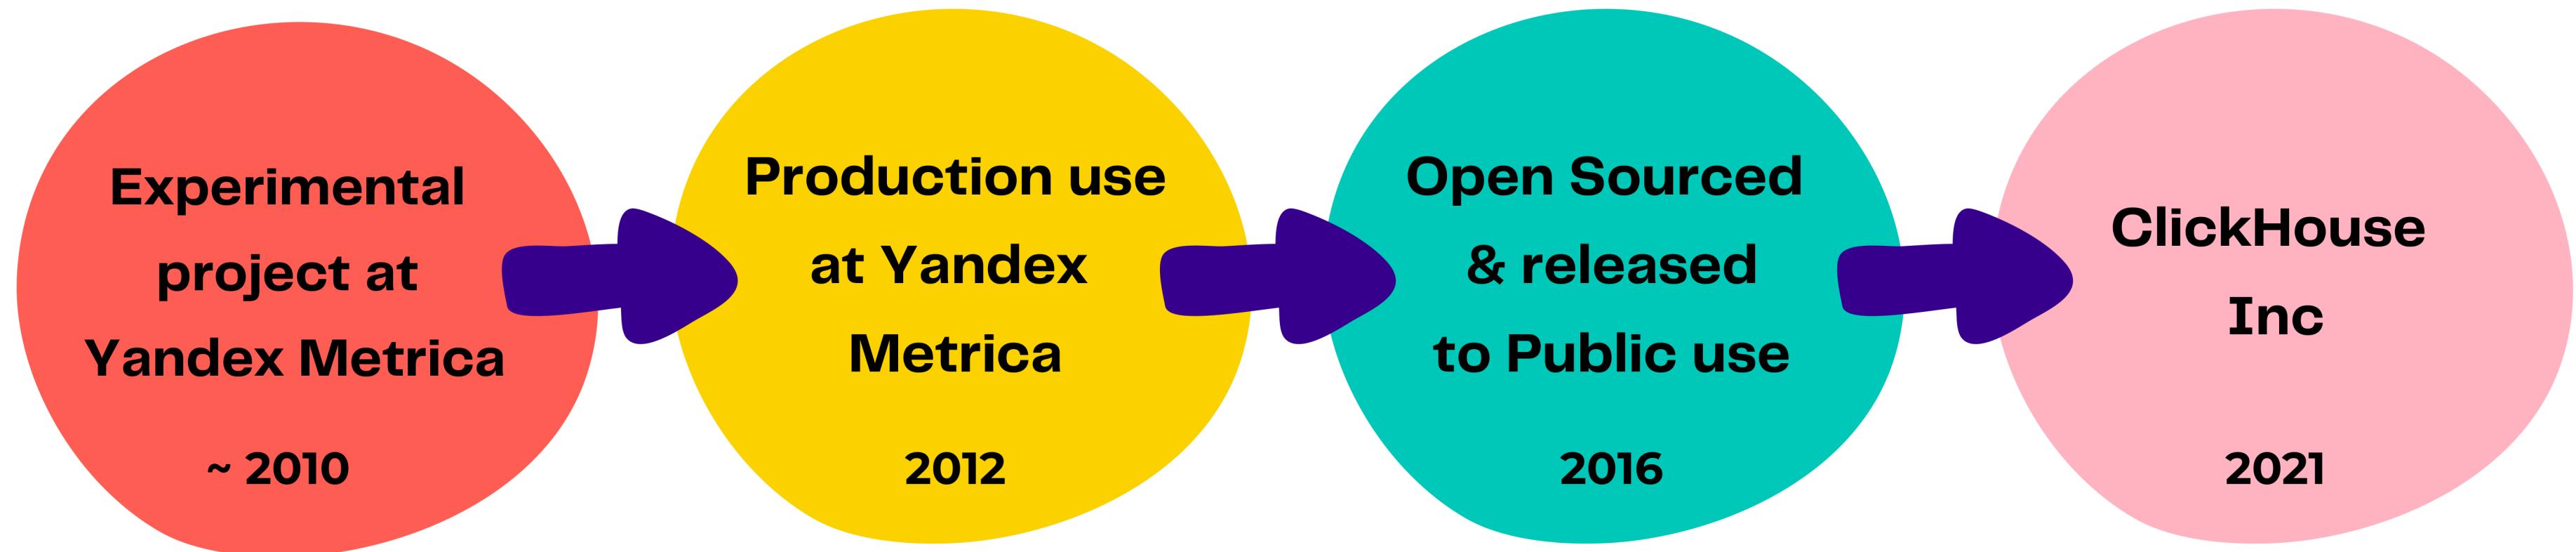
GETTING STARTED WITH

clickHouse

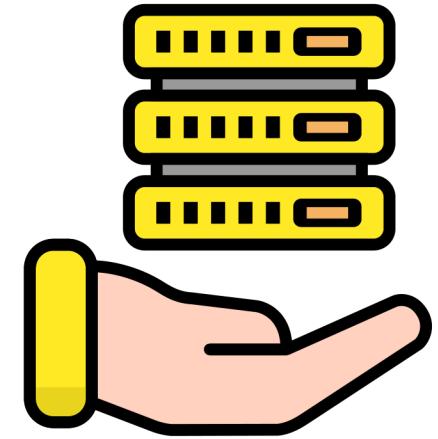
GoldenClover
Education

www.learn-olapdb.com

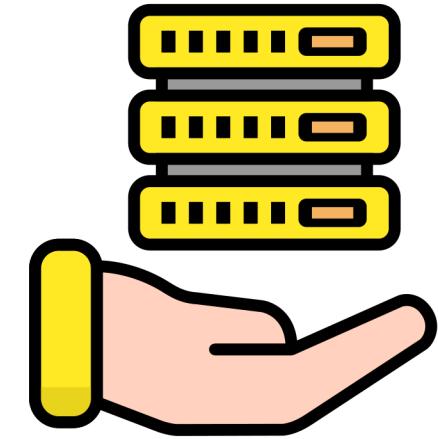
HISTORY



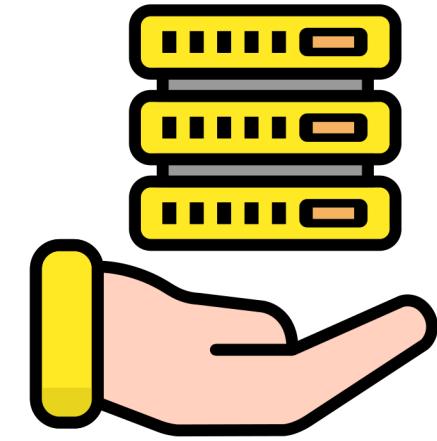
YANDEX METRICA



630 Nodes



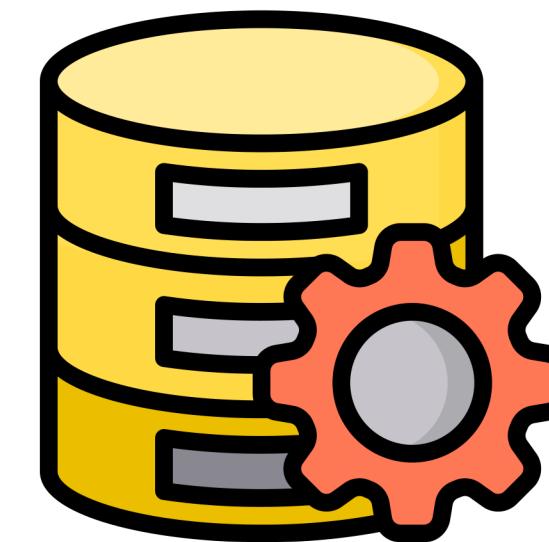
360 Nodes



1860+ Nodes



118 PB



120 Trillion Rows

Row vs Column-Oriented Data Stores

Row	Name	ID	City	Age
1	Name 1	1	City 1	32
2	Name 2	2	City 2	25
3	Name 3	3	City 3	28
x

1. Tabular data

File 1	Name	Name 1	Name 2	Name 3	...
File 2	ID	1	2	3	...
File 3	City	City 1	City 2	City 3	...
File 4	Age	32	25	28	...

2. Tabular data stored in column oriented data store

KEY FEATURES

01

SPEED

02

SQL SUPPORT

03

SCALABLE

04

RELIABLE

05

INTEGRATIONS

06

COMMUNITY



WHEN TO USE CLICKHOUSE DB?

- Write << Read & No or very small amounts of updates/Deletes
 - Data is inserted in large batches
 - Tables have a large number of columns
 - Data is filtered or aggregated while querying
 - No Transactions
-

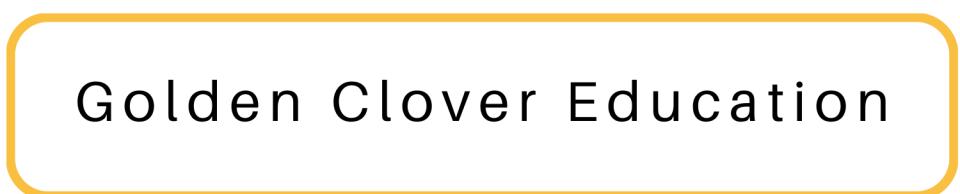
LET'S GET STARTED





INSTALLING

ClickHouse



Golden Clover Education



www.learn-olapdb.com

ClickHouse

Officially Supported on Linux and Mac OS

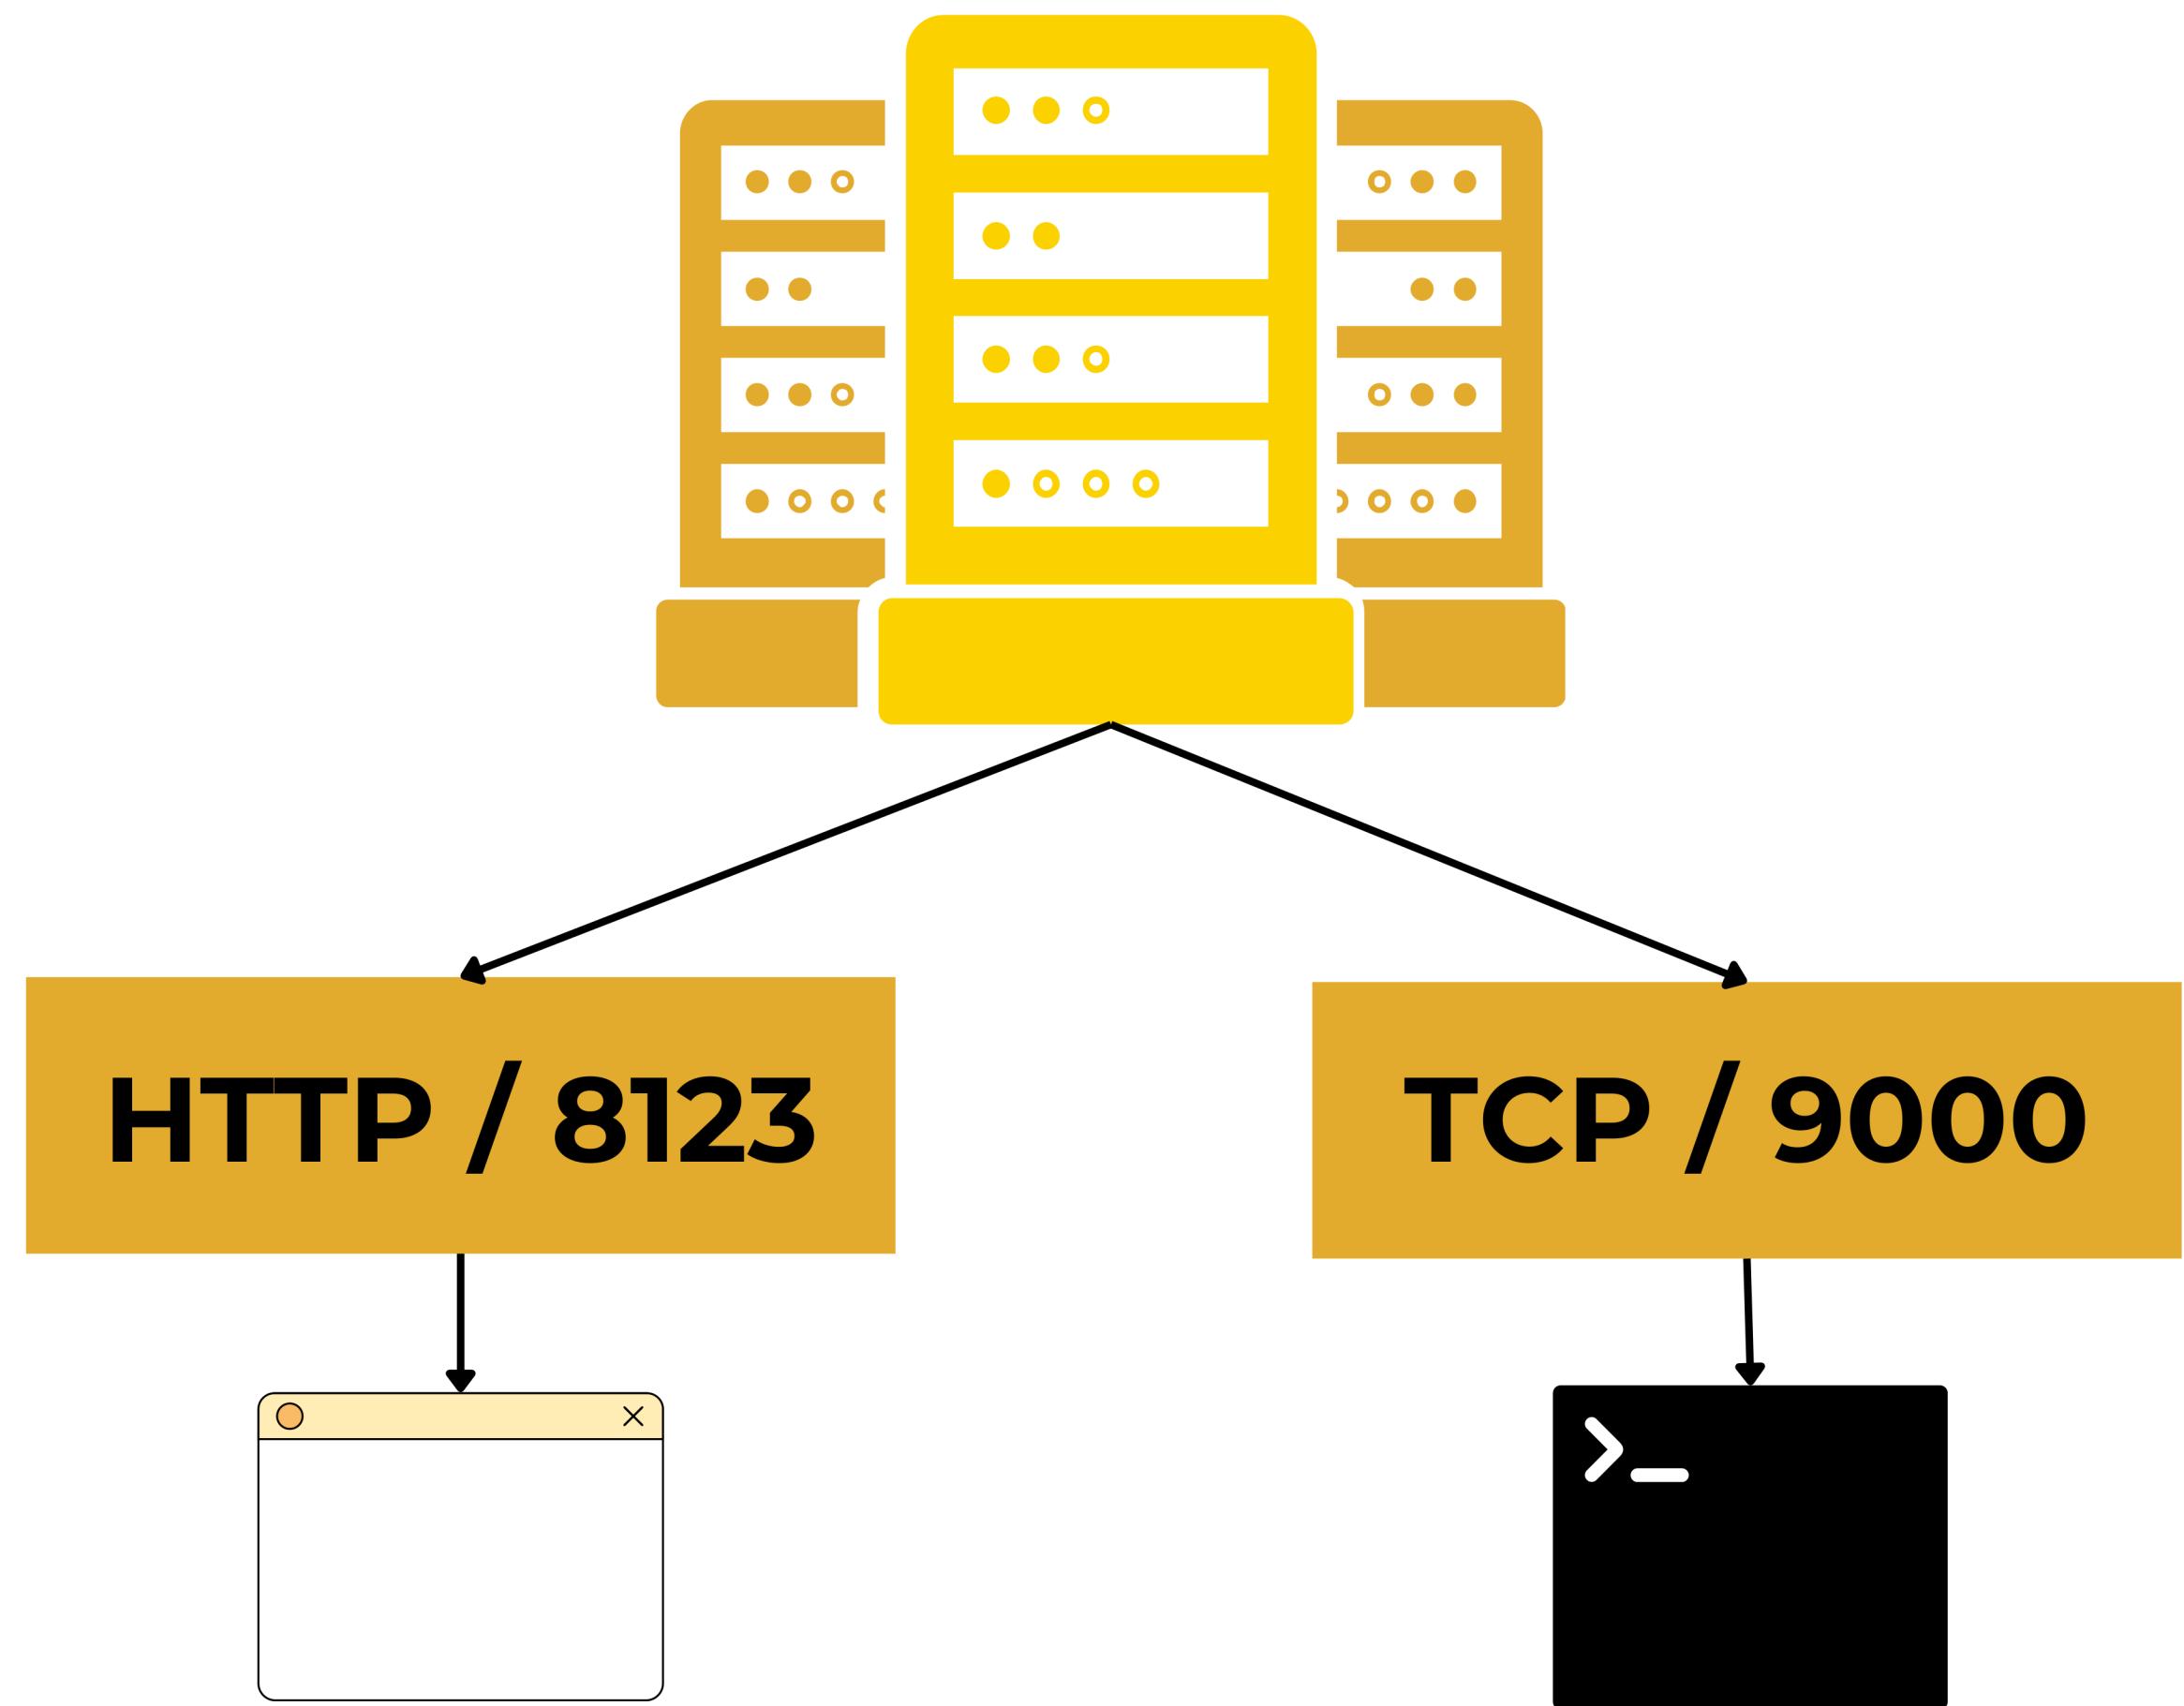
Self managed

- Quick Install
- Install using Precompiled Packages
- Docker
- Kubernetes

Fully Managed

- Altinity Cloud
- ClickHouse Cloud
- ChistaDATA Cloud
- DoubleCloud
- Aiven
- Yandex Cloud

ClickHouse



COMMON PARAMETERS

- host**
- port**
- user**
- password**
- query**
- multiquery**
- multiline**
- database**
- format**
- secure**

clickHouse

Golden Clover Education

www.learn-olapdb.com

ClickHouse SQL

JOIN
SELECT
UPDATE
ALTER
DETACH
WHERE
VIEW
INSERT
DELETE
SET
TRUNCATE
OPTIMIZE
RENAME
DROP
GROUP BY
SHOW
REVOKE
GRANT
ORDER BY

MergeTree Table Engines

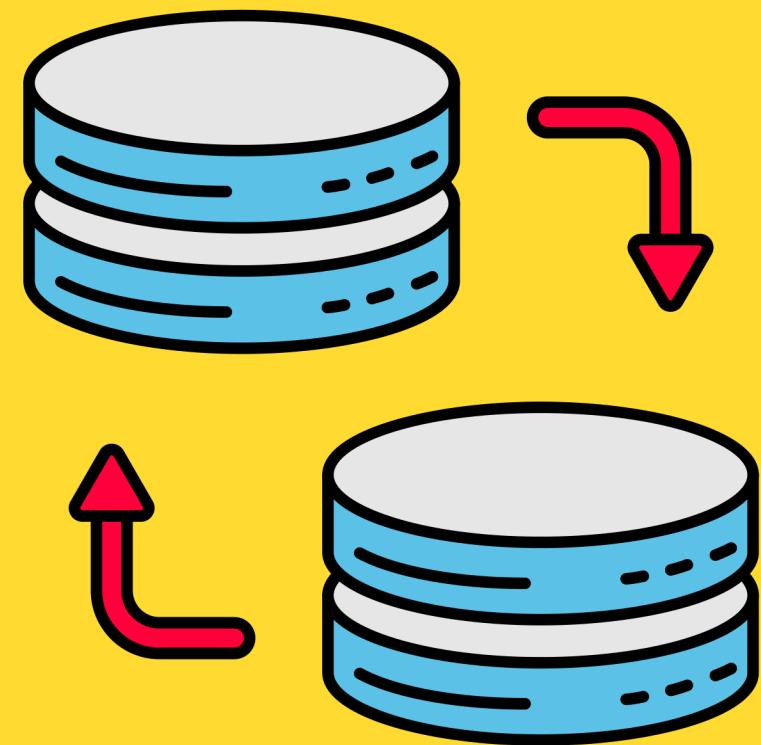
ReplacingMerge
Tree

CollapsingMergeTree

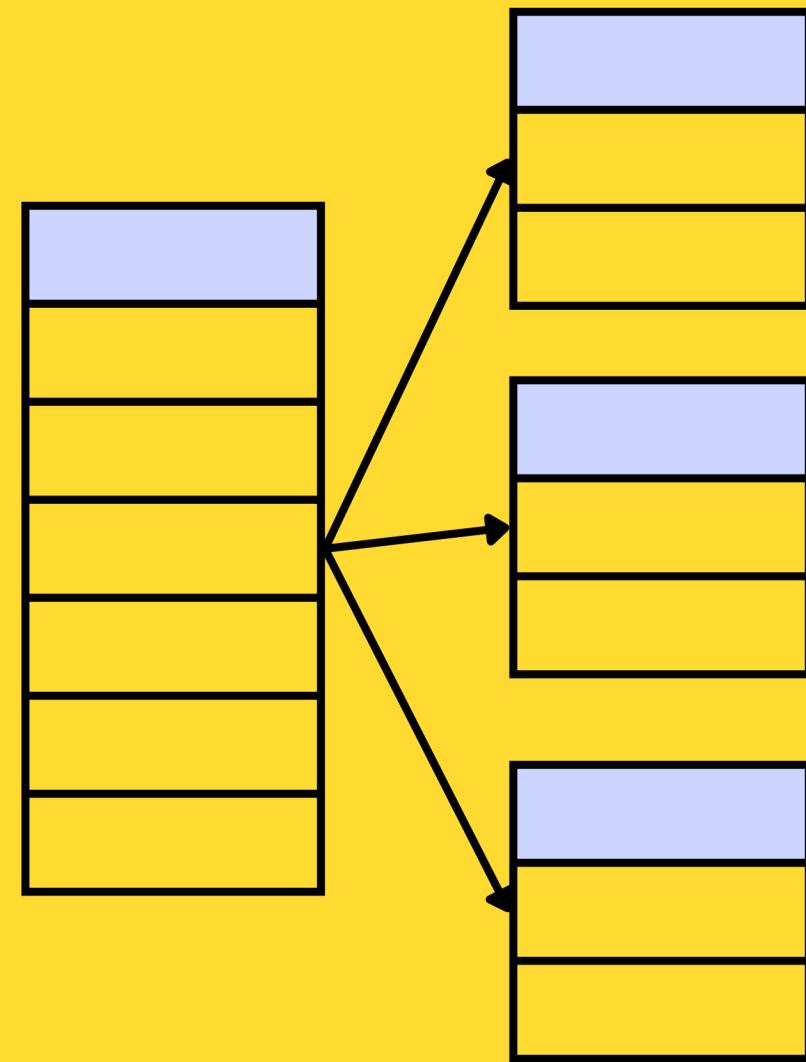
SummingMergeTree

VersionedCollapsing
MergeTree

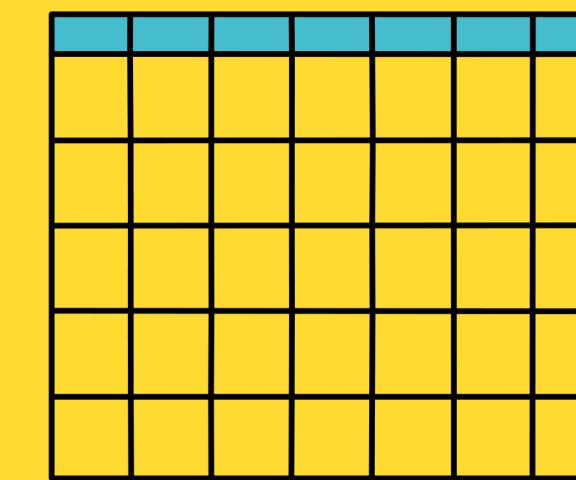
MergeTree



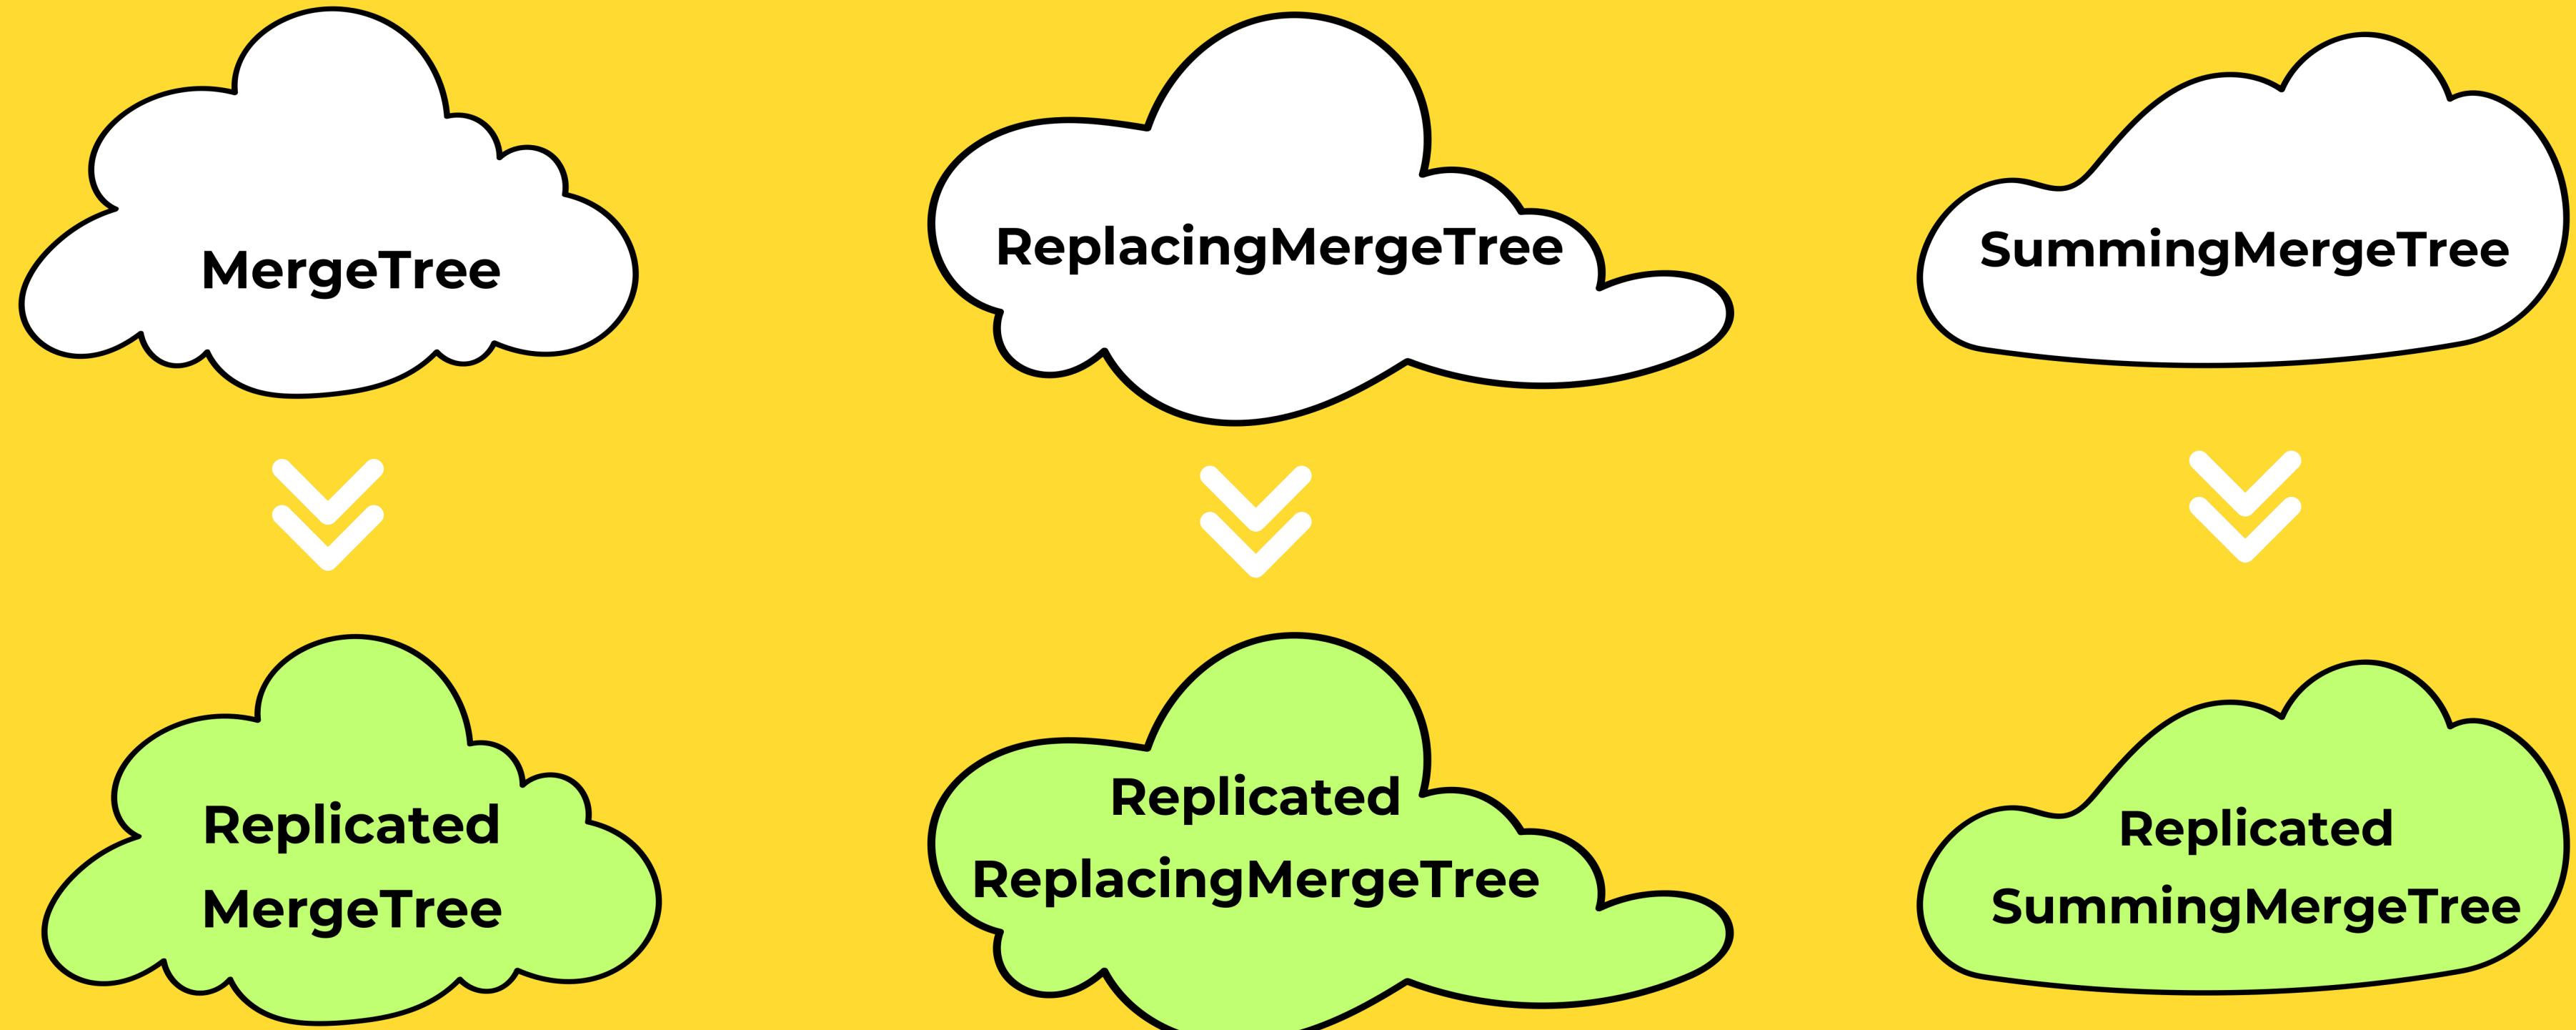
Data Replication



Sharding



Indexing



CollapsingMergeTree

**Versioned
CollapsingMergeTree**

AggregatingMergeTree

**Replicated
CollapsingMergeTree**

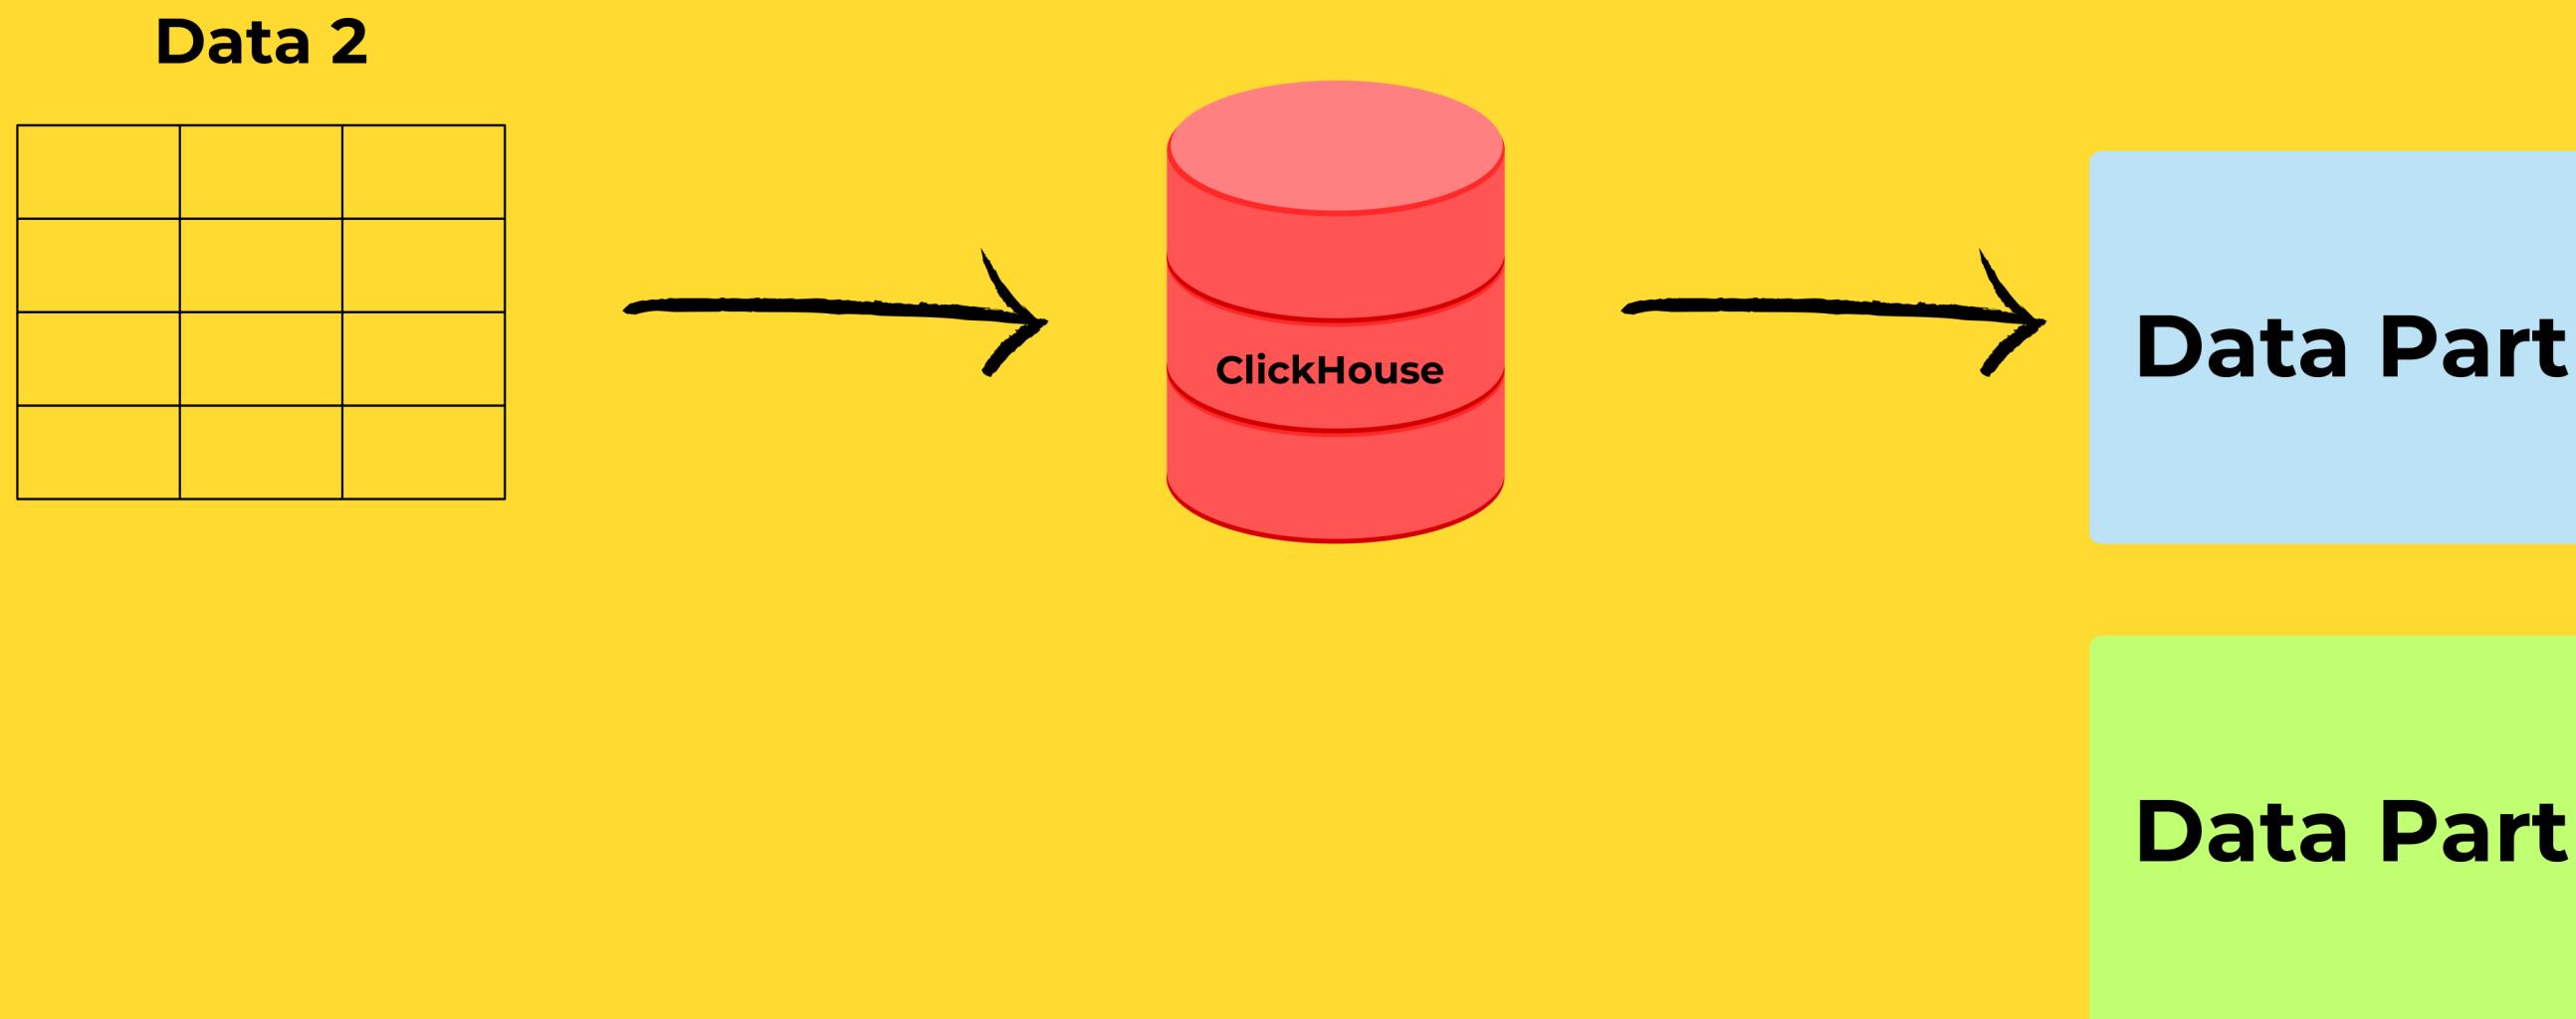
**Replicated
Versioned
Collapsing
MergeTree**

**Replicated
AggregatingMergeTree**

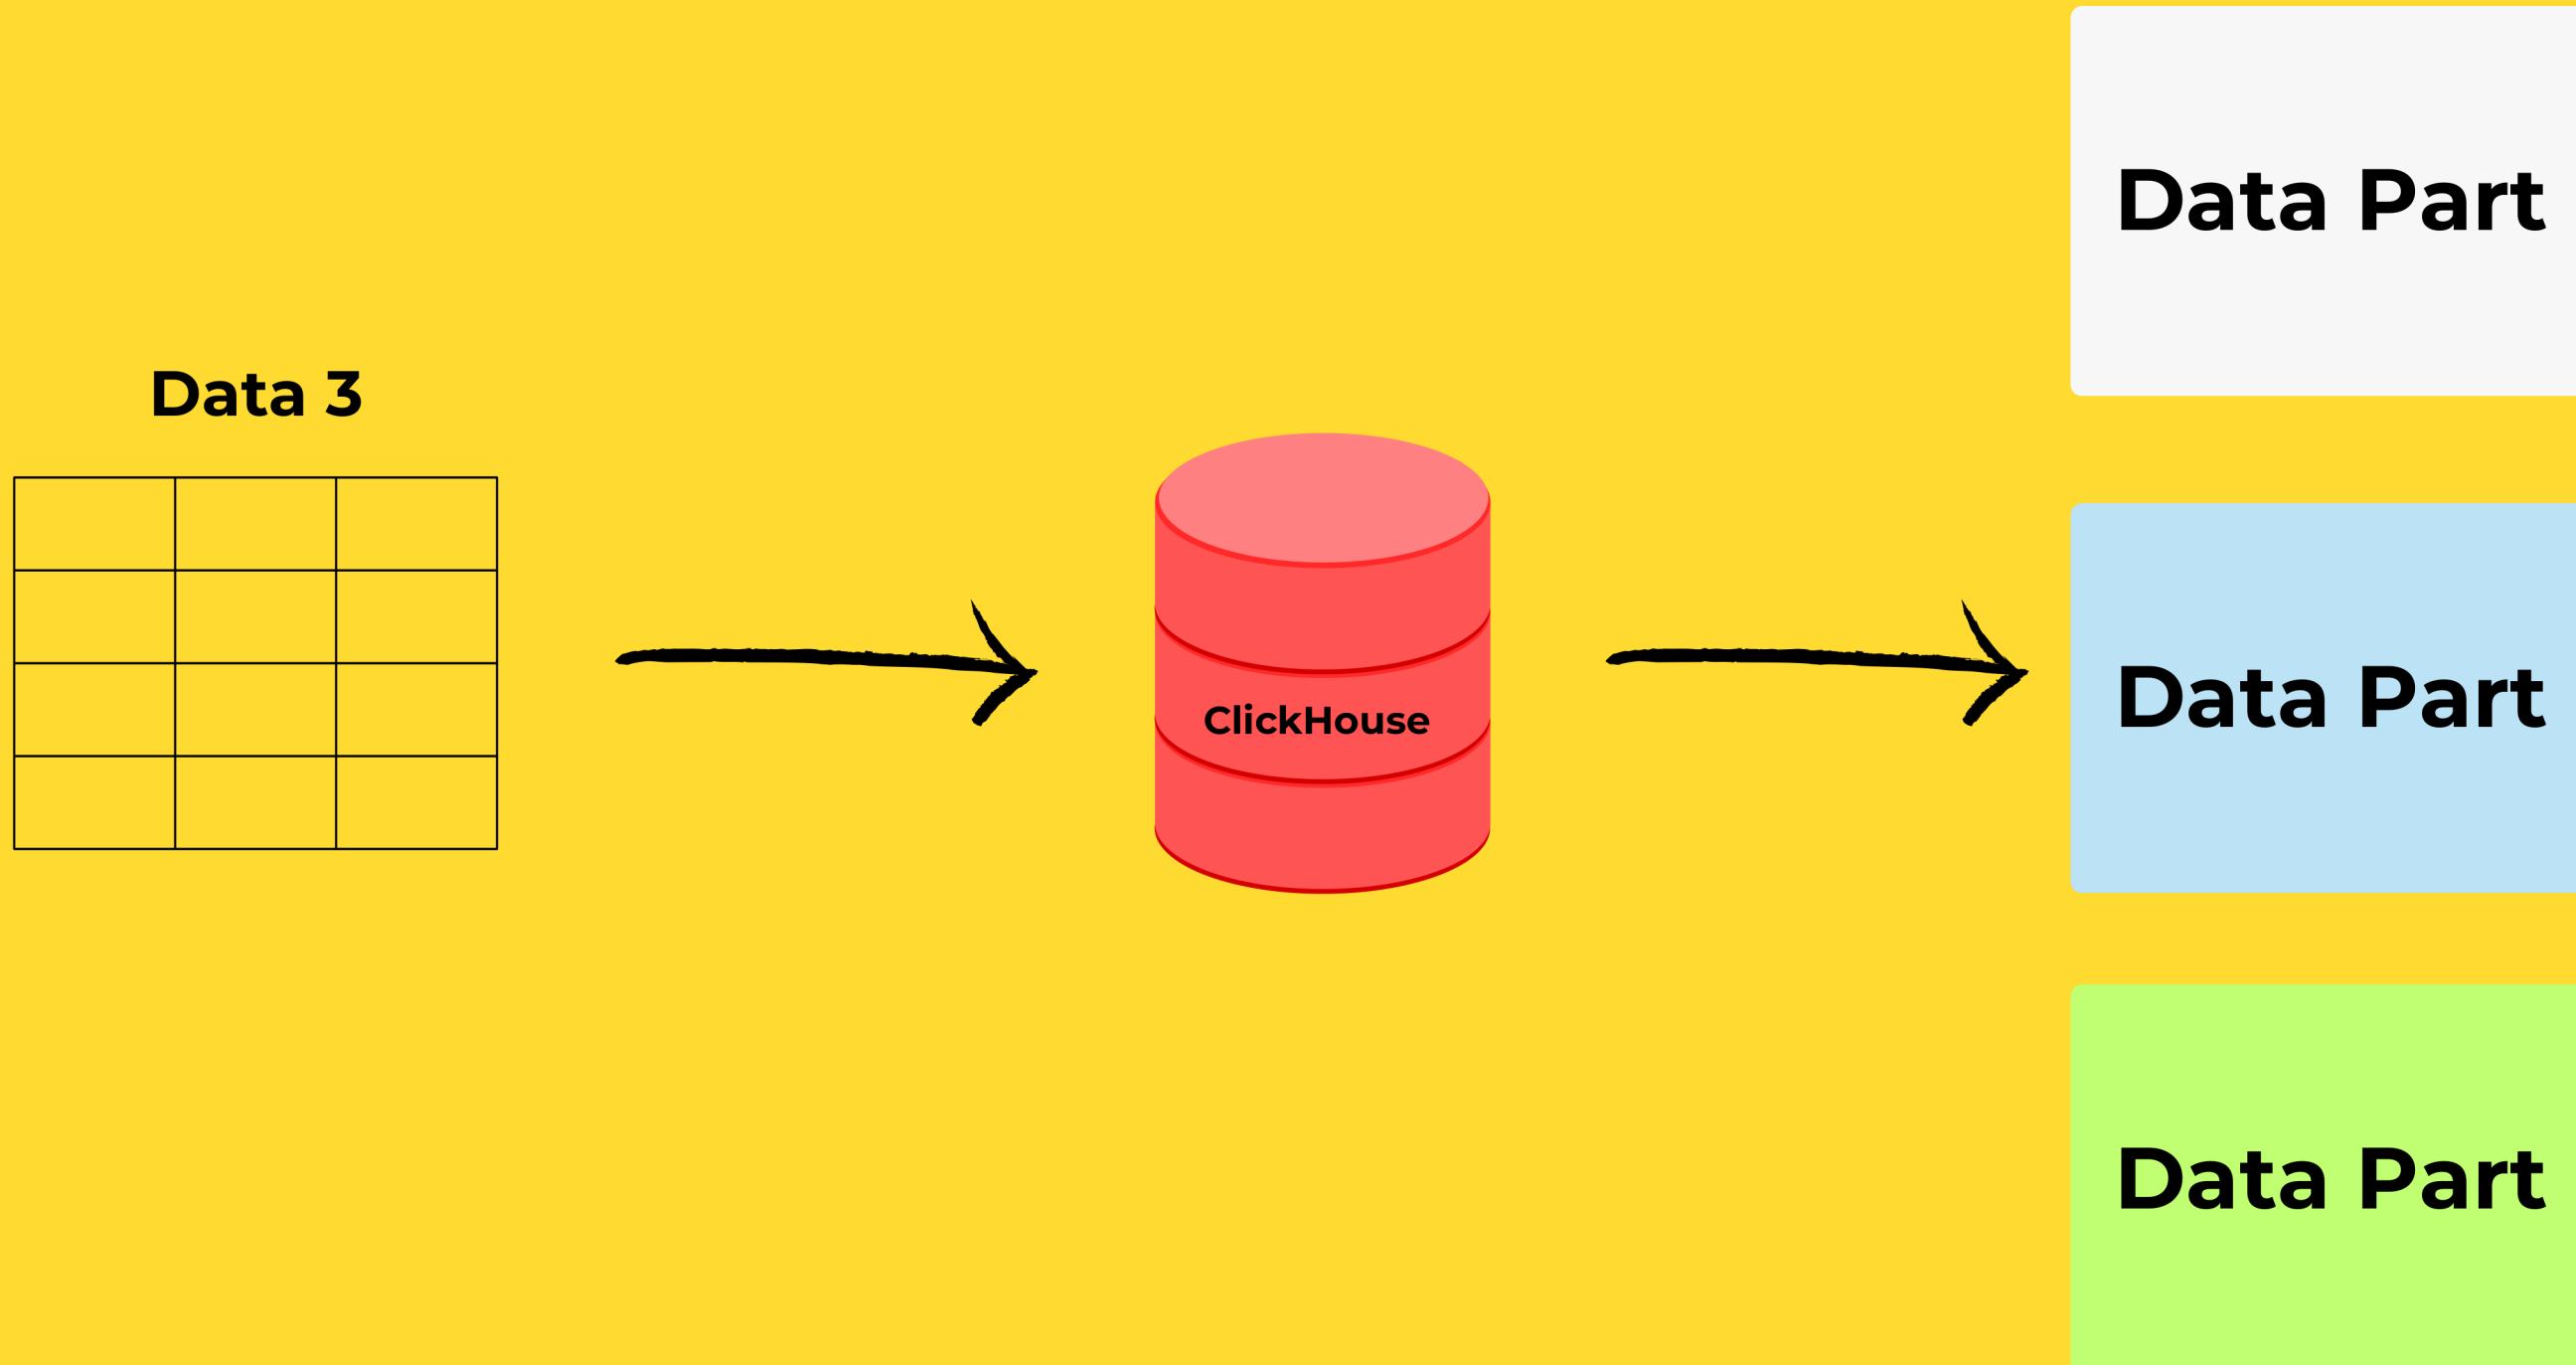
Data Part



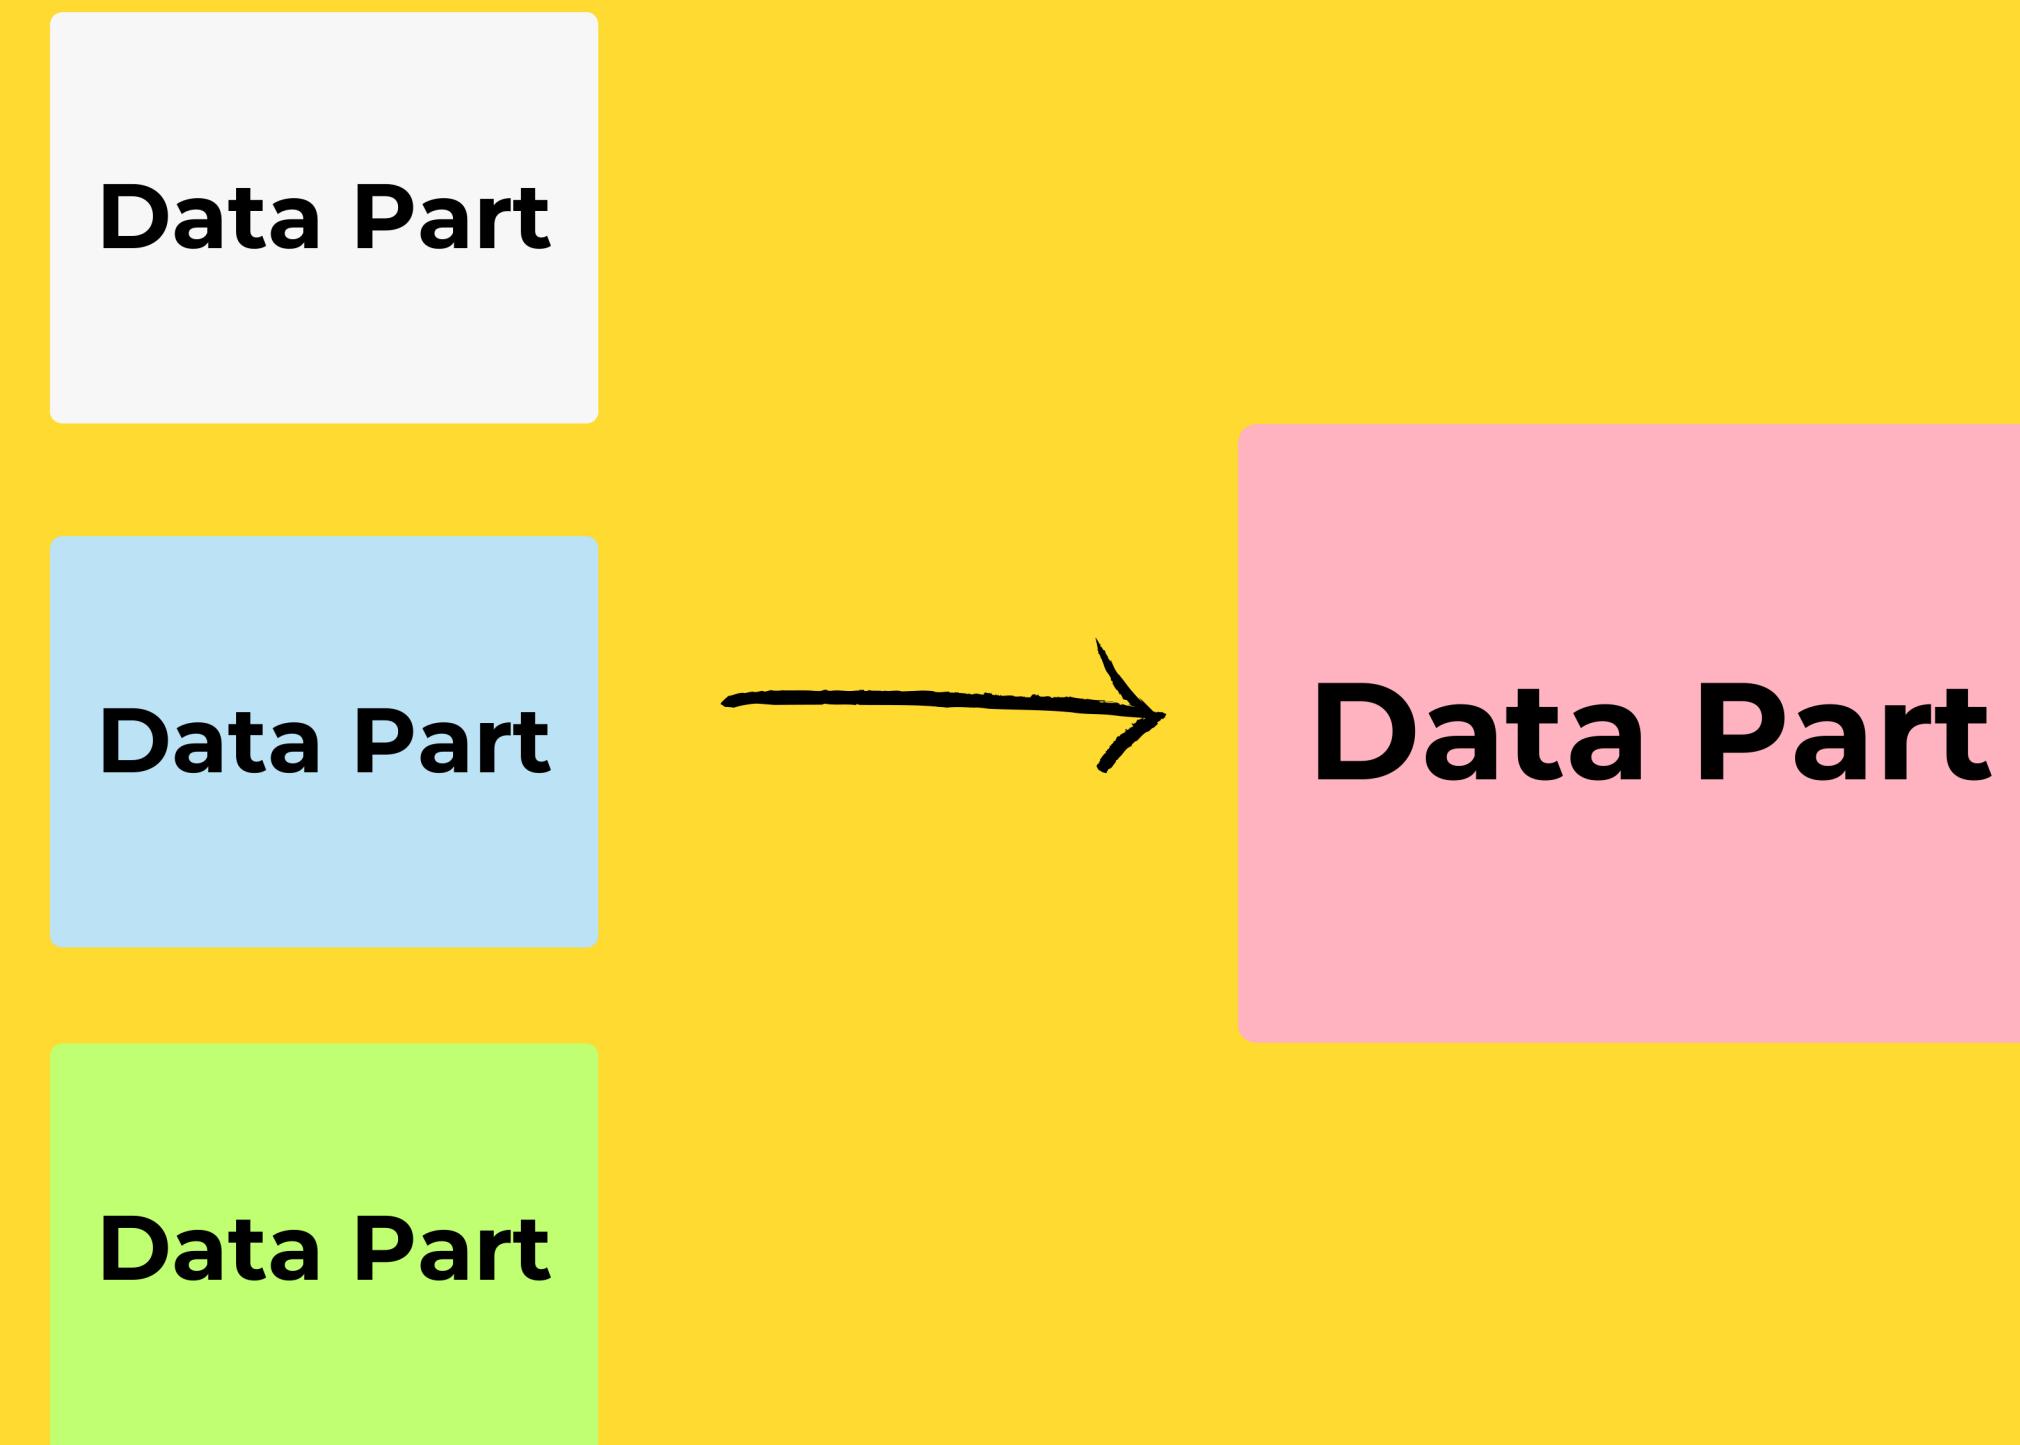
Data parts



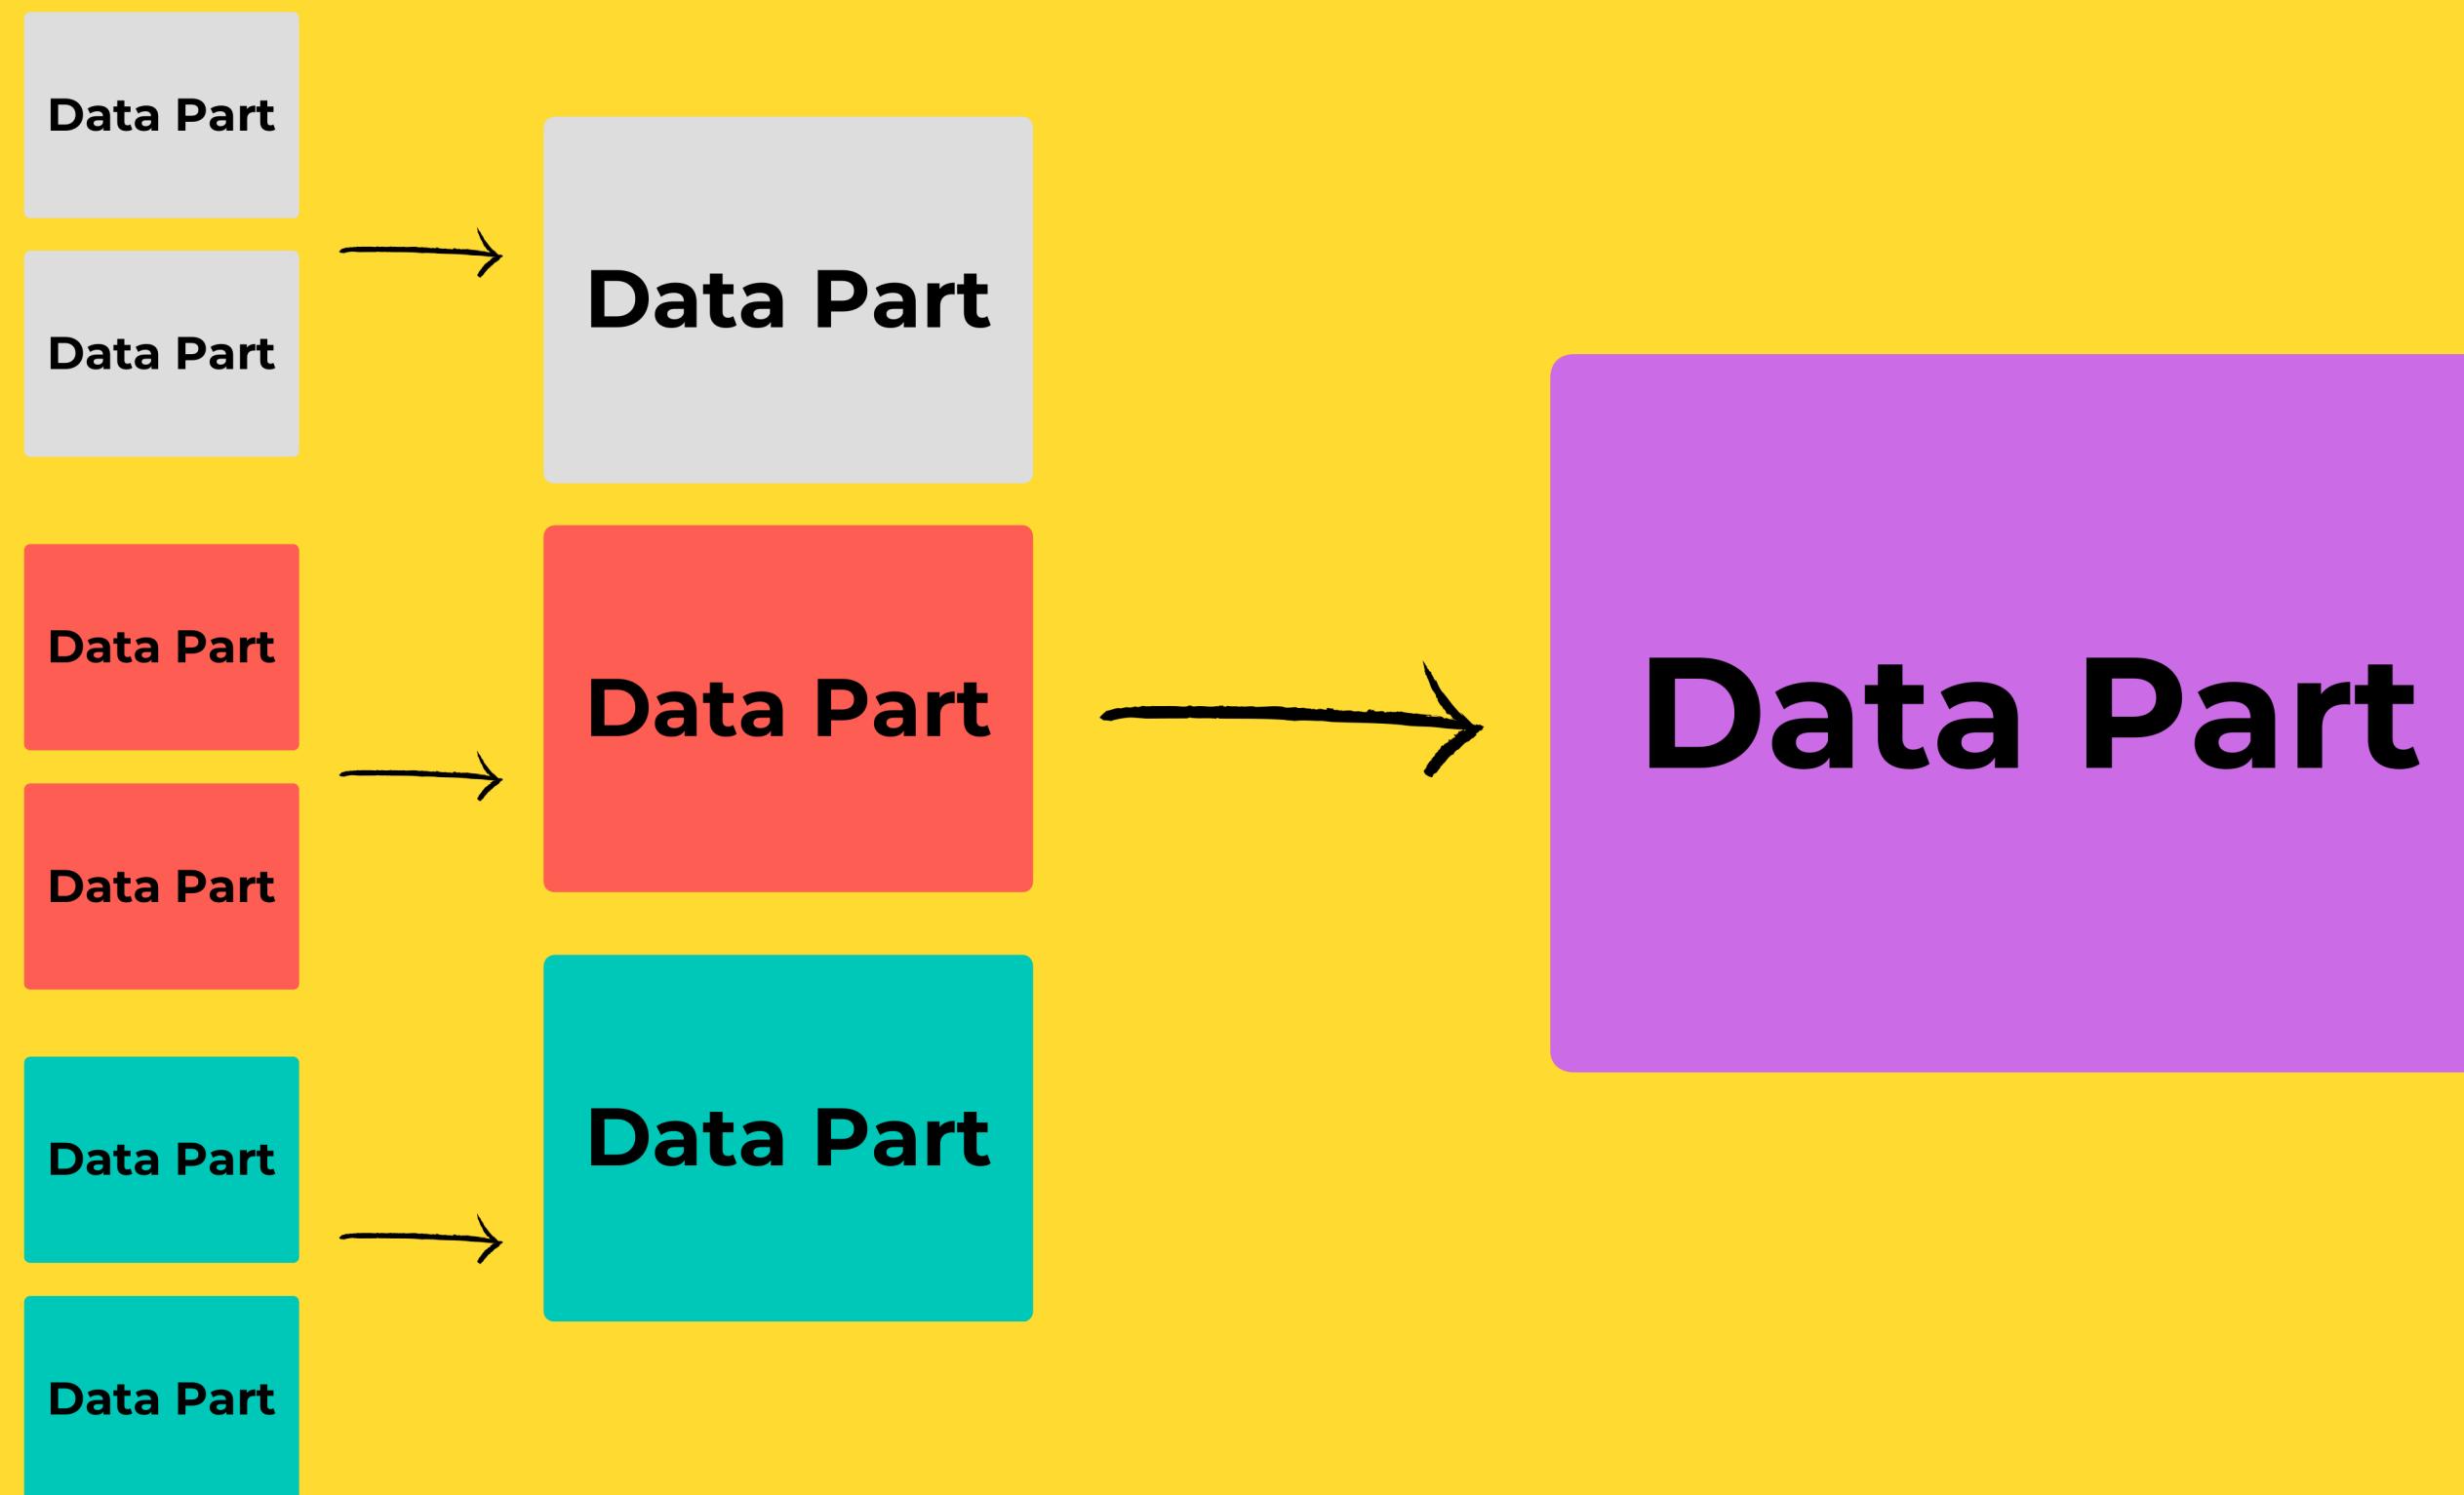
Data parts



Merging the data parts

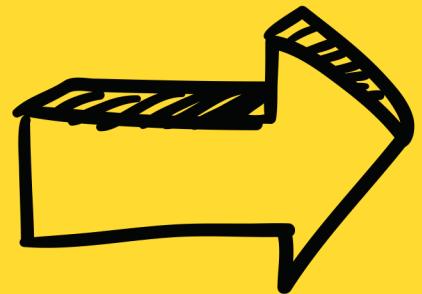


Merging the data parts



MergeTree - Example

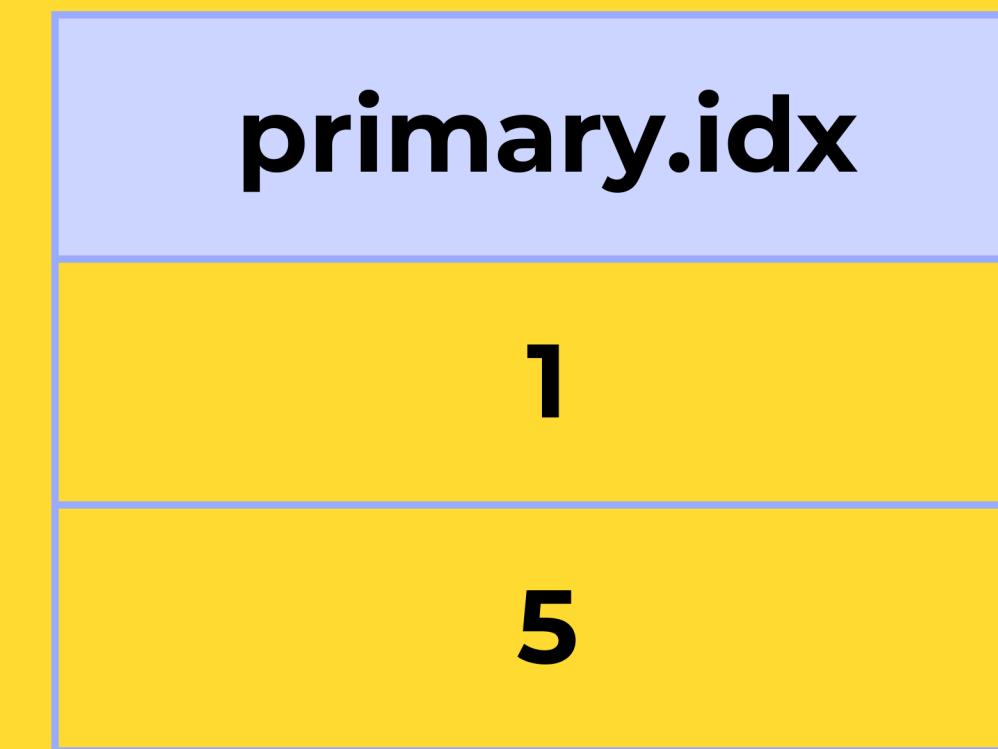
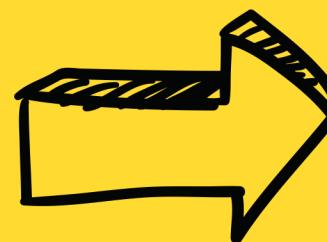
ID	Name	Age



```
CREATE TABLE mergetree_example
(
    ID Int32,
    Name String,
    Age UInt8
)
Engine = MergeTree
PRIMARY KEY (ID)
SETTINGS
index_granularity = 4
```

PRIMARY INDEX & Granules

ID	Name	Age
1	k	15
2	t	8
3	a	10
4	l	17
5	c	12
6	g	7
7	q	4
10	e	2



Inserts & Parts

Insert Data

ID	Name	Age
3	a	10
5	c	12
2	t	8
7	q	4

Sorted Data

ID	Name	Age
2	t	8
3	a	10
5	c	12
7	q	4

Stored in Disk

mergetree_example

Data Part 1

Inserts & Parts

Insert Data

ID	Name	Age
10	e	2
6	g	7
1	k	15
4	l	17

Sorted Data

ID	Name	Age
1	k	15
4	l	17
6	g	7
10	e	2

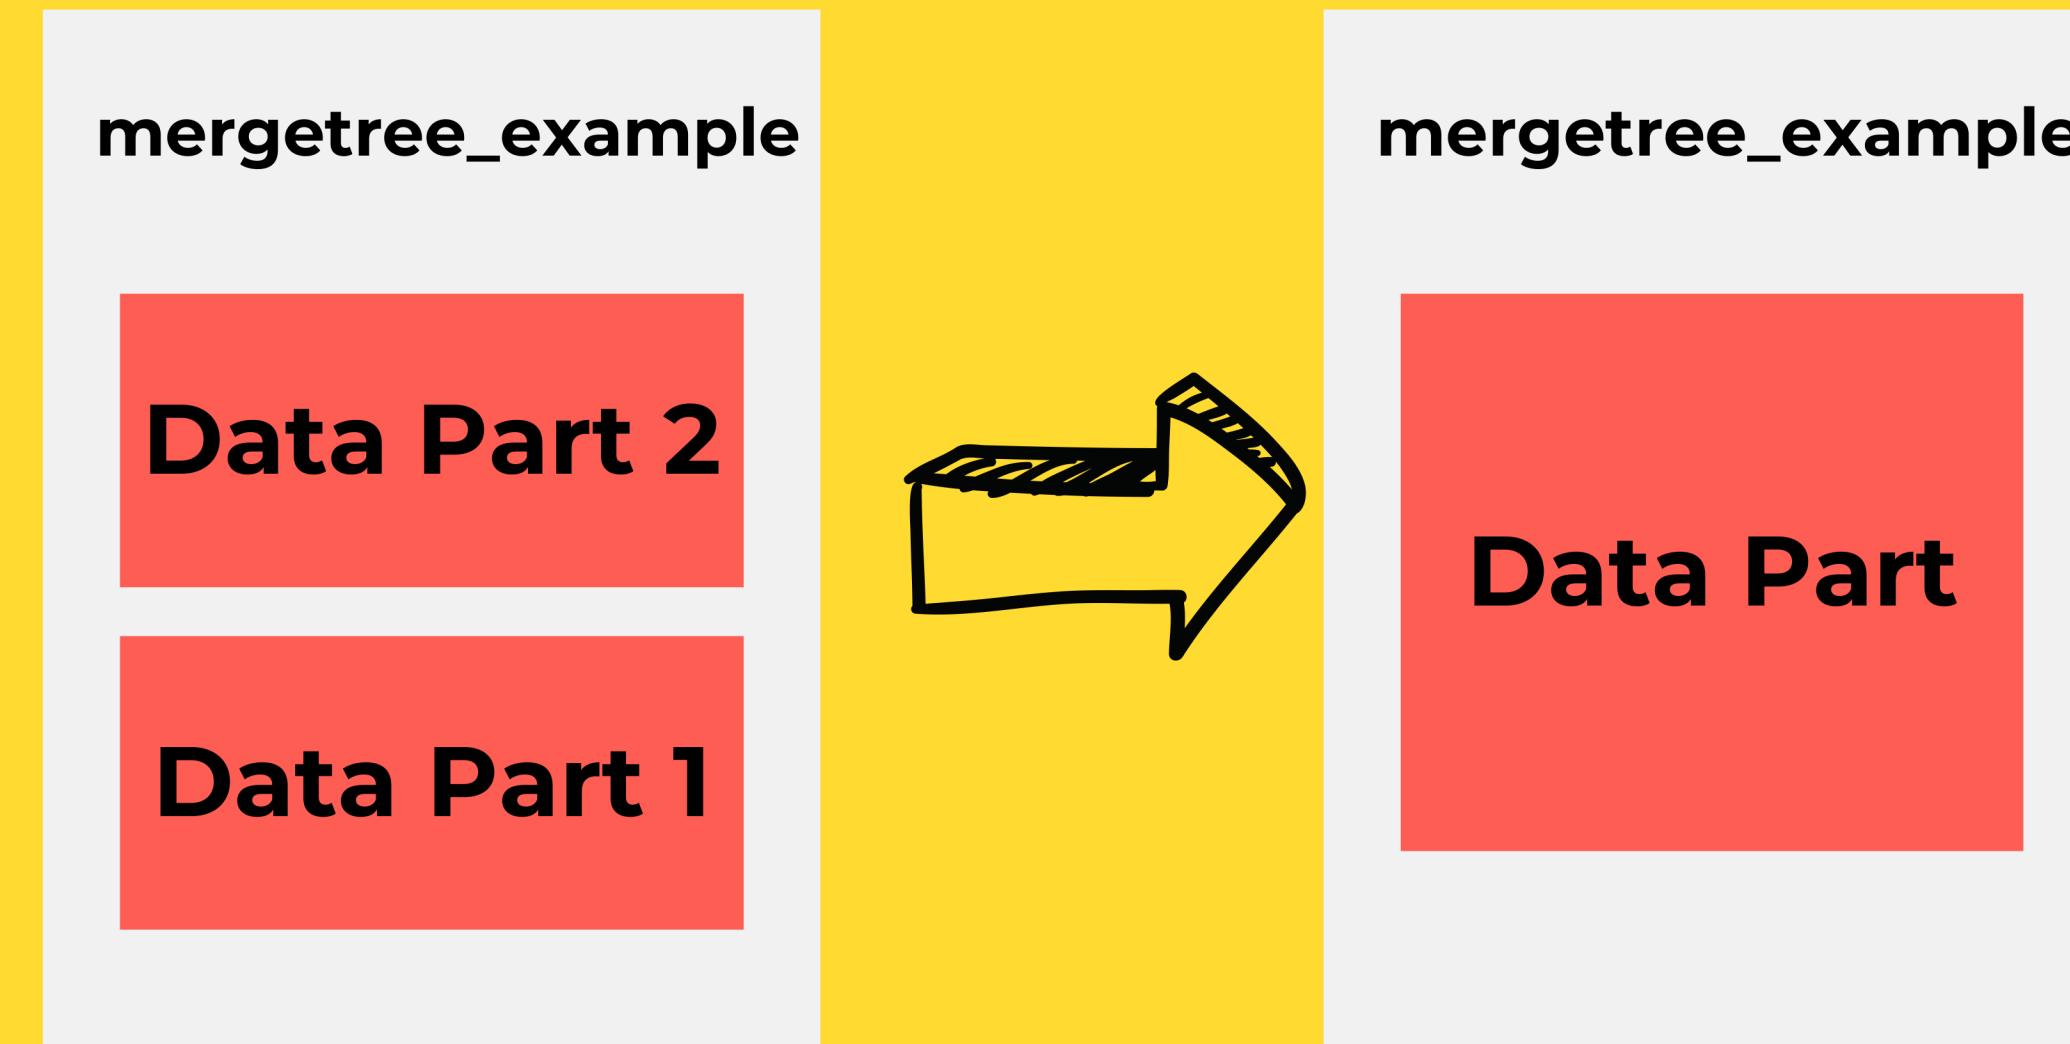
Stored in Disk

mergetree_example

Data Part 2

Data Part 1

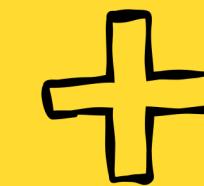
Merging of Data Parts



Merging of Data Parts

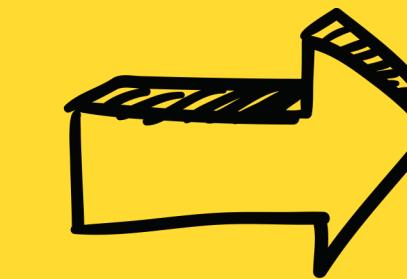
Data part 1

ID	Name	Age
1	k	15
4	l	17
6	g	7
10	e	2



Data part 2

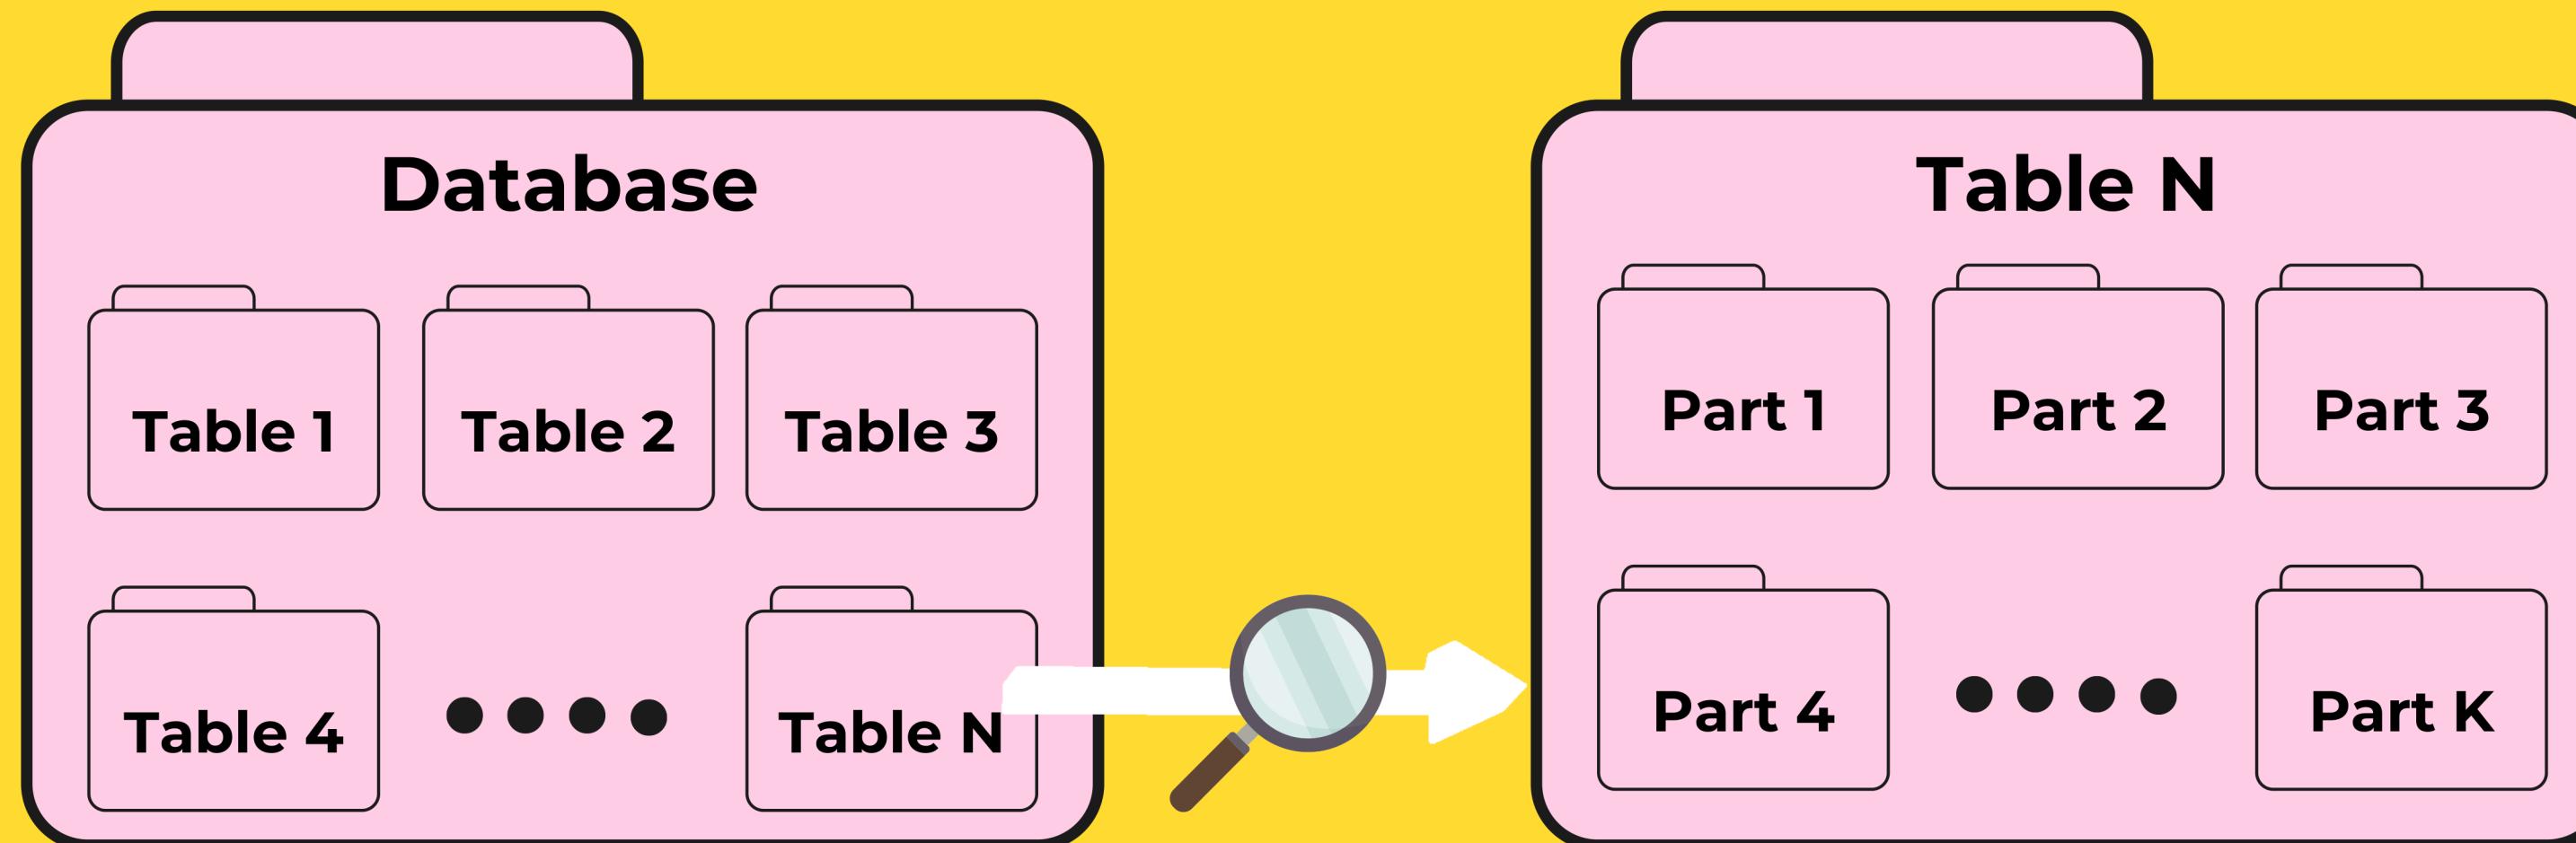
ID	Name	Age
2	t	8
3	a	10
5	c	12
7	q	4



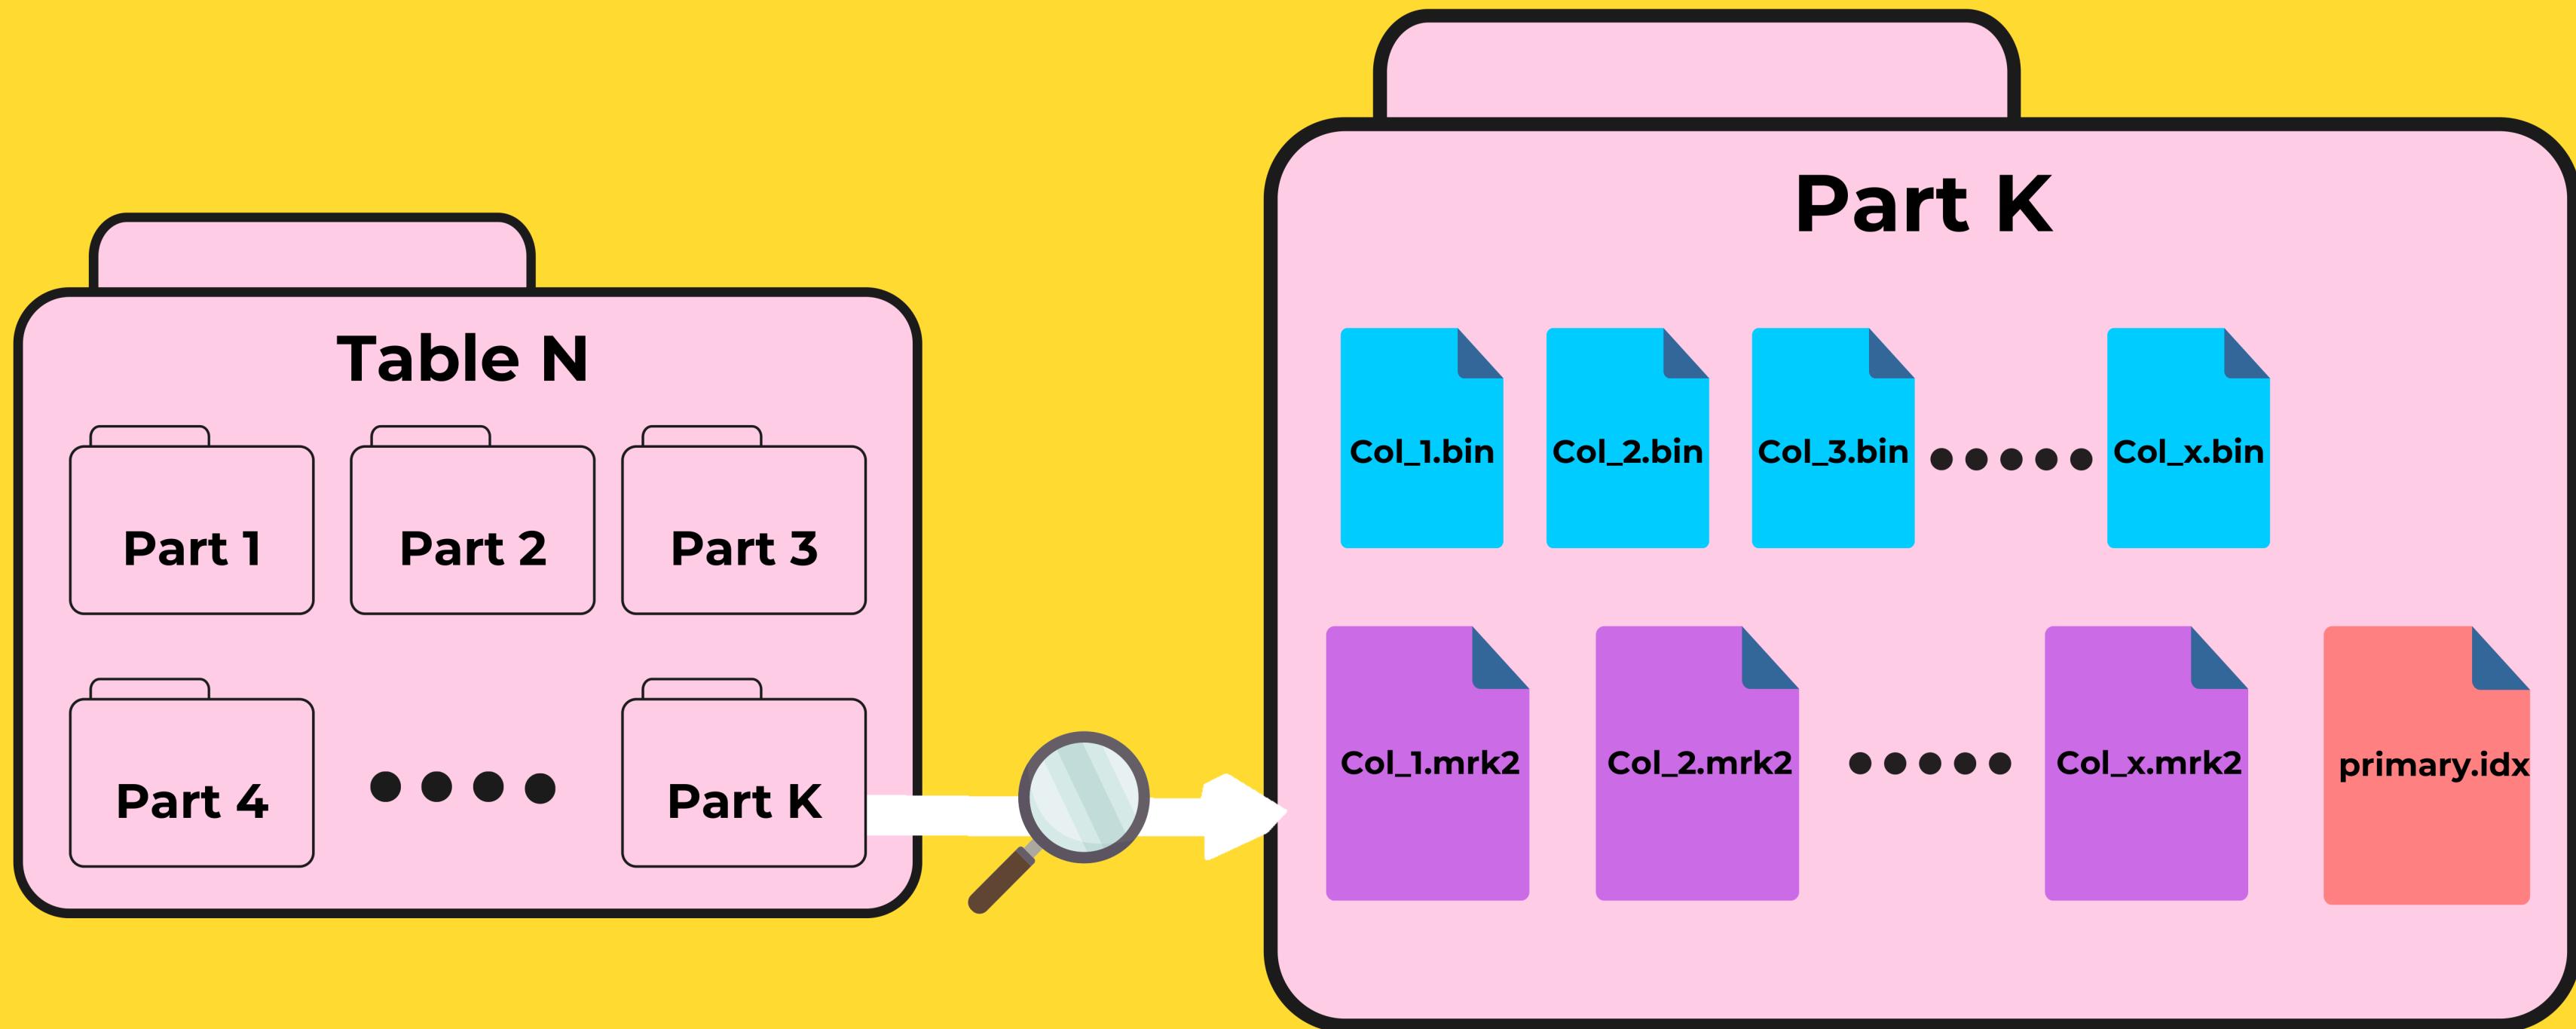
Data Part

ID	Name	Age
1	k	15
2	t	8
3	a	10
4	l	17
5	c	12
6	g	7
7	q	4
10	e	2

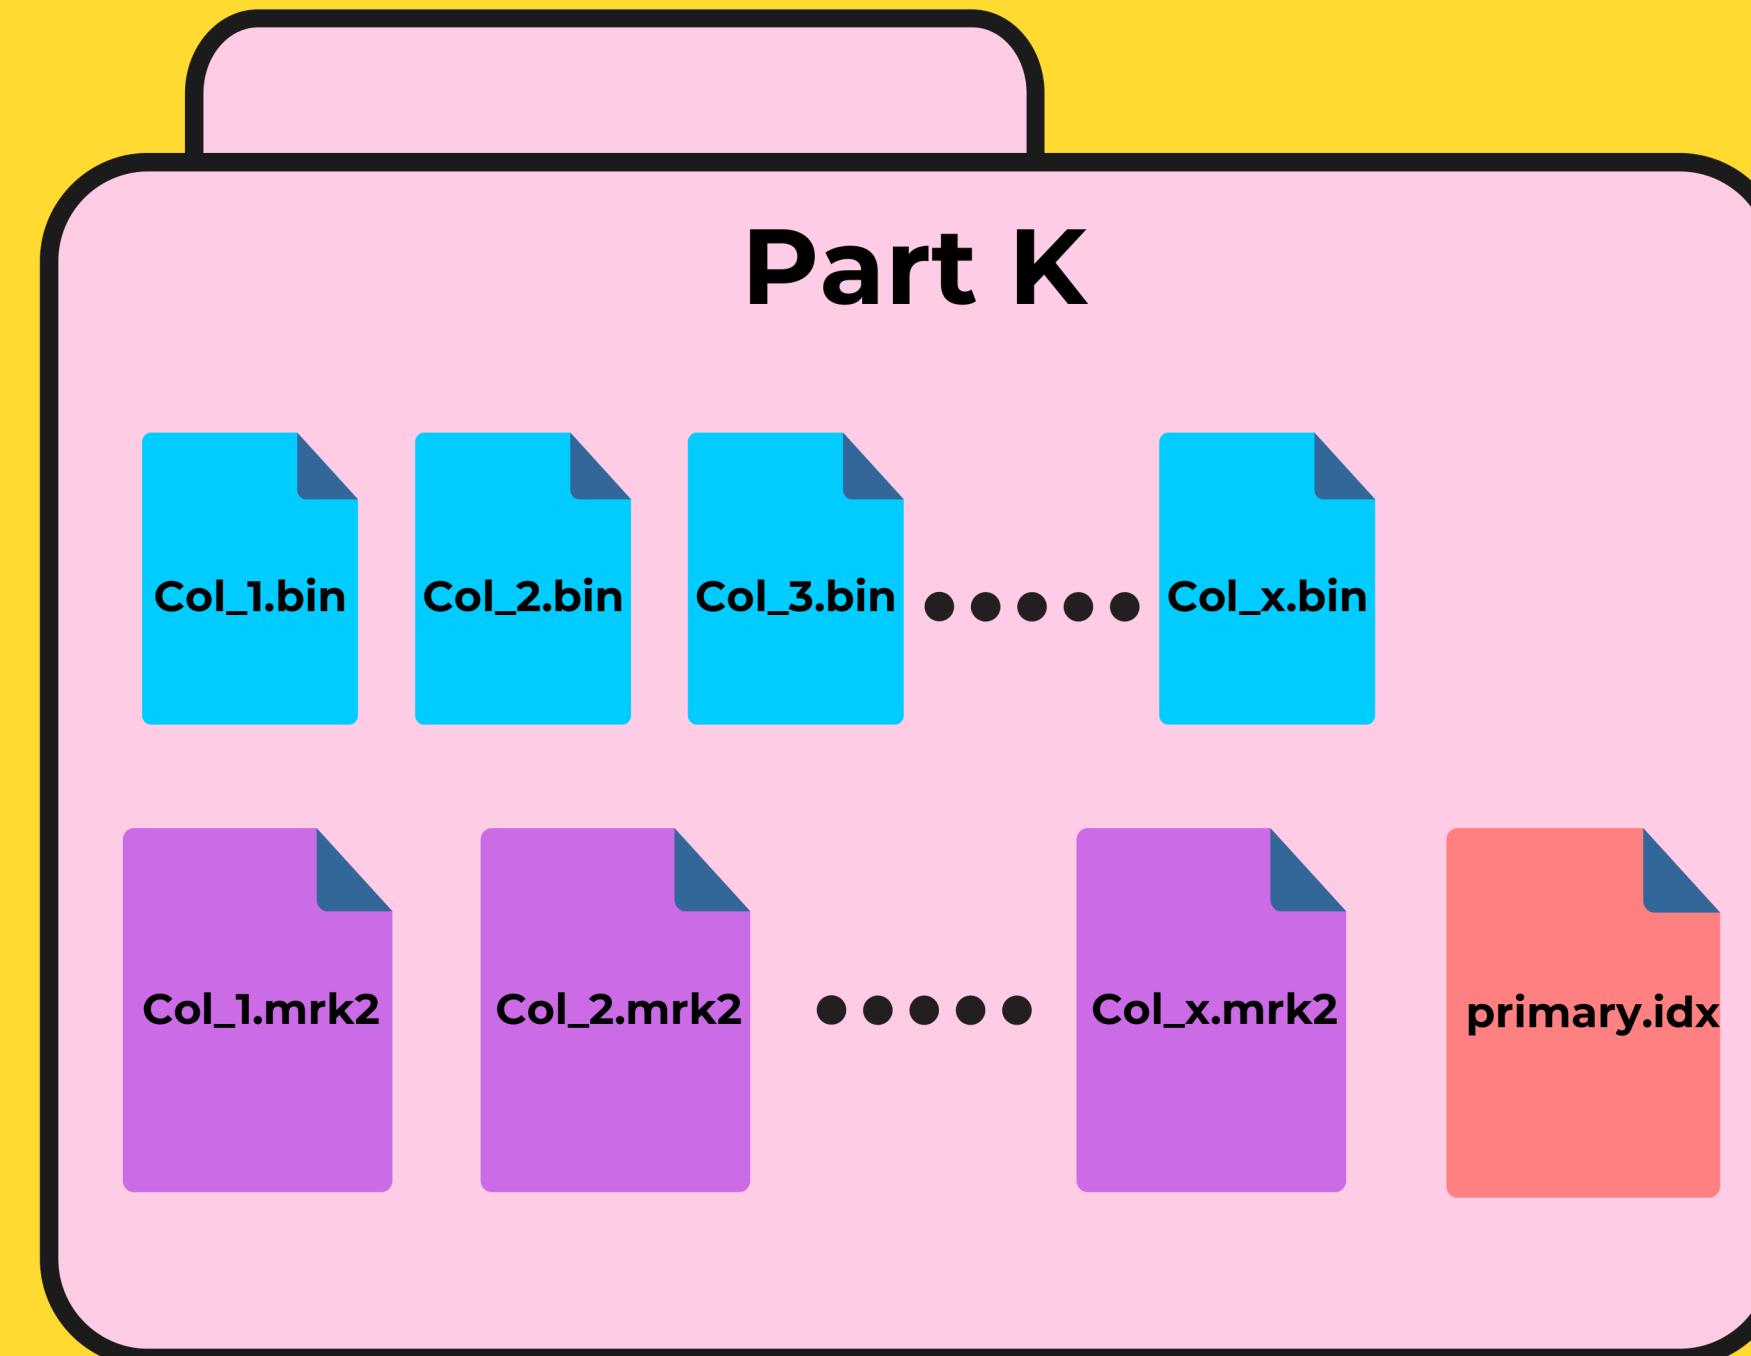
ClickHouse Data Directory



ClickHouse Data Directory



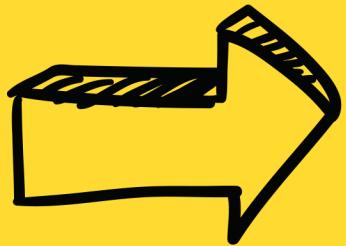
ClickHouse Data Directory



`min_bytes_for_wide_part`

PRIMARY KEY & Granules

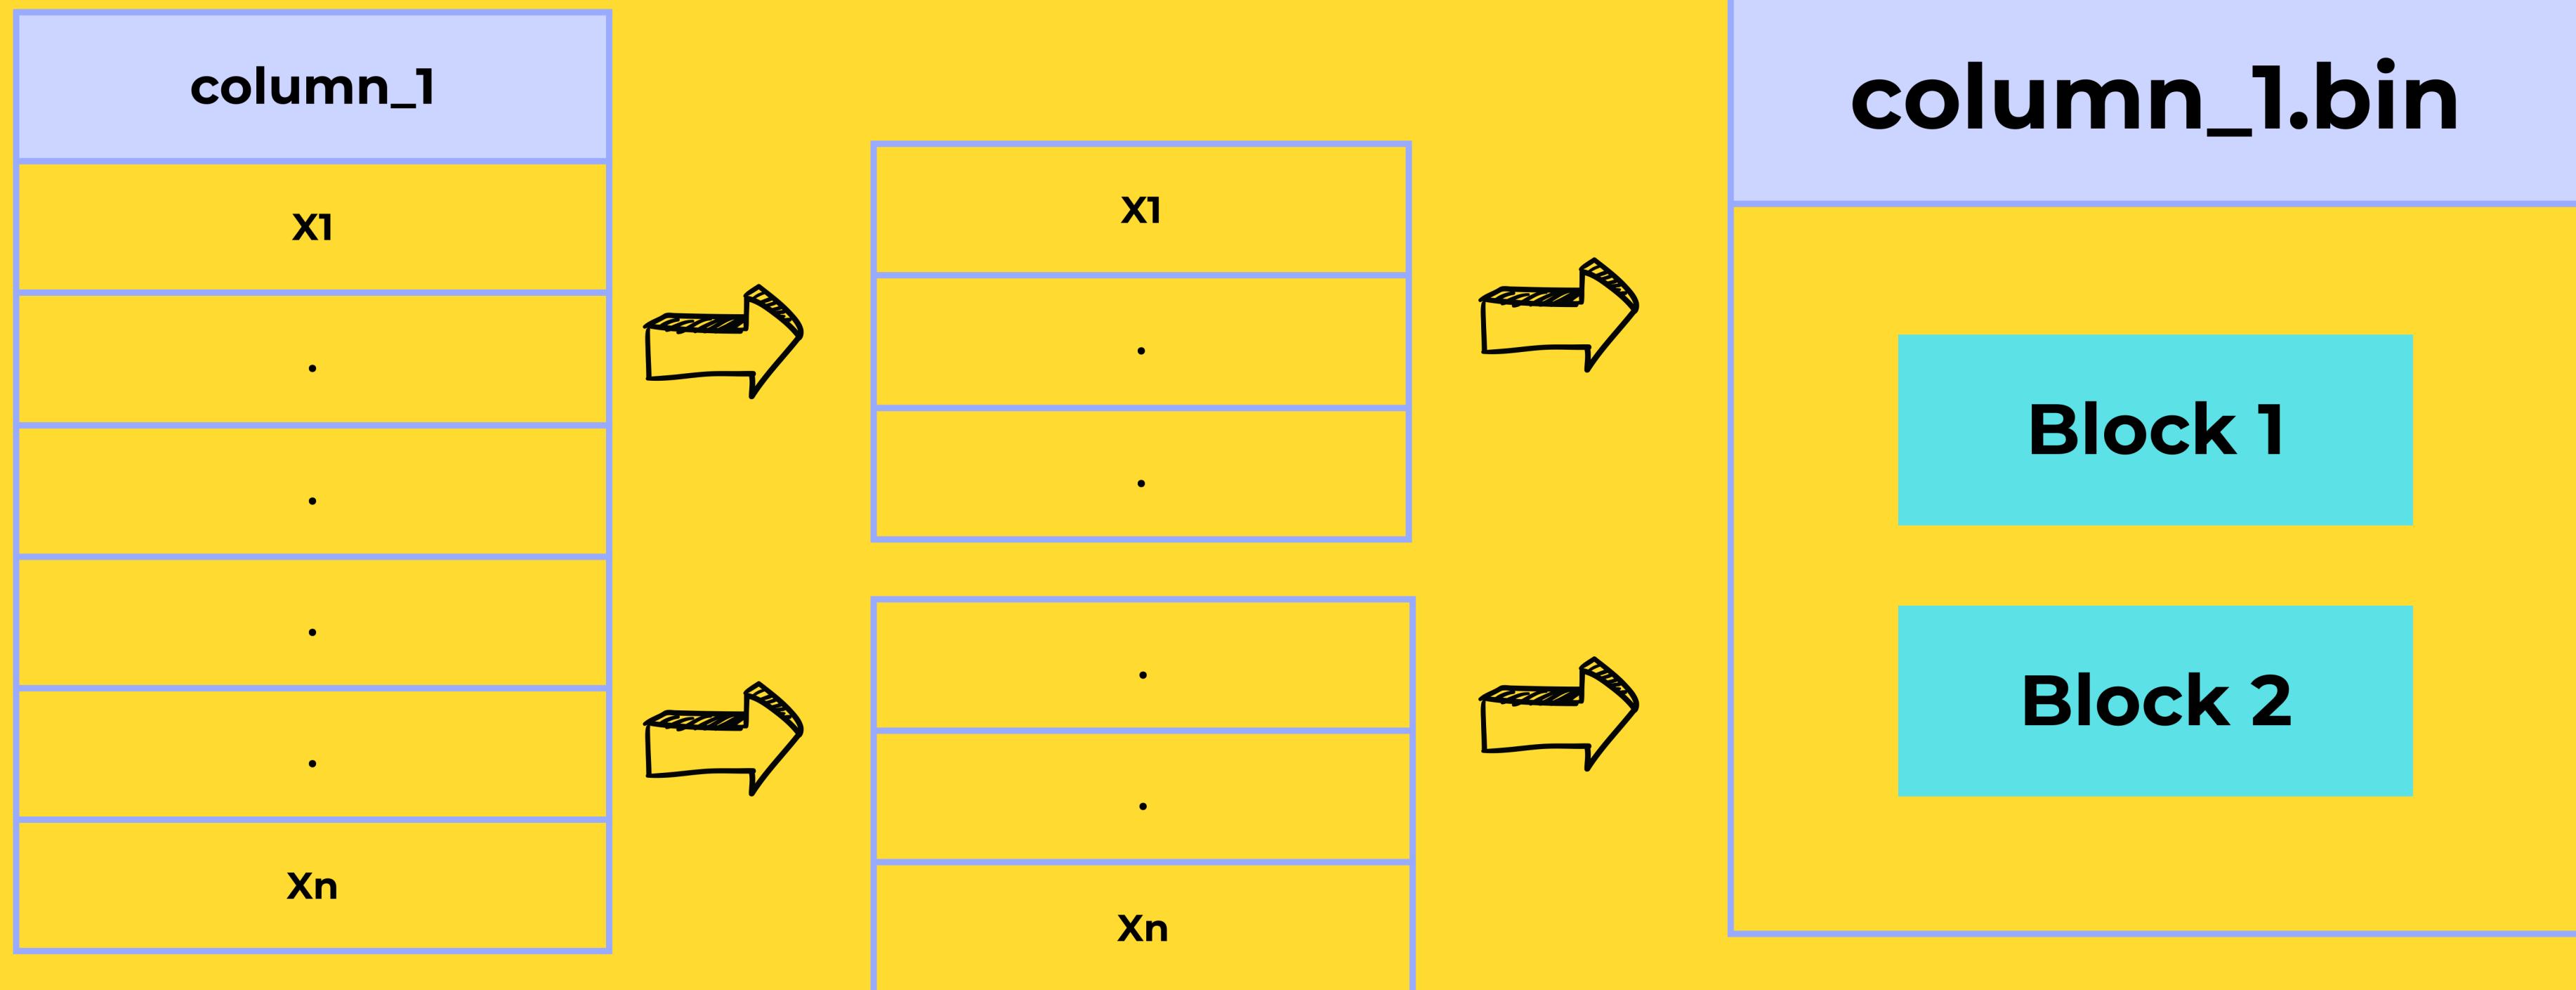
ID	Name	Age
1	k	15
2	t	8
3	a	10
4	l	17
5	c	12
6	g	7
7	q	4
10	e	2



primary.idx
1
5

index_granularity_bytes

<column>.bin & Blocks



ClickHouse Compression Codecs

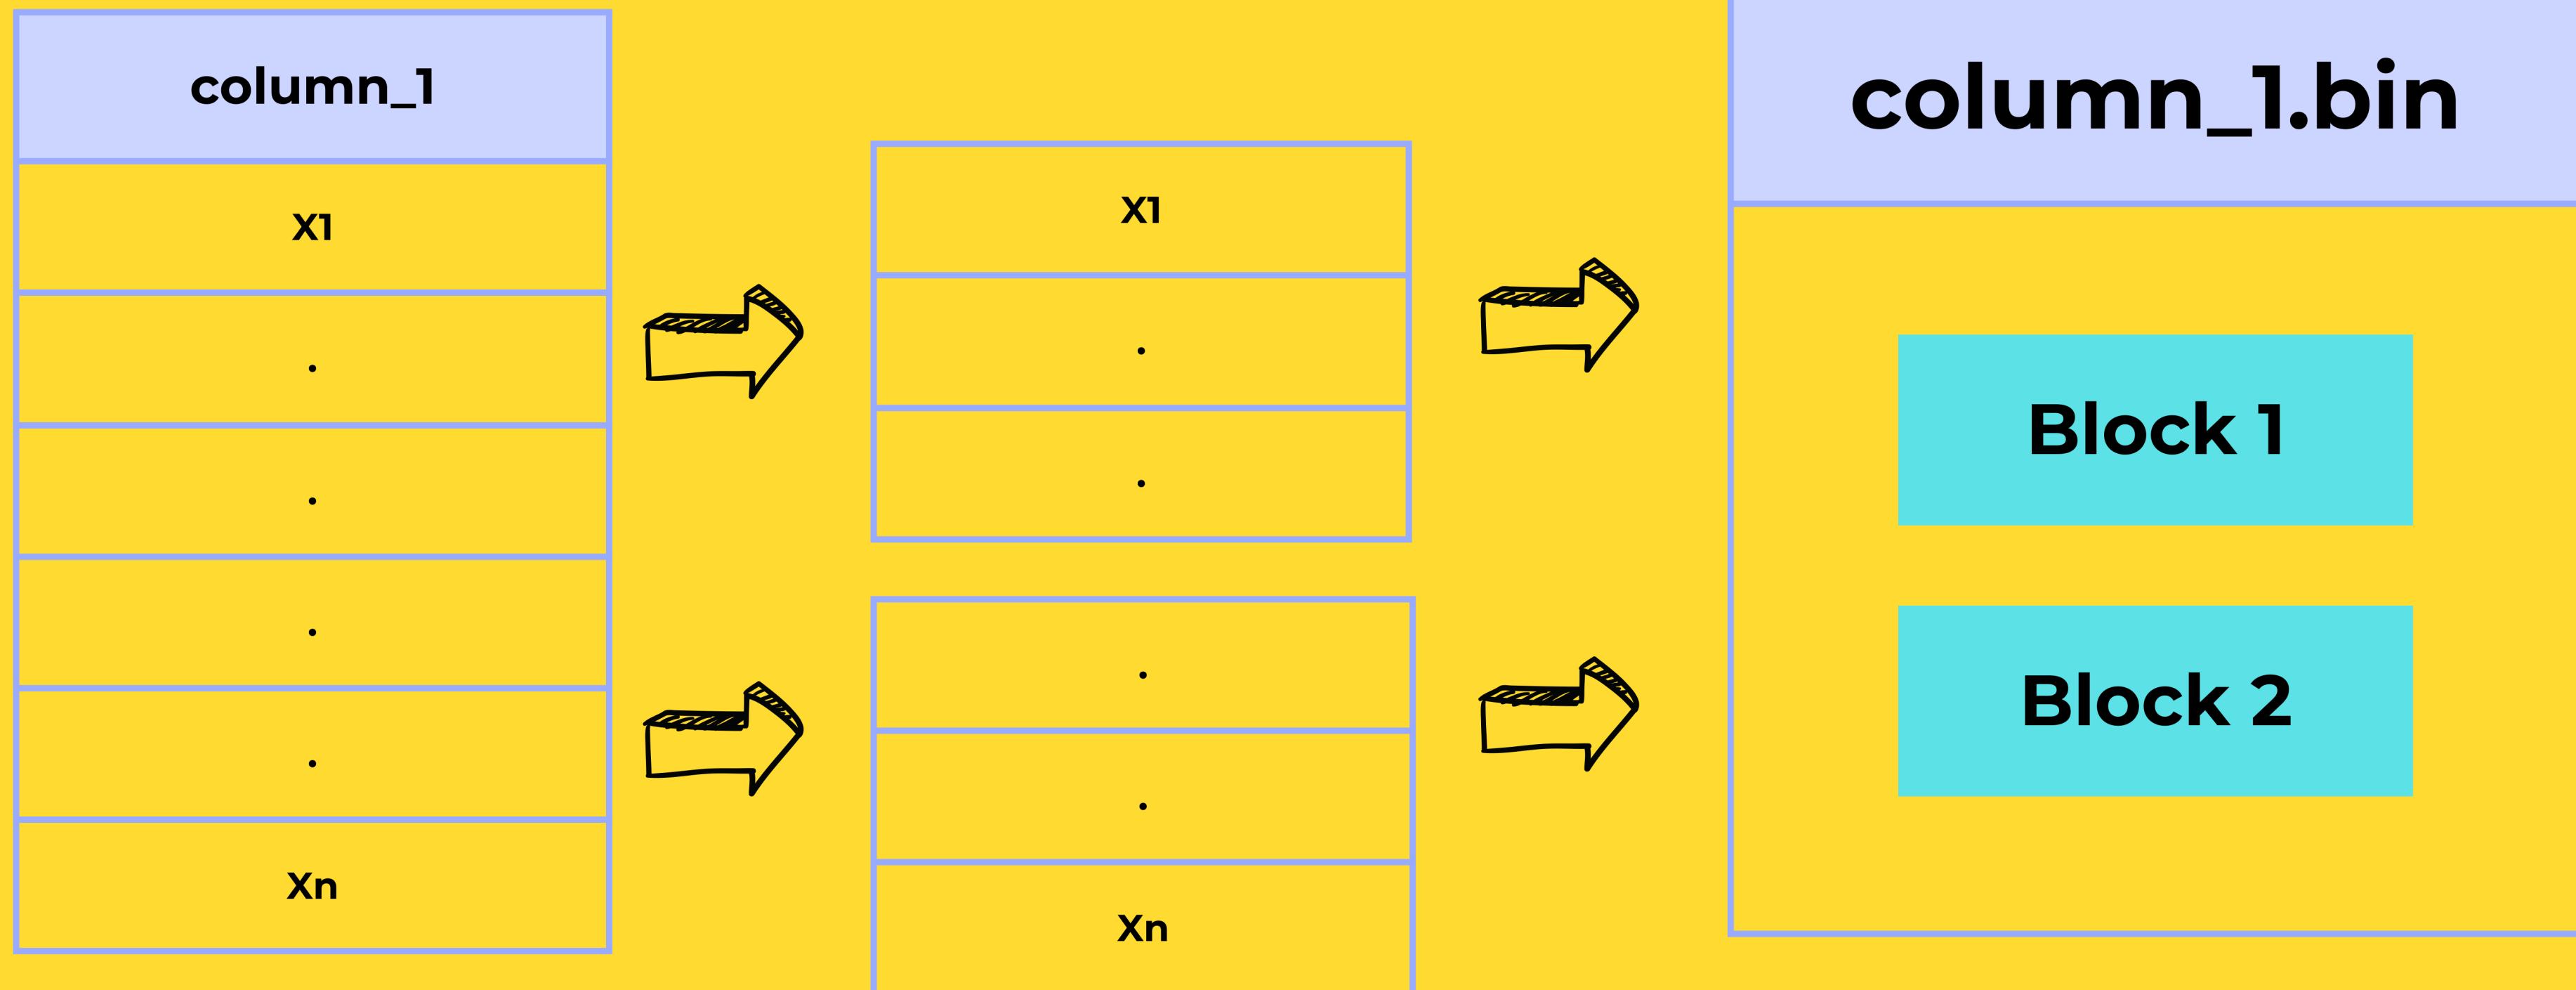
General Purpose Codecs

- LZ4
- ZSTD
- LZ4HC

Special Codecs

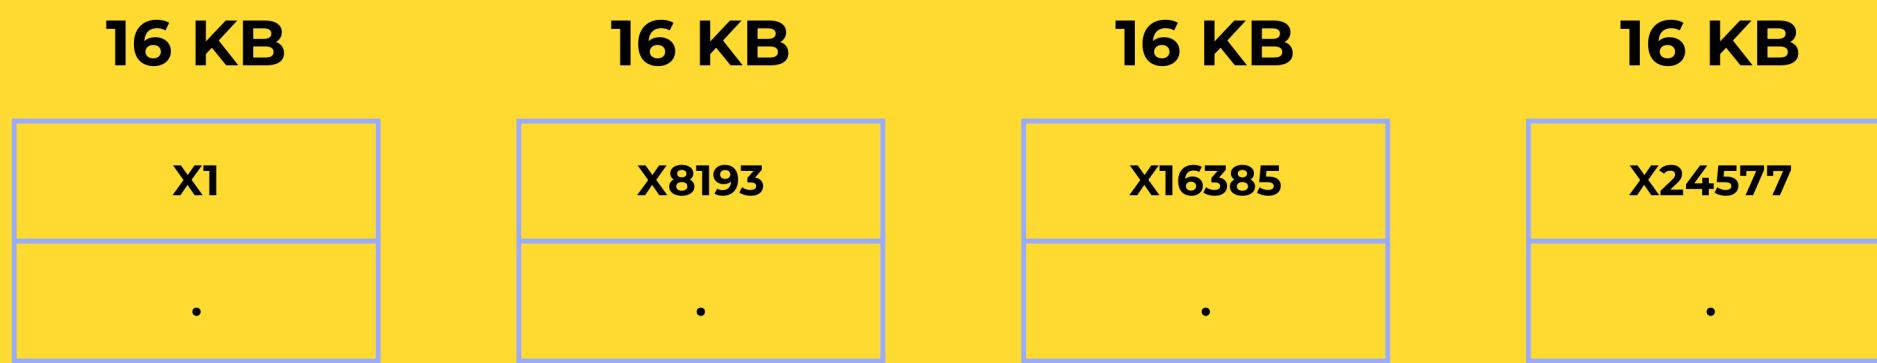
- Delta
- Double Delta
- Gorilla
- GCD
- FPC
- T64

<column>.bin & Blocks



Granules	Blocks
<i>Defined by the number of rows</i>	<i>Defined by the data size</i>
<i>The first row of the primary key for every granule is stored in primary.idx file</i>	<i>Every block of uncompressed data is compressed (by default) and stored in columnar files</i>
<i>index_granularity</i>	<i>max_compress_block_size, min_compress_block_size</i>

Blocks Example - Int16

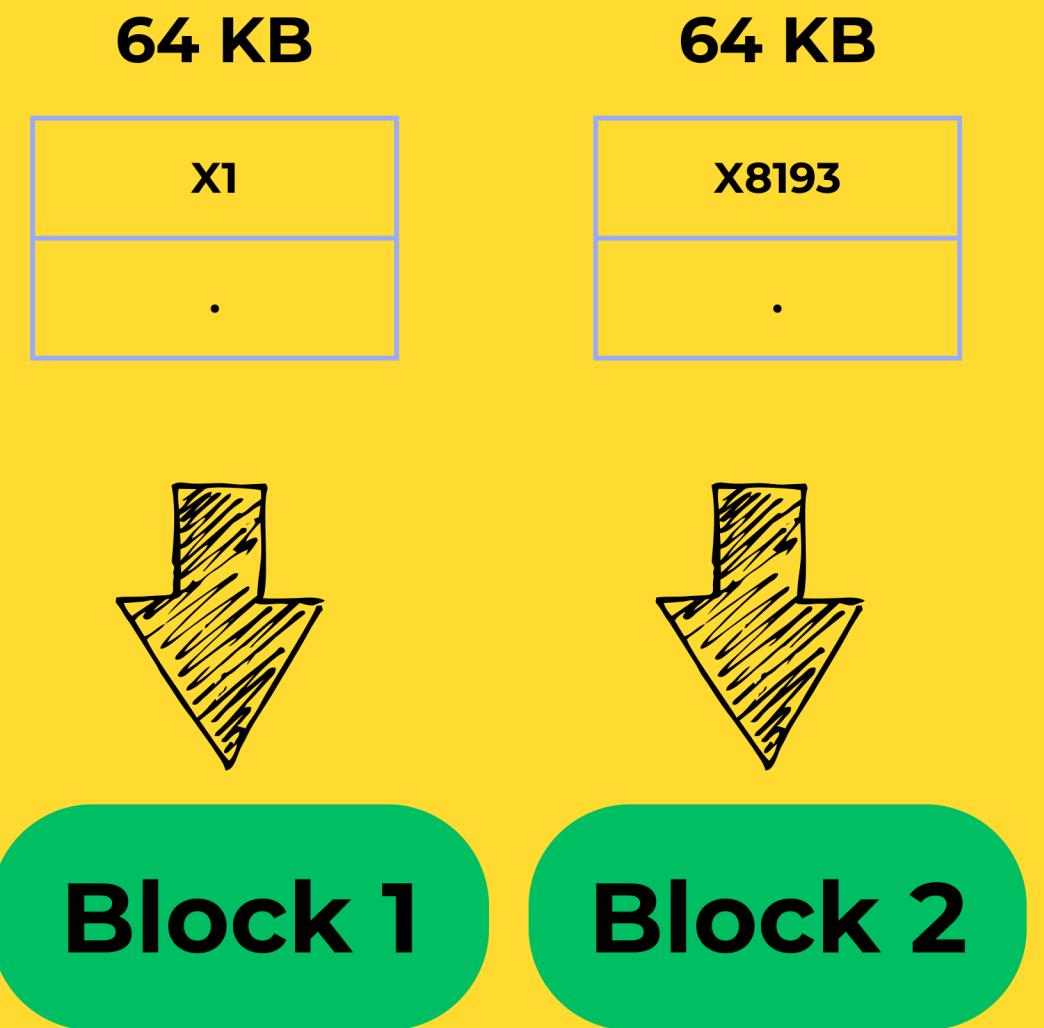


- min_compress_block_size - 64 KB
- max_compress_block_size - 1 MB
- index_granularity - 8192



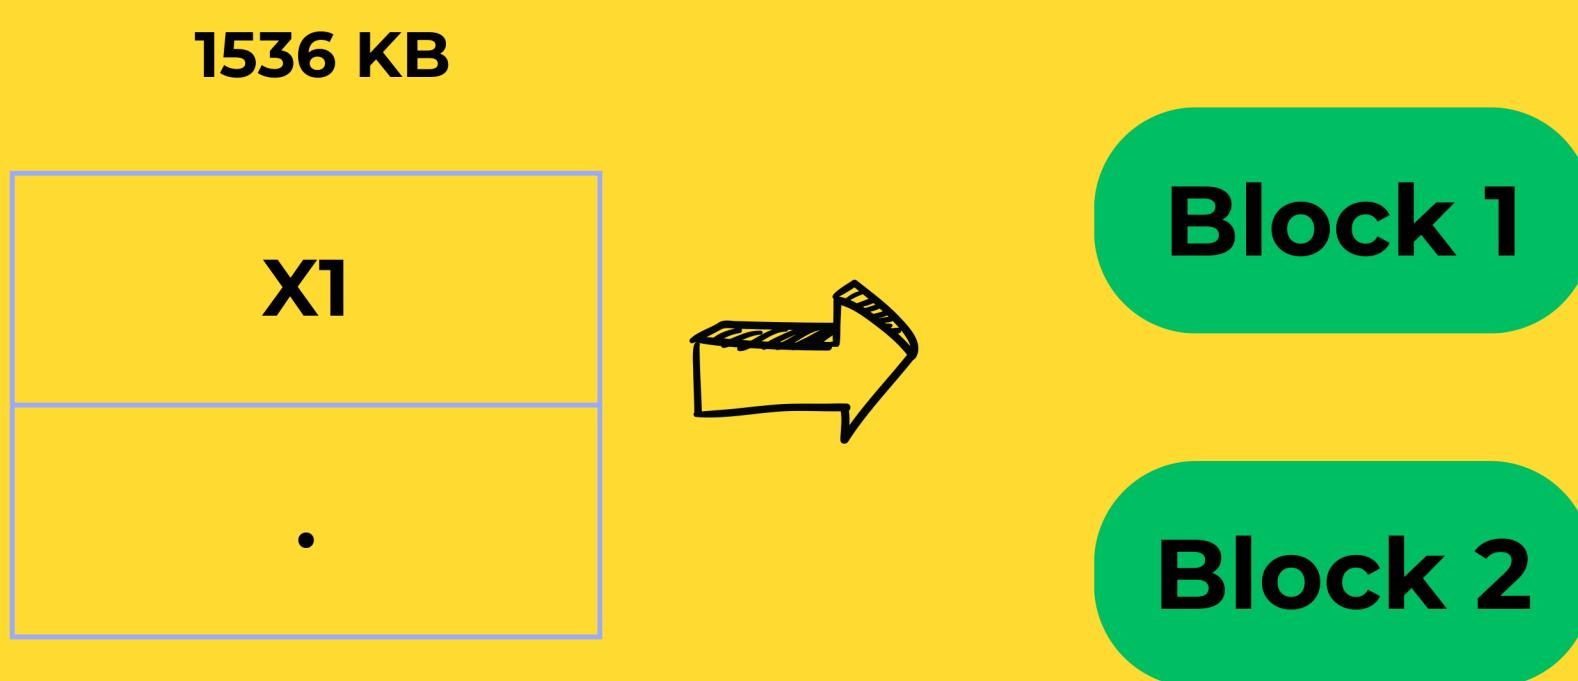
Block 1

Blocks Example - Int64



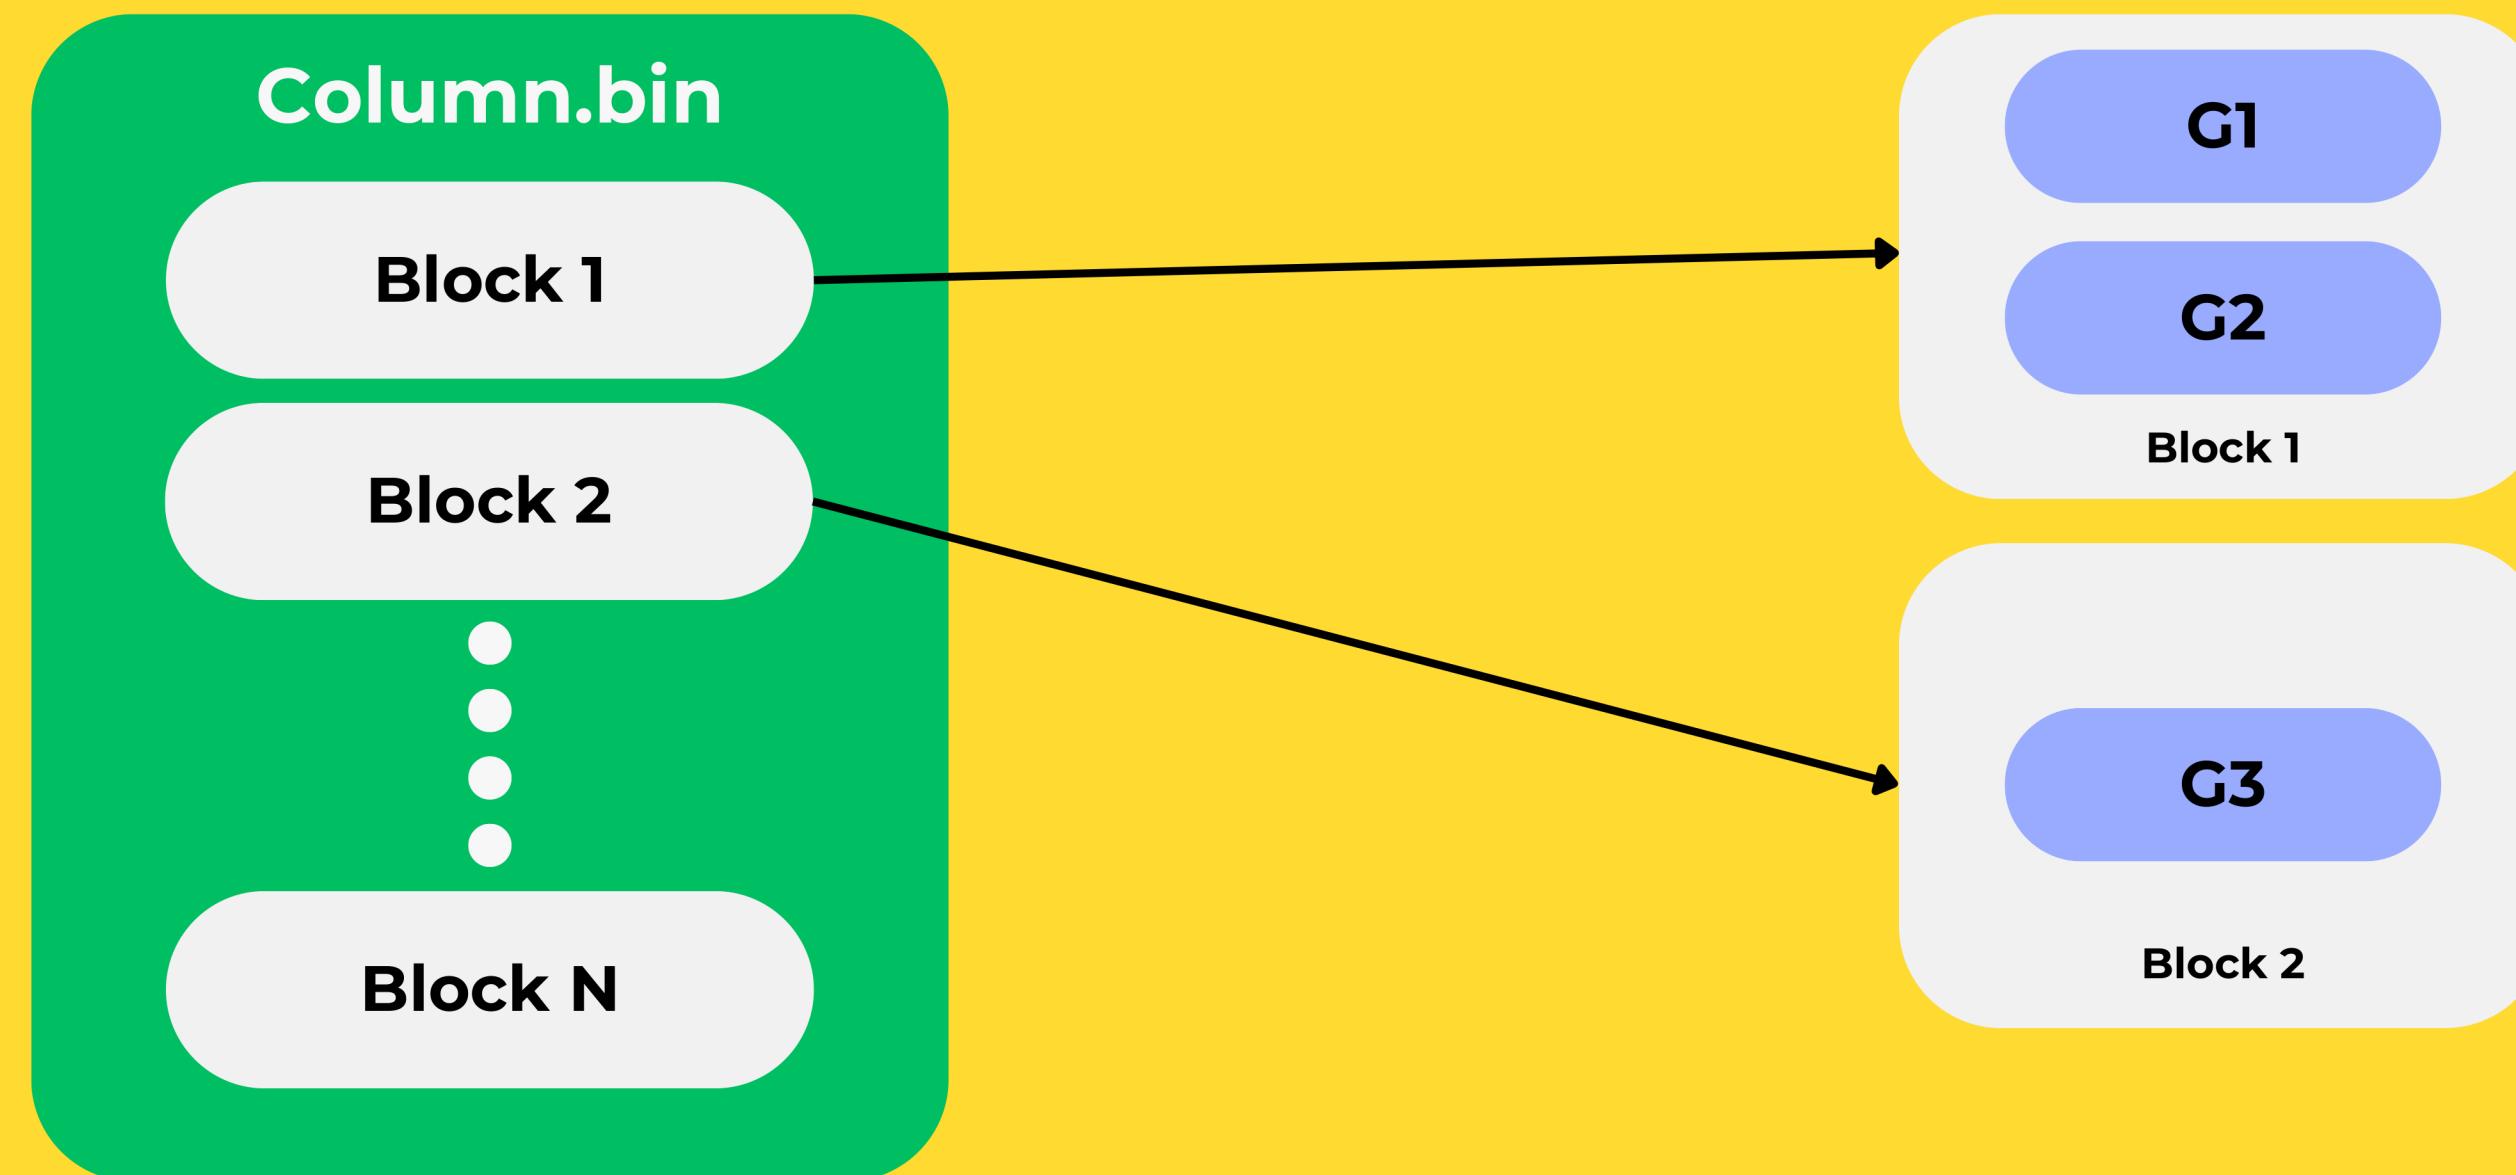
- min_compress_block_size - 64 KB
- max_compress_block_size - 1 MB
- index_granularity - 8192

Blocks Example - String

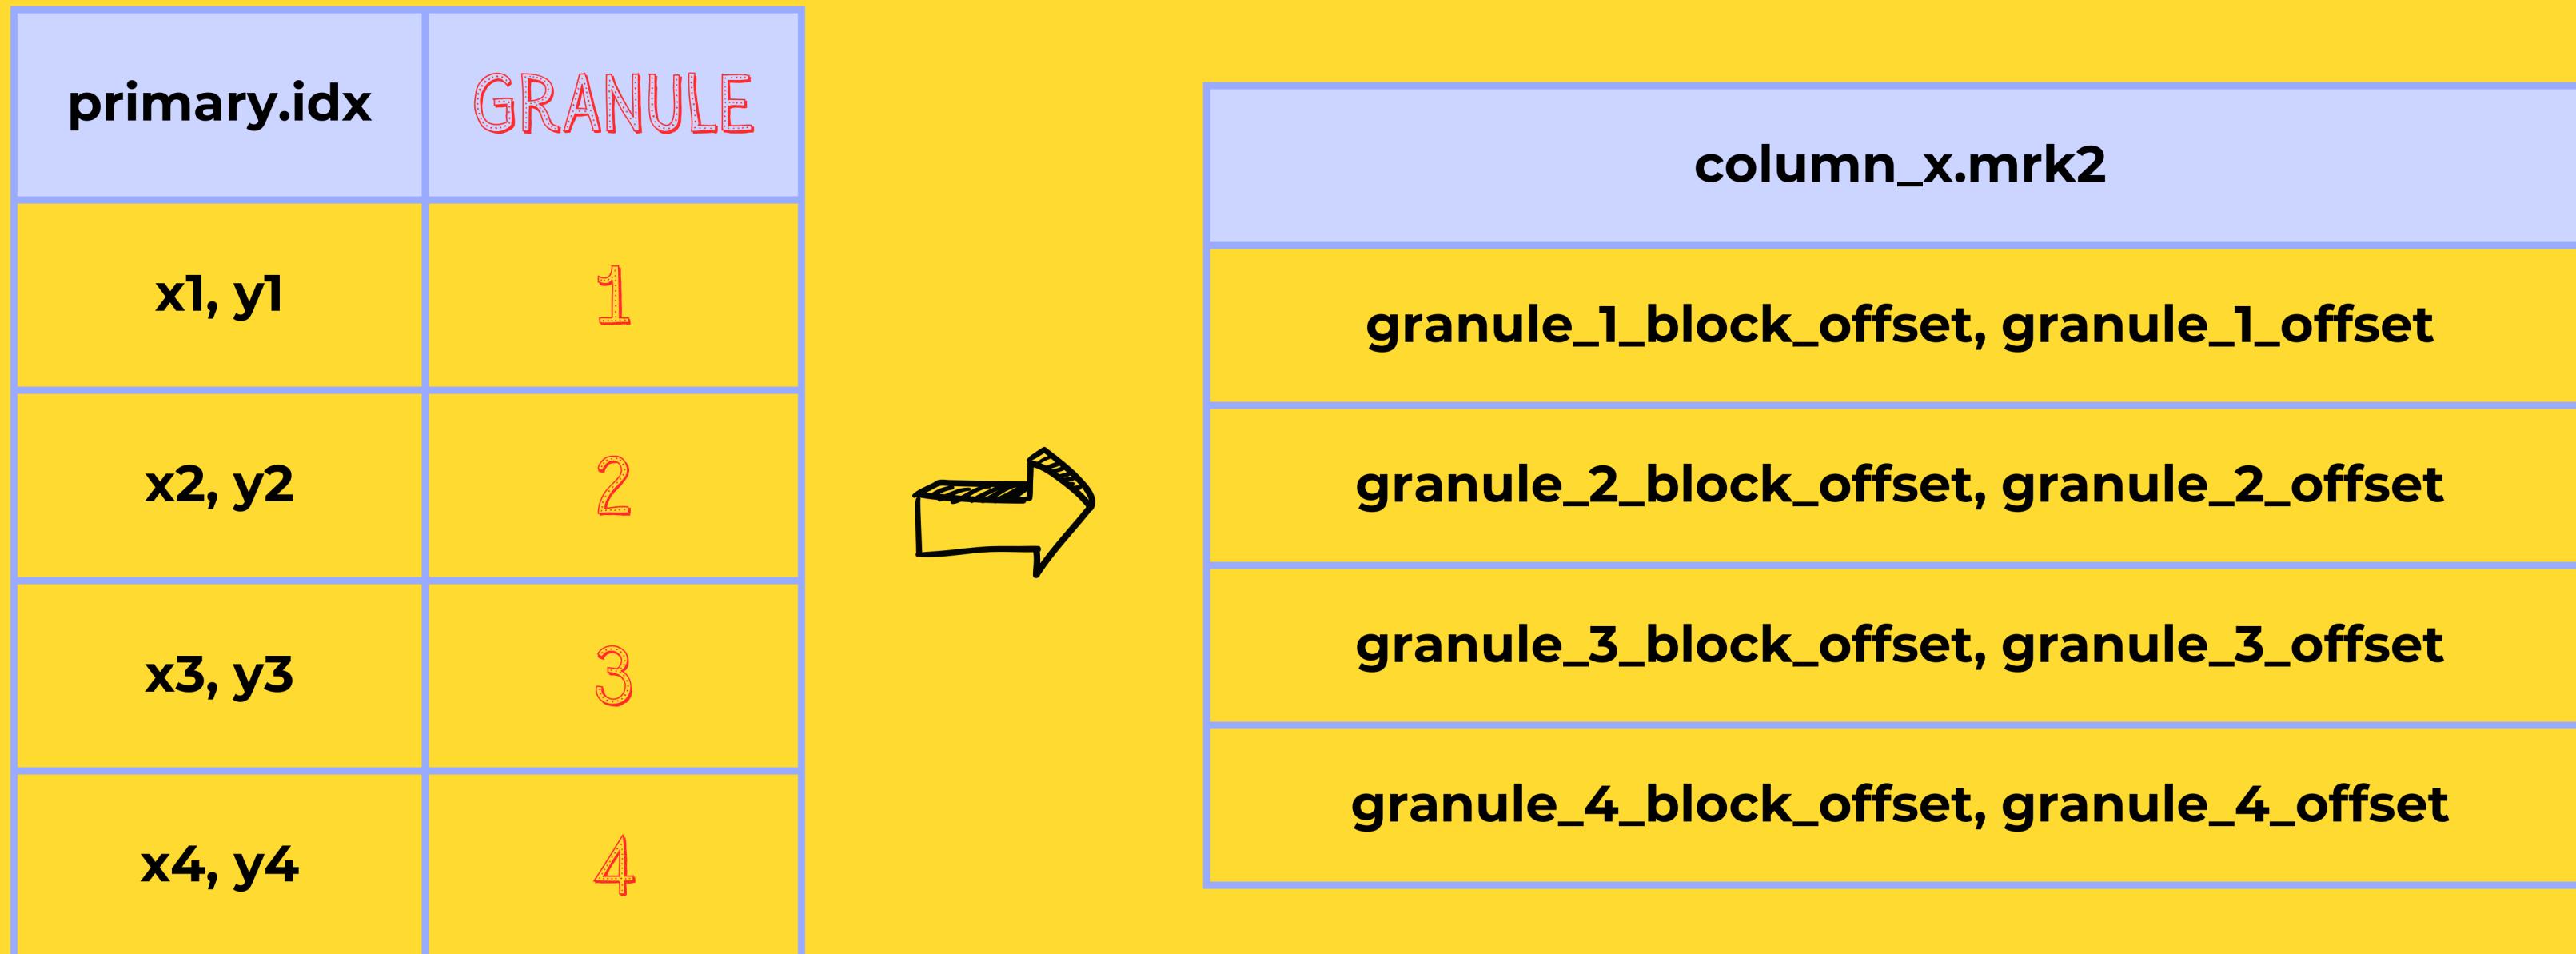


- min_compress_block_size - 64 KB
- max_compress_block_size - 1 MB
- index_granularity - 8192

Offsets in Bin File



Mark File & Offsets

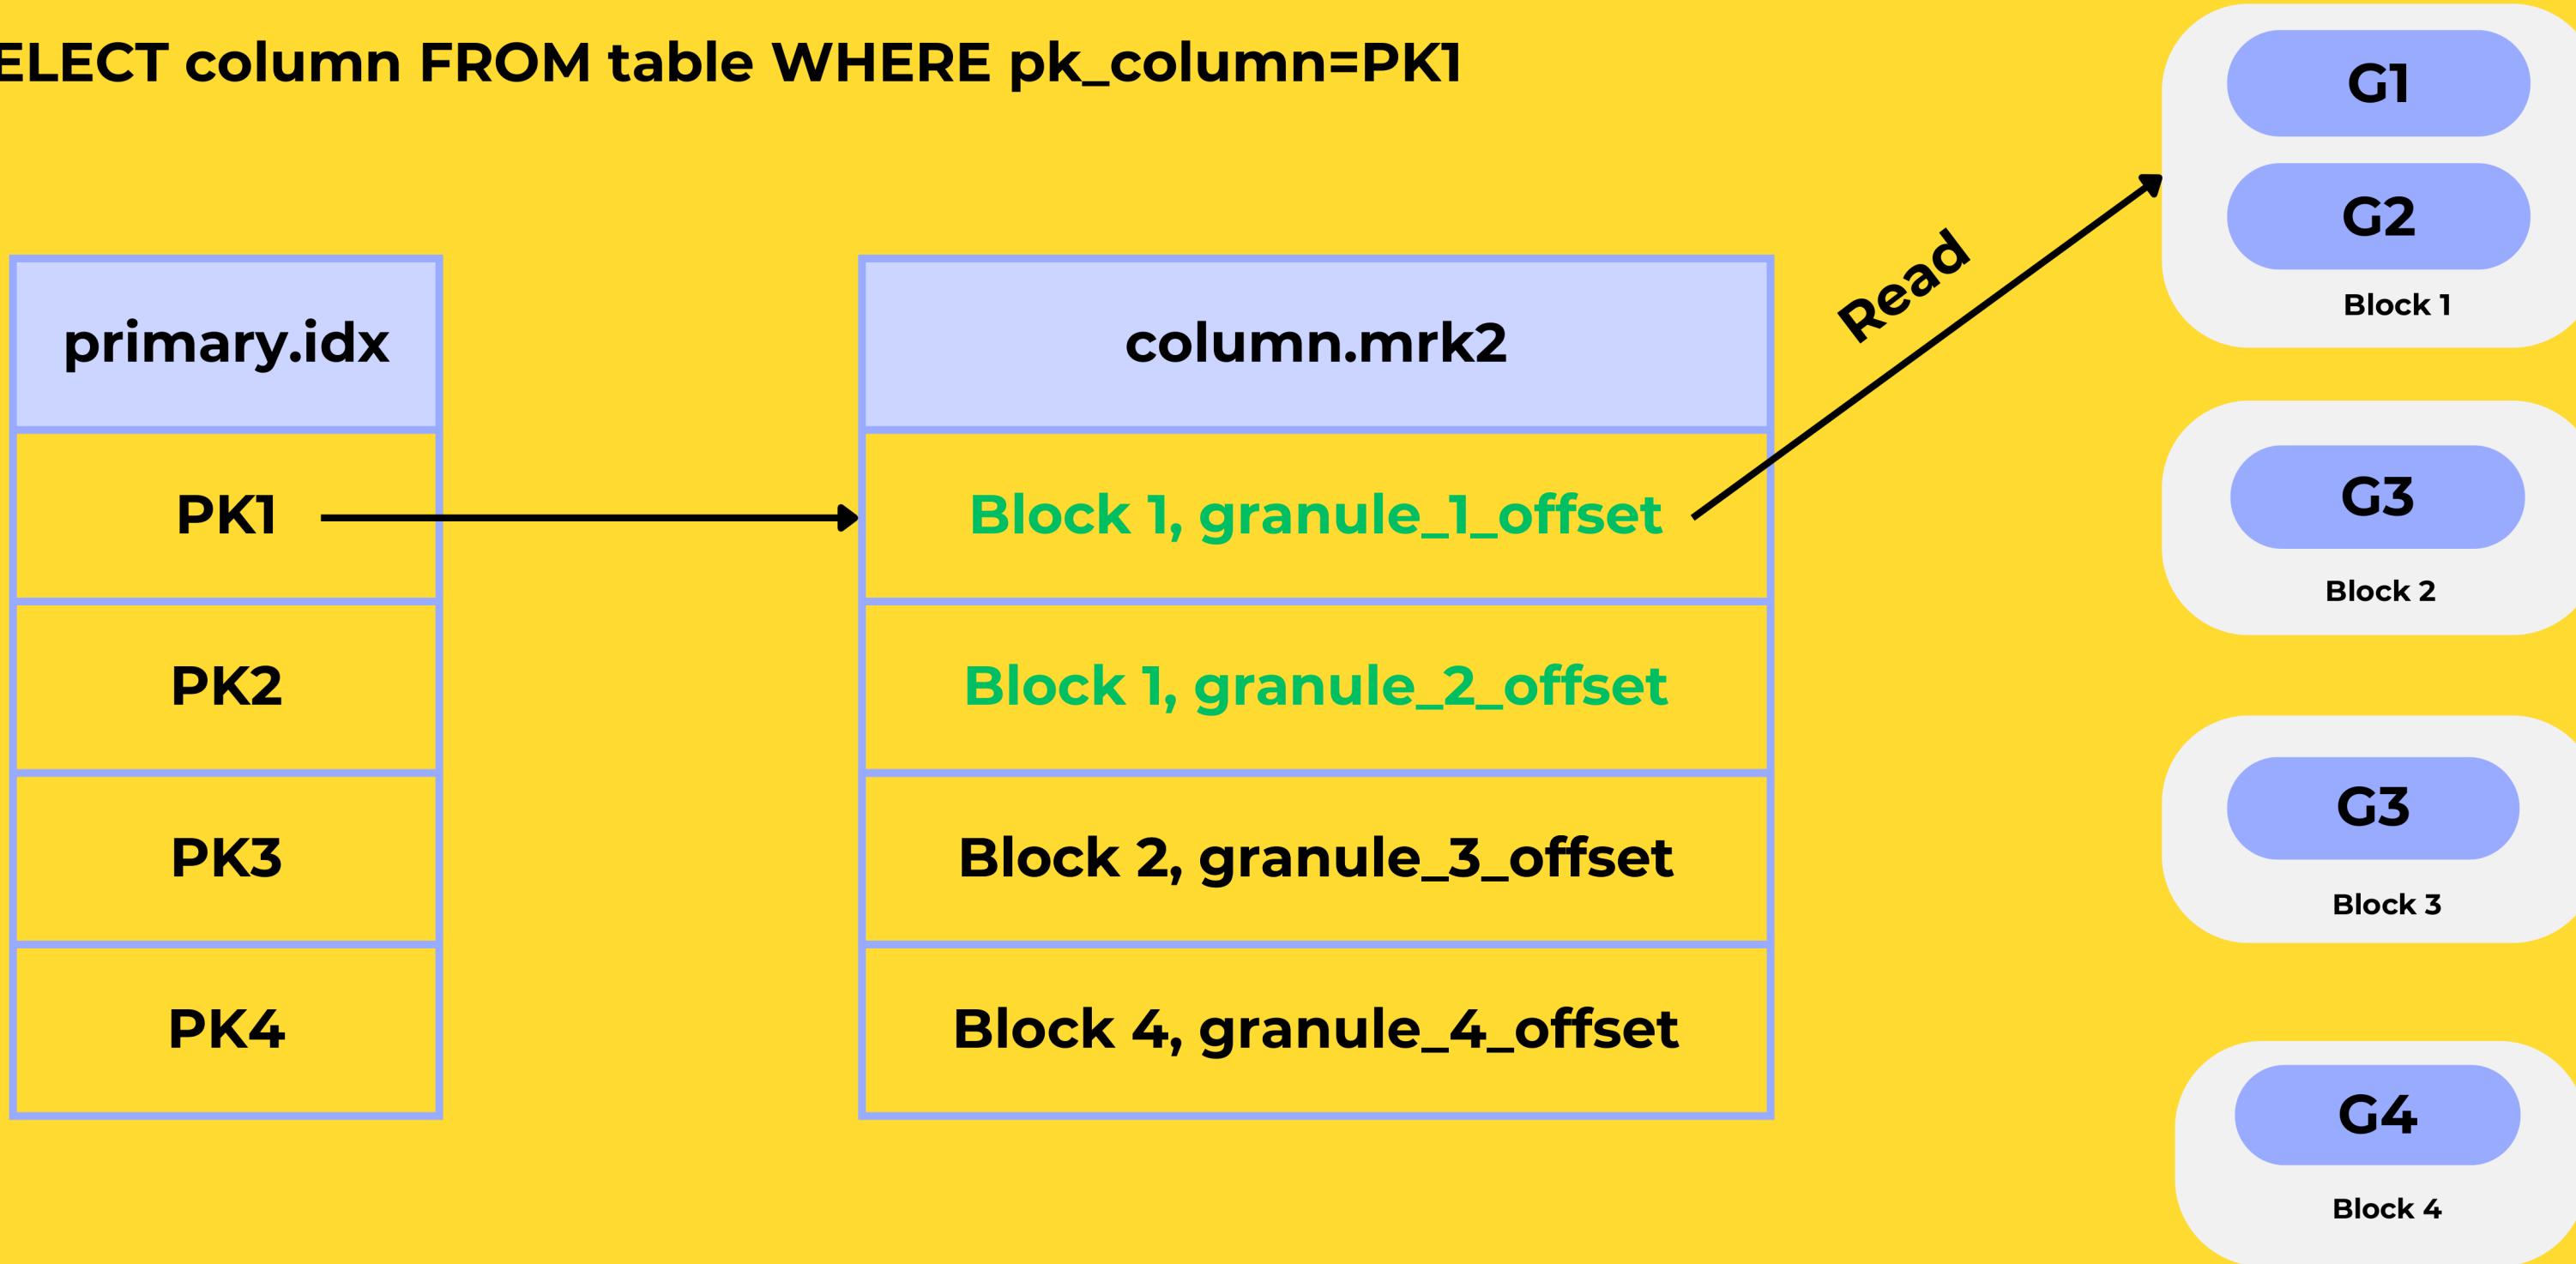


Mark Files

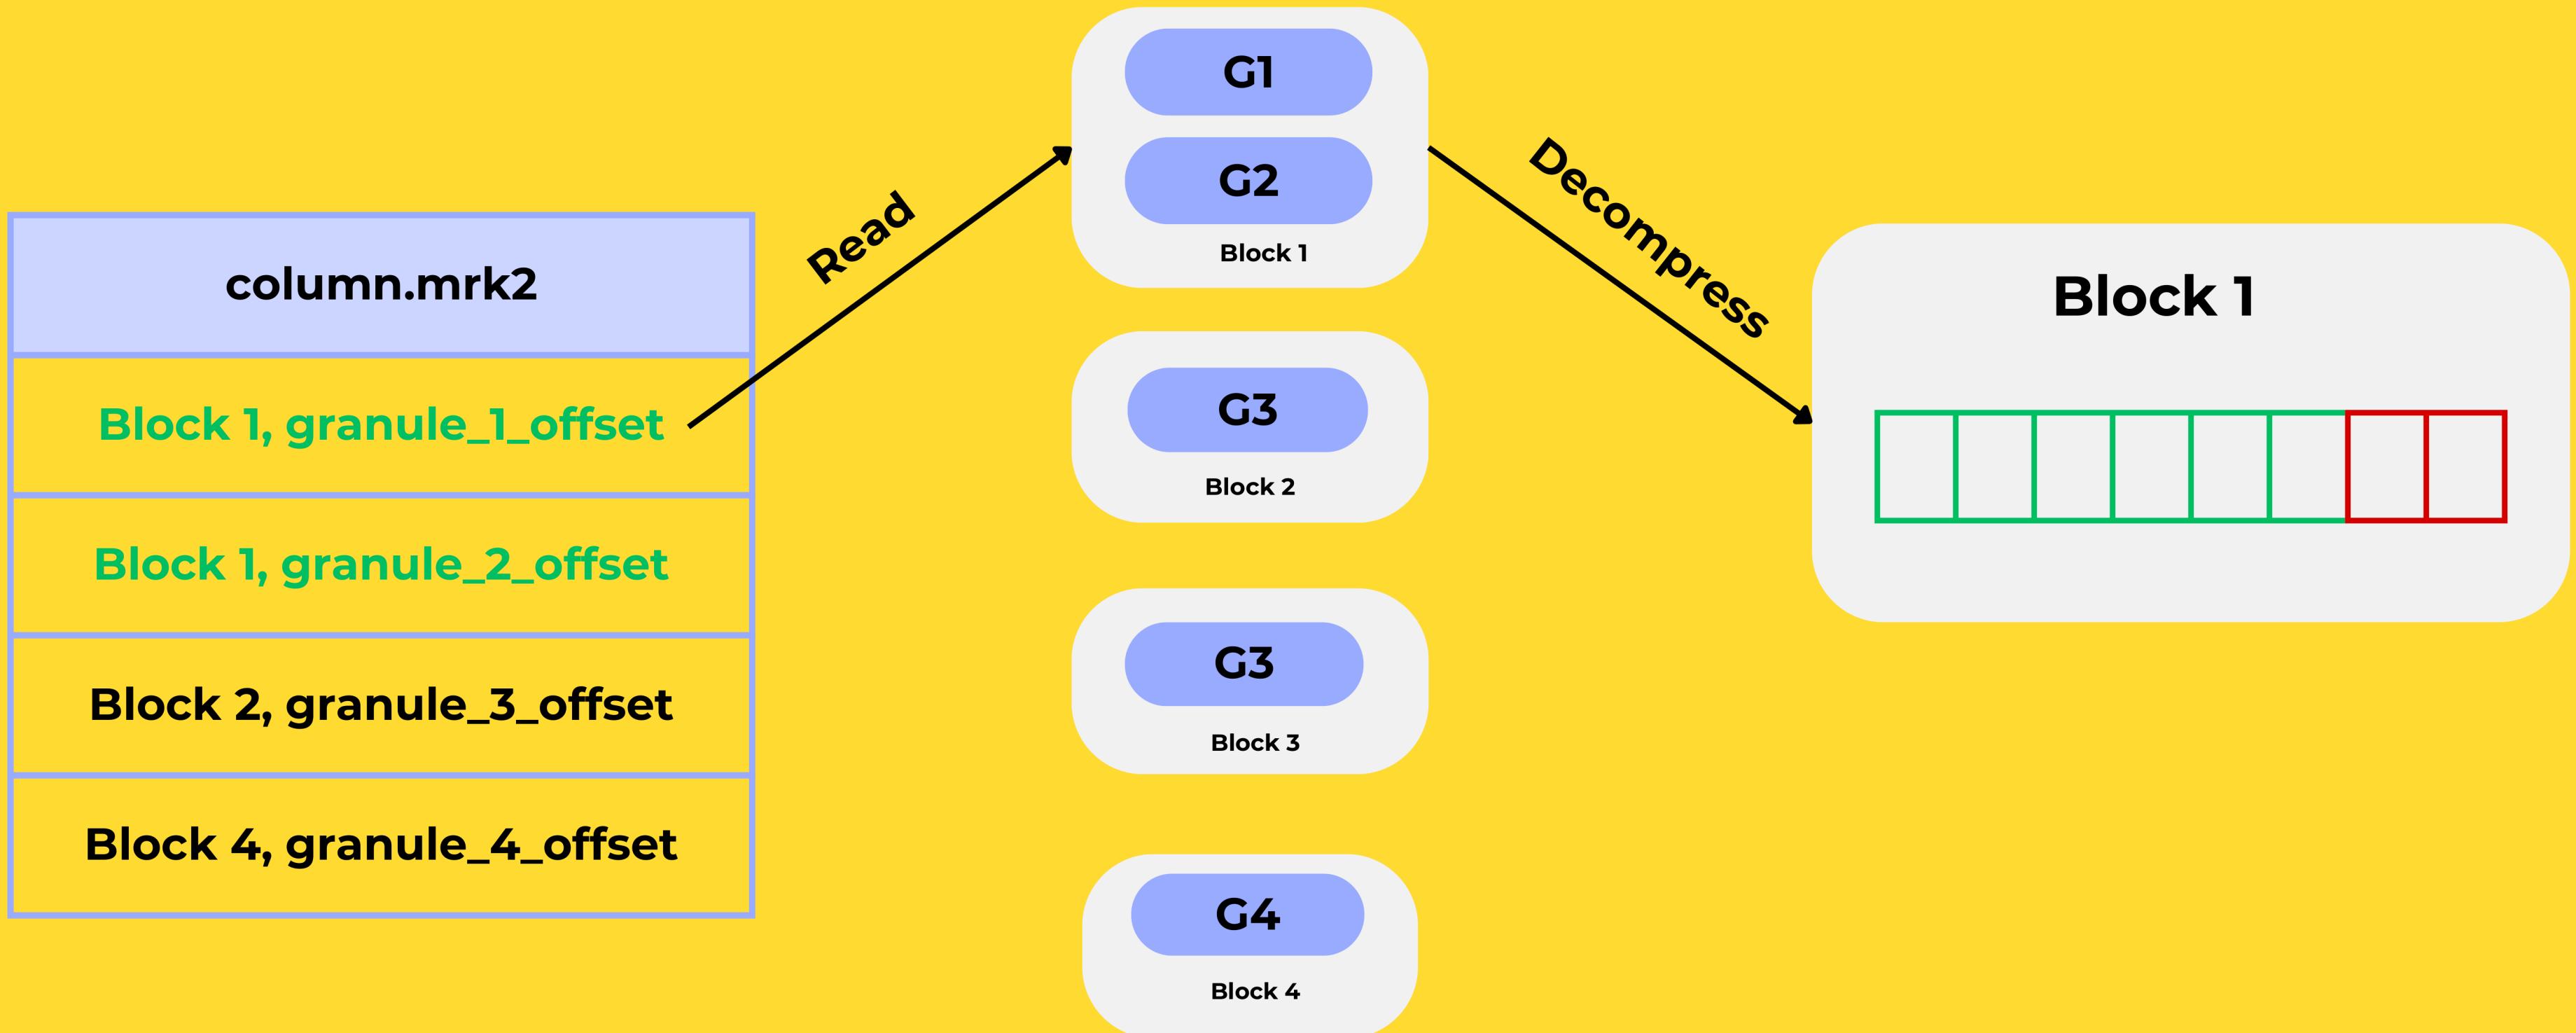
File Type	Data Format	Index Granularity	compress_marks = 1
mrk	Wide	Fixed	cmrk
mrk2	Wide	Adaptive	cmrk2
mrk3	Compact	Fixed or Adaptive	cmrk3

Finding Granule 1

SELECT column FROM table WHERE pk_column=PK1

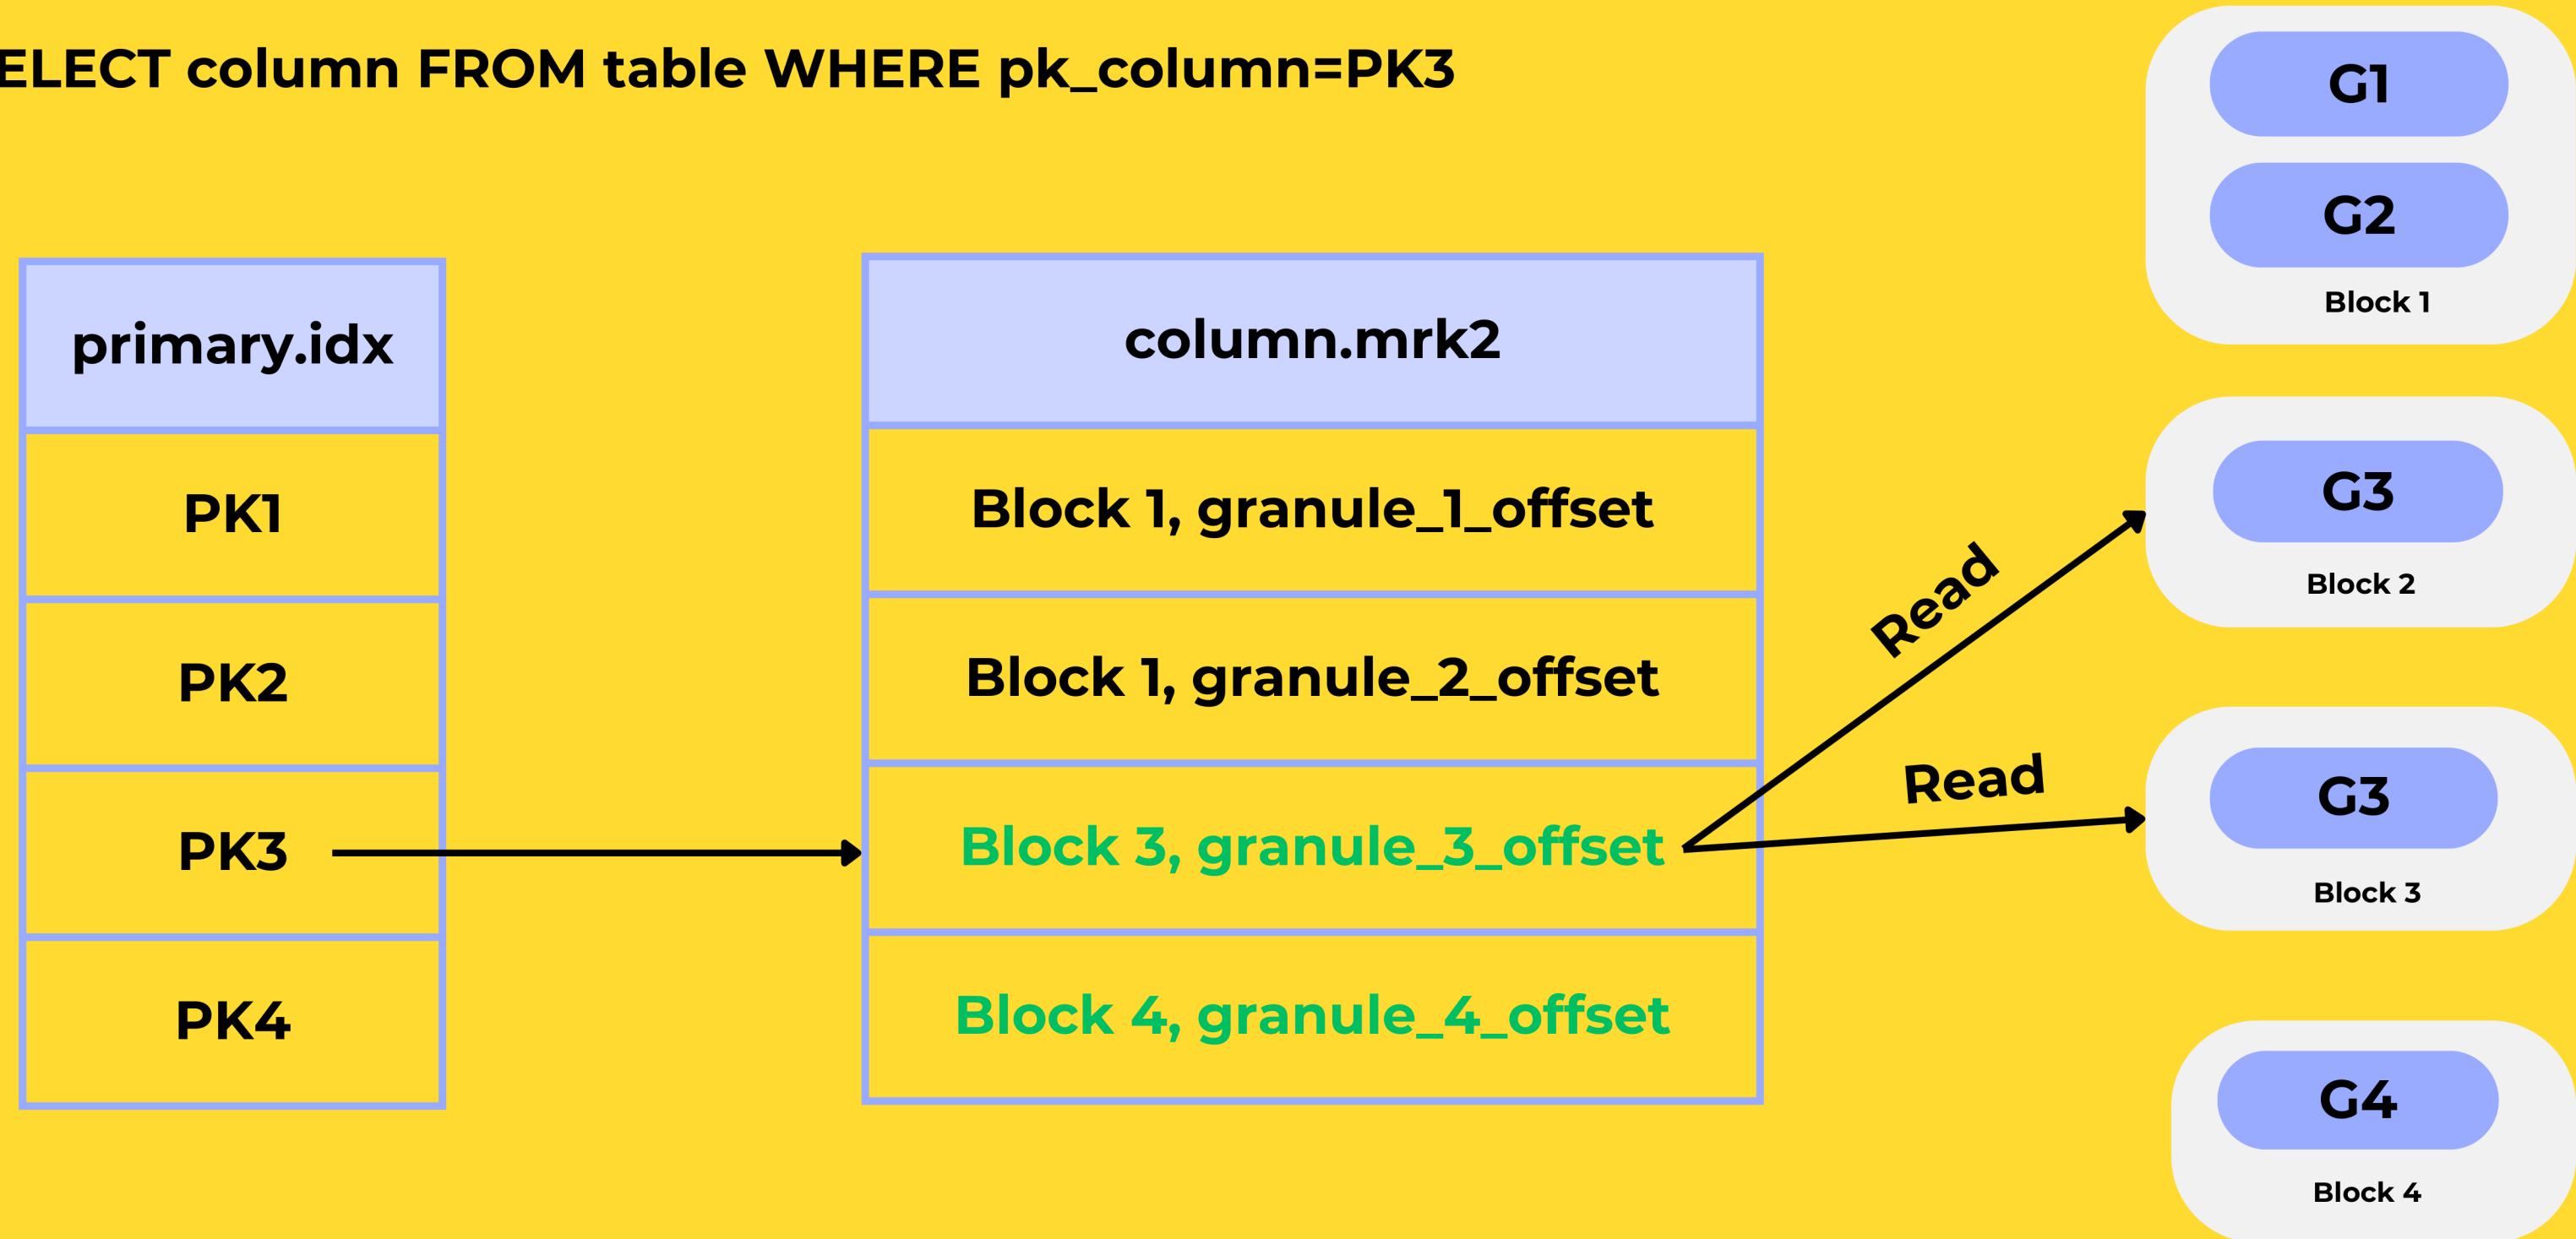


Finding Granule 1

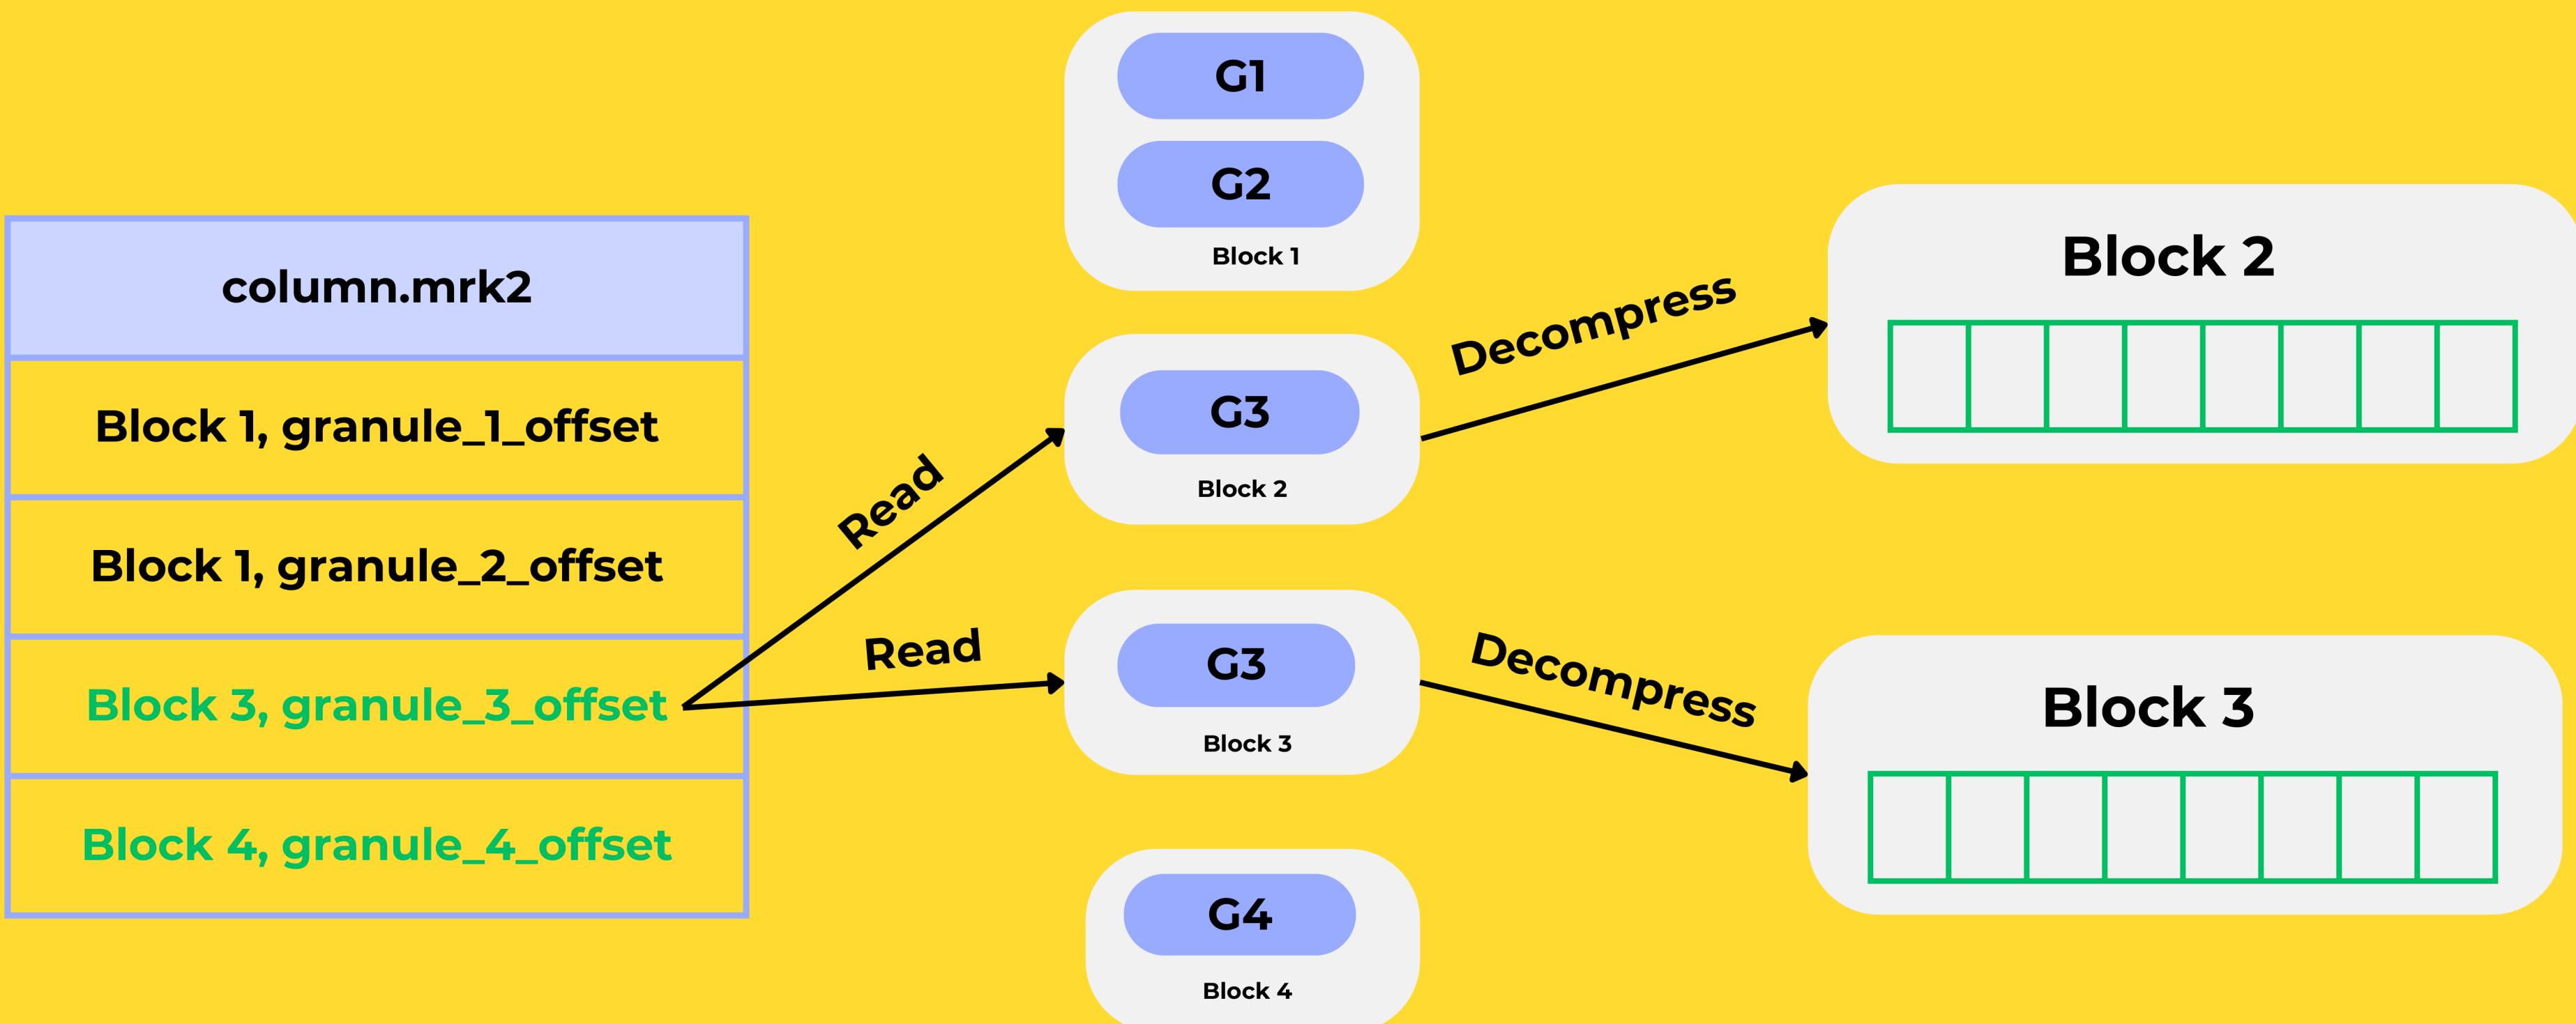


Finding Granule 3

SELECT column FROM table WHERE pk_column=PK3

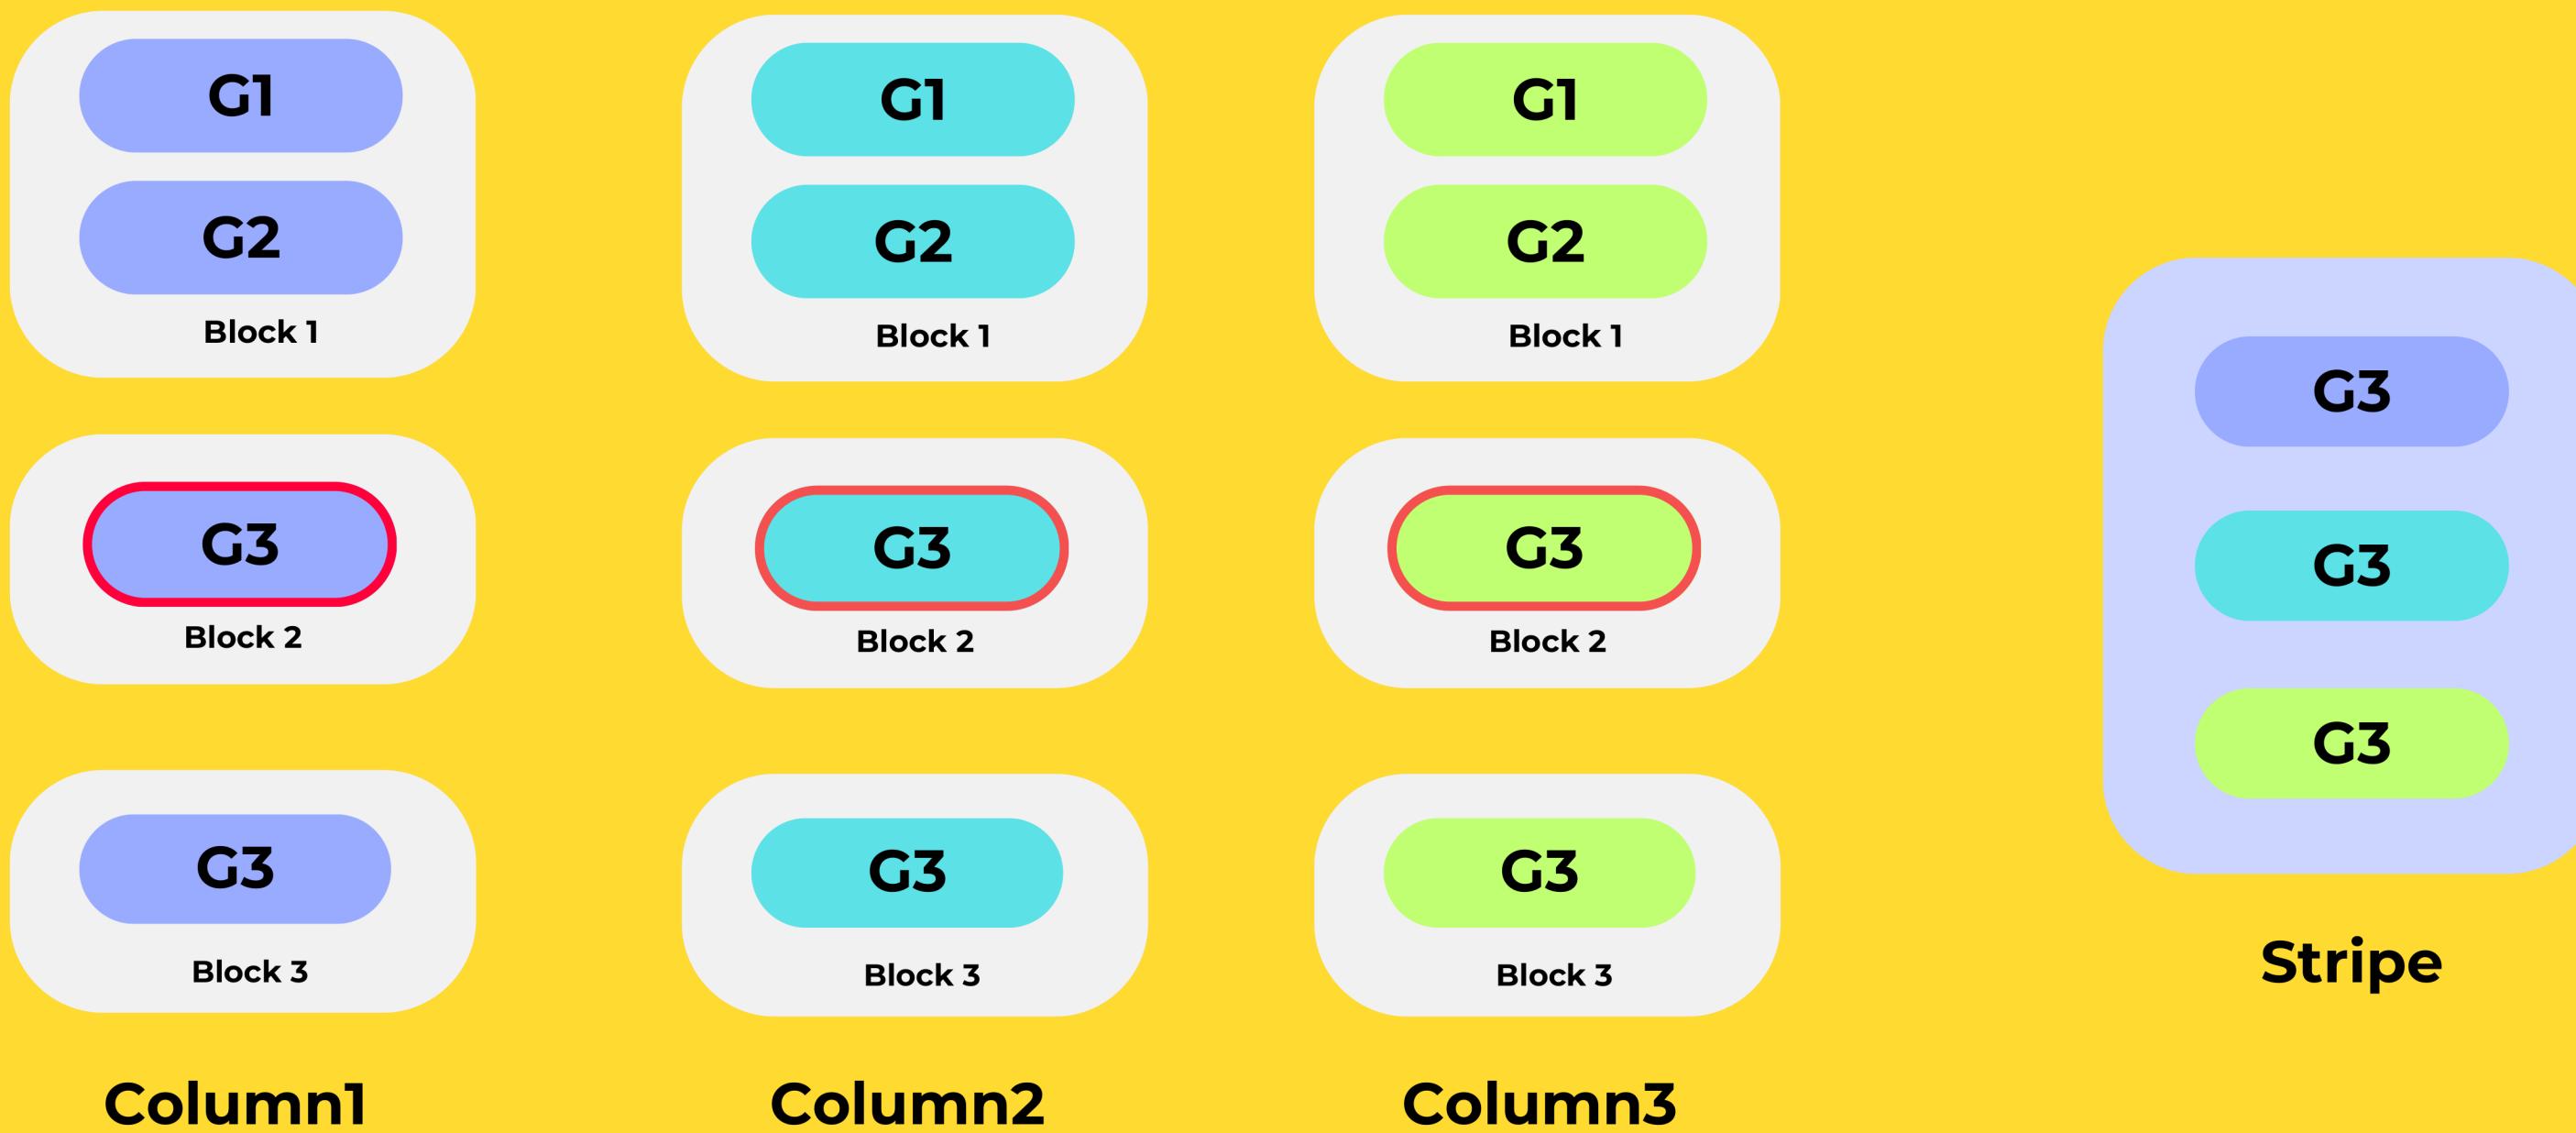


Finding Granule 3



Stripe of granules

SELECT Column1, Column2, Column3 FROM table WHERE pk_column=PK



Primary key vs ORDER BY

PRIMARY KEY

```
CREATE TABLE  
mergetree_example_pk  
(  
    C1 String,  
    C2 Int32,  
    C3 Float64  
    C4 DateTime  
)  
Engine = MergeTree()  
PRIMARY KEY (C1, C2, C3)
```

ORDER BY

```
CREATE TABLE  
mergetree_example_order_by  
(  
    C1 String,  
    C2 Int32,  
    C3 Float64  
    C4 DateTime  
)  
Engine = MergeTree()  
ORDER BY (C1, C2, C3)
```

Primary key vs ORDER BY

Condition	Data Sorted By	Primary Index
PRIMARY KEY (C1, C2)	C1 and C2	(C1, C2)
ORDER BY (C3, C2)	C3 and C2	(C3, C2)
PRIMARY KEY (C1, C2) ORDER BY (C1, C2, C3)	C1, C2 and C3	(C1, C2)

Mergetree CREATE TABLE statement - Syntax

```
CREATE TABLE [IF NOT EXISTS] [database_name.]table_name [ON CLUSTER  
cluster_name]  
(  
    column1 [datatype1] [DEFAULT|MATERIALIZED|ALIAS expr1] [TTL expr1],  
    column2 [datatype2] [DEFAULT|MATERIALIZED|ALIAS expr2] [TTL expr2],  
    ...  
    ...  
)
```

ENGINE = MergeTree()

```
[PARTITION BY expr]  
[ORDER BY expr]  
[PRIMARY KEY expr]  
[SAMPLE BY expr]  
[TTL expr]  
[SETTINGS name=value, ...]
```

Mergetree settings

Setting	Default	Unit
index_granularity	8192	Rows
index_granularity_bytes	1024,000	Bytes
min_index_granularity_bytes	1024	Bytes
merge_with_ttl_timeout	14400	Seconds
write_final_mark	1	NA
merge_max_block_size	8192	Rows

Mergetree settings

Setting	Default	Unit
min_bytes_for_wide_part	1,024,000	Bytes
max_compress_block_size	1,048,576	Bytes
min_compress_block_size	65,536	Bytes
max_partitions_to_read	-1	Unlimited

Mergetree settings

Setting	Default	Unit
max_suspicious_broken_parts	100	No Unit
parts_to_throw_insert	3000	No Unit
parts_to_delay_insert	1000	No Unit
max_delay_to_insert	1	Millisecond
max_parts_in_total	100000	No Unit

Mergetree settings

Setting	Default	Unit
max_suspicious_broken_parts_bytes	1,073,741,824	Bytes
max_files_to_modify_in_alter_columns	75	No Unit
max_files_to_remove_in_alter_columns	50	No Unit
old_parts_lifetime	480	Seconds
clean_deleted_rows	Never	No Unit
max_part_loading_threads	CPU Cores	No Unit

Operations on Partitions

- Freeze and Unfreeze
- Attach and Detach
- Move to table
- Replace
- Update
- Delete
- Drop

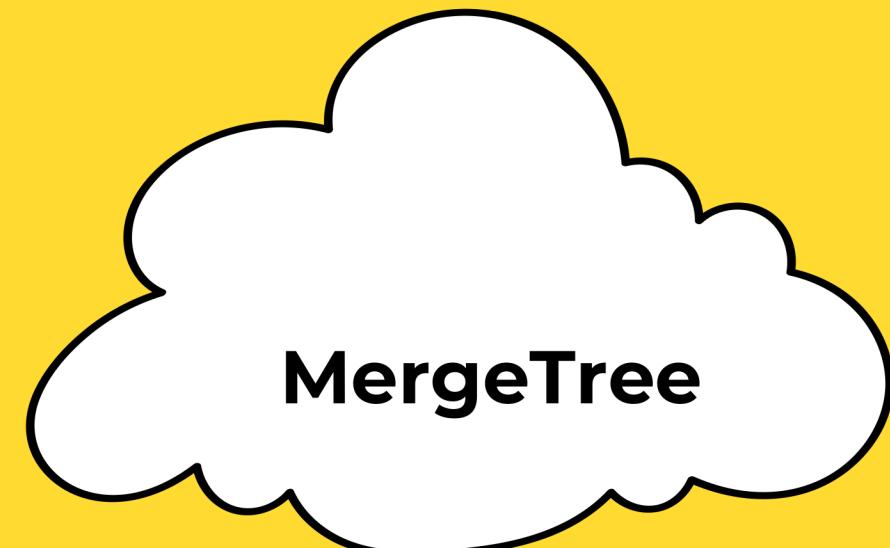
MergeTree Table Engines

ReplacingMerge
Tree

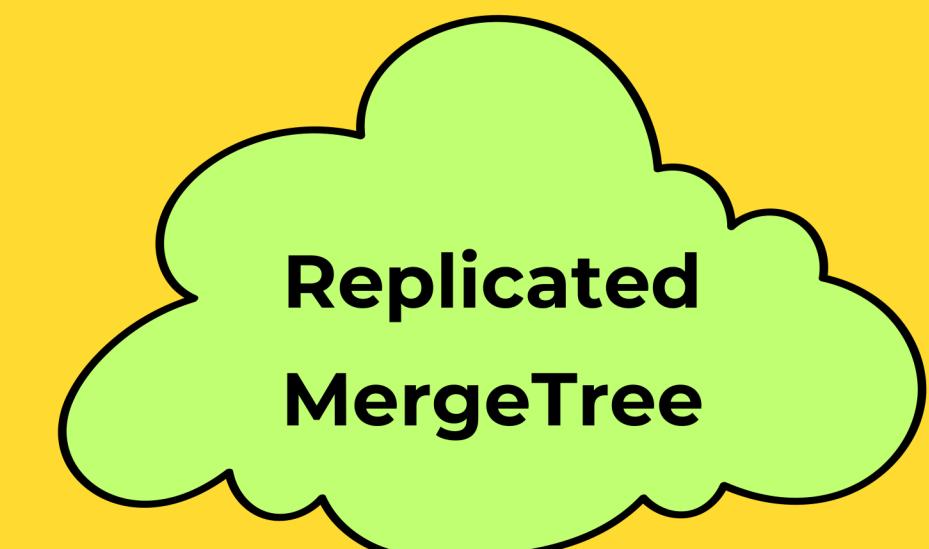
CollapsingMergeTree

SummingMergeTree

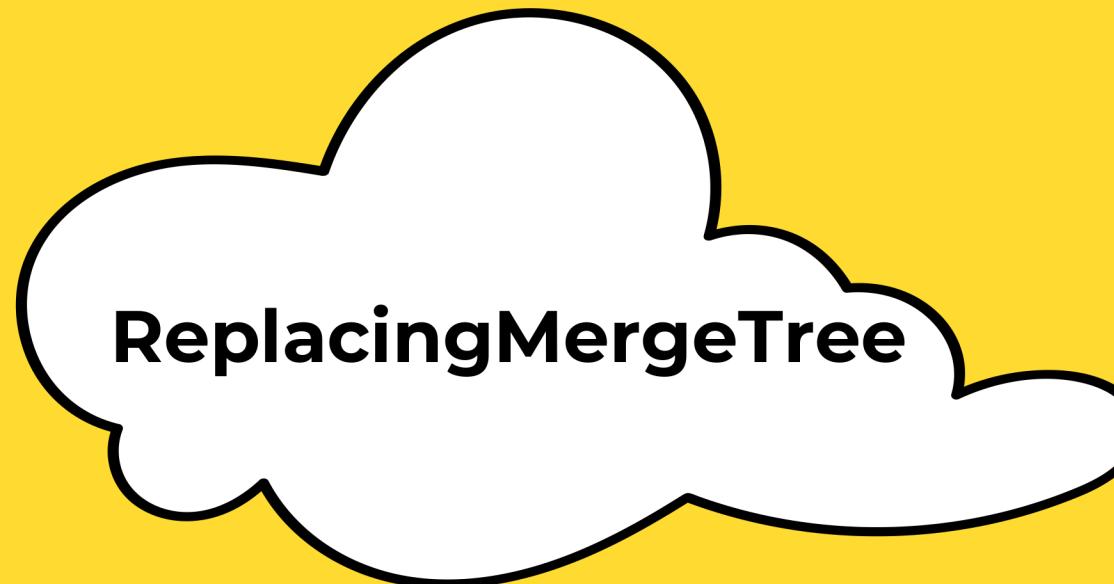
VersionedCollapsing
MergeTree



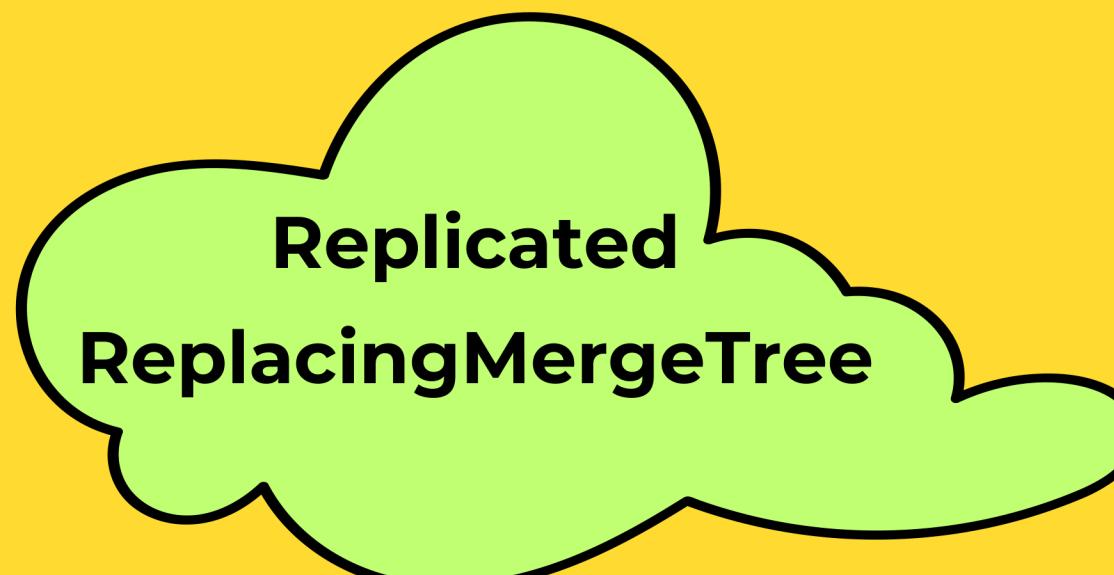
MergeTree



Replicated
MergeTree



ReplacingMergeTree



Replicated
ReplacingMergeTree



SummingMergeTree



Replicated
SummingMergeTree

CollapsingMergeTree

**Versioned
CollapsingMergeTree**

AggregatingMergeTree

**Replicated
CollapsingMergeTree**

**Replicated
Versioned
Collapsing
MergeTree**

**Replicated
AggregatingMergeTree**